
Data File Partitioning

Prof. Hyuk-Yoon Kwon


Data File Partitioning

In this chapter you will learn

- How to improve query performance with data file partitioning
- How to create and populate partitioned tables in Impala and Hive

Data Storage Partitioning (1)

- By default, all files in a data set are stored in a single HDFS directory
 - All files in the directory are read during analysis or processing
 - “Full table scan” each file can be partitioned in HDFS




Data Storage Partitioning (2)

- **Partitioning** subdivides the data

- Analysis can be done on only the relevant subset of data
- Potentially much faster!

- Hadoop partitions using subdirectories



Hadoop Partitioning

- **Partitioning is involved at two phases**
 - Storage – putting the data into correct partition (subdirectory) *loading data*
 - Retrieval – getting the data out of the correct partition based on the query or analysis being done *querying data*
- **Hadoop with built-in support for partitioning**
 - Hive and Impala (covered in next section)
 - Sqoop – When using the `--hive-import` option you can specify flags `--hive-partition-key` and `--hive-partition-value`
- **Other tools can be used to store partitioned data**
 - Spark and MapReduce
 - Flume (at ingestion)

Example: Impala/Hive Partitioning Accounts By State (1)

- Example: accounts is a non-partitioned table

```
CREATE EXTERNAL TABLE accounts(
    cust_id INT,
    fname STRING,
    lname STRING,
    address STRING,
    city STRING,
    state STRING,
    zipcode STRING)
ROW FORMAT DELIMITED
FIELDS TERMINATED BY ','
LOCATION '/loudacre/accounts';
```

Example: Impala/Hive Partitioning Accounts By State (2)

- What if most of Loudacre's analysis on the customer table was done by state? For example:

```
SELECT fname, lname  
      FROM accounts  
 WHERE state='NY' ;
```

- By default, all queries have to scan *all* files in the directory
- Use **partitioning** to store data in *separate files by state*
 - State-based queries scan only the relevant files

Example: Impala/Hive Partitioning Accounts By State (3)

- Create a partitioned table using PARTITIONED BY

```
CREATE EXTERNAL TABLE accounts_by_state(  
    cust_id INT,  
    fname STRING,  
    lname STRING,  
    address STRING,  
    city STRING,  
    state STRING,  
    zipcode STRING)  
PARTITIONED BY (state STRING)  
ROW FORMAT DELIMITED  
FIELDS TERMINATED BY ','  
LOCATION '/loudacre/accounts_by_state';
```

Partition Columns

- The partition column is displayed if you DESCRIBE the table


```
DESCRIBE accounts_by_state;
+-----+-----+-----+
| name      | type      | comment |
+-----+-----+-----+
| cust_id   | int       |          |
| fname     | string    |          |
| lname     | string    |          |
| address   | string    |          |
| city      | string    |          |
| zipcode   | string    |          |
| state     | string    |          |
+-----+-----+-----+
```

A partition column is a “virtual column”; data is not stored in the file

Nested Partitions

- You can also create nested partitions ... no more than 2 levels due to performance

```
... PARTITIONED BY (state STRING, / zipcode STRING)
```



1002794 Joseph Gallardo 1895 Hamilton Street Anchorage
1009777 Tree van Nilson 1331 Village Lane Anchorage
...

1002443 Rachel Crawford 2202 West 10th Street Akiak
1003232 Penny Lane 233 West 5th Street Akiak
...

Impala & Hive \Rightarrow good at handling large files
Too many nested partitions \Rightarrow many small files
performance degraded

Loading Data Into a Partitioned Table

2 options depending on source & nature

- **Dynamic partitioning**

- Impala/Hive add new partitions automatically as needed at load time
 - Data is stored into the correct partition (subdirectory) based on column value

- **Static partitioning**

- You define new partitions using ADD PARTITION
 - When loading data, you specify which partition to store data in

Dynamic Partitioning

- We can create new partitions dynamically from existing data

```
INSERT OVERWRITE TABLE accounts_by_state
  PARTITION(state)
    SELECT cust_id, fname, lname, address,
          city, zipcode, state FROM accounts;
```

- Partitions are automatically created based on the value of the *last column*
 - If the partition does not already exist, it will be created
 - If the partition does exist, it will be overwritten

Practice – Dynamic Partitioning

1. Import the **accounts** table from MySQL directly into the Hive

- Hint1: use Sqoop command
- Hint2: use the option “--hive-import”

2. Create a partitioned table by the **city** from **accounts** table

3. Insert data into the partitioned table from **accounts**

4. Create a partitioned table **accounts_areacode** by the **areacode** from **accounts** table

- areacode could be defined by the first three digits of the phone number

5. Insert data into the partitioned table **accounts_areacode** from **accounts**

- areacode could be obtained by substr(phone_number, 1, 3)

6. Create a nested partitioned table **accounts_nested** by (state, areacode) from accounts table

7. Insert data into the nested partitioned table **accounts_nested** from **accounts**

```
sqoop import --connect jdbc:mysql://localhost/loudacre --username training  
--password training --fields-terminated-by '\t' --table accounts --hive-import  
--warehouse-dir=/user/hive/warehouse  
  
create external table accounts_by_areacode (fname string, lname string,  
address string, city string, zipcode string) partitioned by (areacode  
string) row format delimited fields terminated by ',' location '/loudacre/  
accounts_by_areacode';  
  
insert overwrite table accounts_by_areacode partition (areacode) select  
first_name, last_name, address, city, zipcode, substr(phone_number, 1, 3)  
as areacode from accounts;  
  
create external table accounts_by_nested (fname string, lname string,  
address string, city string, zipcode string) partitioned by (state string,  
areacode string) row format delimited fields terminated by ',' location '/  
loudacre/accounts_by_nested';  
  
insert overwrite table accounts_by_nested partition (state, areacode)  
select first_name, last_name, address, city, zipcode, state,  
substr(phone_number, 1, 3) as areacode from accounts;
```

Static Partitioning Example: Partition Calls by Day (1)

- Loudacre's customer service phone system generates logs detailing calls received
 - Analysts use this data to summarize previous days' calls
 - For example:

```
SELECT event_type, COUNT(event_type)
  FROM call_log
 WHERE call_date = '2014-10-01'
 GROUP BY event_type;
```

Static Partitioning Example: Partition Calls by Day (2)

- Logs are generated daily, e.g.

call-20141001.log

```
19:45:19,312-555-7834,CALL_RECEIVED  
19:45:23,312-555-7834,OPTION_SELECTED,Shipping  
19:46:23,312-555-7834,ON_HOLD  
19:47:51,312-555-7834,AGENT_ANSWER,Agent ID N7501  
19:48:37,312-555-7834,COMPLAINT,Item not received  
19:48:41,312-555-7834,CALL_END,Duration: 3:22  
...
```

call-20141002.log

```
03:45:01,505-555-2345,CALL_RECEIVED  
03:45:09,505-555-2345,OPTION_SELECTED,Billing  
03:56:21,505-555-2345,AGENT_ANSWER,Agent ID A1503  
03:57:01,505-555-2345,QUESTION  
...
```

Static Partitioning Example: Partition Calls by Day (3)

- In the previous example, existing data was partitioned dynamically based on a column value
- This time we use static partitioning
 - Because the data files do not include the partitioning data

Static Partitioning Example: Partition Calls by Day (4)

- The partitioned table is defined the same way

```
CREATE TABLE call_logs (
    call_time STRING,
    phone STRING,
    event_type STRING,
    details STRING)
PARTITIONED BY (call_date STRING)
ROW FORMAT DELIMITED
FIELDS TERMINATED BY ',';
```

Loading Data Into Static Partitions (1)

- With static partitioning, you **create new partitions as needed**
- e.g. For each new day of call log data, add a partition:

```
ALTER TABLE call_logs  
ADD PARTITION (call_date='2014-10-02');
```

- This command
 - Adds the partition to the table's metadata
 - Creates subdirectory
`/user/hive/warehouse/call_logs/
call_date=2014-10-02`

Loading Data Into Static Partitions (2)

- Then load the day's data into the correct partition

```
LOAD DATA INPATH '/mystaging/call-20141002.log'  
  INTO TABLE call_logs  
  PARTITION(call_date='2014-10-02');
```

- This command moves the HDFS file `call-20141002.log` to the partition subdirectory
- To overwrite all data in a partition

```
LOAD DATA INPATH '/mystaging/call-20141002.log'  
  INTO TABLE call_logs OVERWRITE  
  PARTITION(call_date='2014-10-02');
```

Practice – Static Partitioning

1. Create the following partitioned table

```
create external table temp(id int, name string, sal int)
partitioned by(city string)
location '/loudacre/test';
```

2. Make static partitions for city='hyd' and city='sec' using alter table

3. Make two sample text files that contains the following contents in HDFS

[test1.txt]

```
1 ravi 100 hyd
2 krishna 200 hyd
```

[test2.txt]

```
3 fff 300 sec
```

```
create external table temp (id int, name string, sal int) partitioned by (city
string) location '/loudacre/test';
```

```
alter table temp add partition (city='hyd');
```

```
alter table temp add partition (city='sec');
```

```
load data inpath '/loudacre/test1.txt' into table temp partition (city='hyd');
```

```
load data inpath '/loudacre/test2.txt' into table temp partition (city='sec');
```

4. Load two sample files to the table for the created static partitions