

Some useful Colab keyboard shorcuts

- Show keyboard shortcuts: `ctrl+M+H`
-

We'll begin this lecture with a practical example of a neural network to understand essential components and concepts.

A first look at a neural network

- An example of a neural network to learn to classify handwritten digits
- We will use the MNIST dataset.
 - 28*28 grayscale images, 10 categories
 - 60,000 training images, 10,000 test images
 - Refer to <http://yann.lecun.com/exdb/mnist/>
 - Solving MNIST is like the "hello world" of deep learning
- Note
 - In ML, a category in a classification problem is called a *class*.
 - Data points are called *samples*.
 - The class associated with a specific sample is called a *label, target* or *ground truth*.
- First, let's look into the MNIST dataset.

```
In [1]: import torch  
torch.__version__
```

```
Out[1]: '2.8.0+cu126'
```

```
In [2]: !pip install --upgrade torch
```

```
Requirement already satisfied: torch in /usr/local/lib/python3.12/dist-packages (2.8.0+cu126)
Requirement already satisfied: filelock in /usr/local/lib/python3.12/dist-packages (from torch) (3.19.1)
Requirement already satisfied: typing-extensions>=4.10.0 in /usr/local/lib/python3.12/dist-packages (from torch) (4.15.0)
Requirement already satisfied: setuptools in /usr/local/lib/python3.12/dist-packages (from torch) (75.2.0)
Requirement already satisfied: sympy>=1.13.3 in /usr/local/lib/python3.12/dist-packages (from torch) (1.13.3)
Requirement already satisfied: networkx in /usr/local/lib/python3.12/dist-packages (from torch) (3.5)
Requirement already satisfied: jinja2 in /usr/local/lib/python3.12/dist-packages (from torch) (3.1.6)
Requirement already satisfied: fsspec in /usr/local/lib/python3.12/dist-packages (from torch) (2025.3.0)
Requirement already satisfied: nvidia-cuda-nvrtc-cu12==12.6.77 in /usr/local/lib/python3.12/dist-packages (from torch) (12.6.77)
Requirement already satisfied: nvidia-cuda-runtime-cu12==12.6.77 in /usr/local/lib/python3.12/dist-packages (from torch) (12.6.77)
Requirement already satisfied: nvidia-cuda-cupti-cu12==12.6.80 in /usr/local/lib/python3.12/dist-packages (from torch) (12.6.80)
Requirement already satisfied: nvidia-cudnn-cu12==9.10.2.21 in /usr/local/lib/python3.12/dist-packages (from torch) (9.10.2.21)
Requirement already satisfied: nvidia-cublas-cu12==12.6.4.1 in /usr/local/lib/python3.12/dist-packages (from torch) (12.6.4.1)
Requirement already satisfied: nvidia-cufft-cu12==11.3.0.4 in /usr/local/lib/python3.12/dist-packages (from torch) (11.3.0.4)
Requirement already satisfied: nvidia-curand-cu12==10.3.7.77 in /usr/local/lib/python3.12/dist-packages (from torch) (10.3.7.77)
Requirement already satisfied: nvidia-cusolver-cu12==11.7.1.2 in /usr/local/lib/python3.12/dist-packages (from torch) (11.7.1.2)
Requirement already satisfied: nvidia-cusparse-cu12==12.5.4.2 in /usr/local/lib/python3.12/dist-packages (from torch) (12.5.4.2)
Requirement already satisfied: nvidia-cusparseelt-cu12==0.7.1 in /usr/local/lib/python3.12/dist-packages (from torch) (0.7.1)
Requirement already satisfied: nvidia-nccl-cu12==2.27.3 in /usr/local/lib/python3.12/dist-packages (from torch) (2.27.3)
Requirement already satisfied: nvidia-nvtx-cu12==12.6.77 in /usr/local/lib/python3.12/dist-packages (from torch) (12.6.77)
Requirement already satisfied: nvidia-nvjitlink-cu12==12.6.85 in /usr/local/lib/python3.12/dist-packages (from torch) (12.6.85)
Requirement already satisfied: nvidia-cufile-cu12==1.11.1.6 in /usr/local/lib/python3.12/dist-packages (from torch) (1.11.1.6)
Requirement already satisfied: triton==3.4.0 in /usr/local/lib/python3.12/dist-packages (from torch) (3.4.0)
Requirement already satisfied: mpmath<1.4,>=1.1.0 in /usr/local/lib/python3.12/dist-packages (from sympy>=1.13.3->torch) (1.3.0)
Requirement already satisfied: MarkupSafe>=2.0 in /usr/local/lib/python3.12/dist-packages (from jinja2->torch) (3.0.2)
```

In [3]: !nvidia-smi

```
Mon Sep 15 05:50:36 2025
+-----+
-----+
| NVIDIA-SMI 550.54.15           Driver Version: 550.54.15     CUDA Version: 12.
4      |
|-----+-----+-----+
-----+
| GPU  Name           Persistence-M | Bus-Id          Disp.A | Volatile Uncor
r. ECC |
| Fan  Temp    Perf            Pwr:Usage/Cap |          Memory-Usage | GPU-Util  Compu
te M. |
|          |                               |                   |           | GPU-Util  Comp
te M. |
|          |                               |                   |           | M
IG M. |
|=====+=====+=====+=====+=====+=====+=====+=====+
=====|
|   0  Tesla T4             Off  | 00000000:00:04.0 Off |
0 |
| N/A   44C     P8          10W / 70W | 0MiB / 15360MiB | 0%       De
fault |
|          |                   |                   |           |
N/A |
+-----+-----+-----+-----+-----+-----+-----+-----+
-----+
+-----+
-----+
| Processes:
|
| GPU  GI  CI          PID  Type  Process name          GPU M
emory |
|       ID  ID
|
|=====+=====+=====+=====+=====+=====+=====+=====+
=====|
| No running processes found
|
+-----+
-----+
```

In [4]: `import torch`

In [15]: `from torchvision import datasets, transforms`

```
train_data = datasets.MNIST(root='./data', train=True, download=True)
test_data = datasets.MNIST(root='./data', train=False, download=True)
# https://docs.pytorch.org/vision/main/generated/torchvision.datasets.MNIST.html

train_images, train_labels = train_data.data, train_data.targets
test_images, test_labels = test_data.data, test_data.targets
```

- The images are encoded as Numpy arrays, and the labels are an array of digits (0-9).
- The images and labels have a one-to-one correspondence.

In [16]: `print(type(train_images))`

```
<class 'torch.Tensor'>

In [17]: print('Shape of train_images array: ', train_images.shape)
print('# of training samples: ', len(train_images))

print('Shape of test_images array: ', test_images.shape)
print('# of test samples: ', len(test_images))

Shape of train_images array:  torch.Size([60000, 28, 28])
# of training samples:  60000
Shape of test_images array:  torch.Size([10000, 28, 28])
# of test samples:  10000
```

```
In [18]: train_images[10]
```



```
    0,    0,    0,    0,    0,    0,    0,    0,    0,    0,    0,    0,    0,    0,    0]],  
    dtype=torch.uint8)
```

```
In [19]: train_images[10].shape
```

```
Out[19]: torch.Size([28, 28])
```

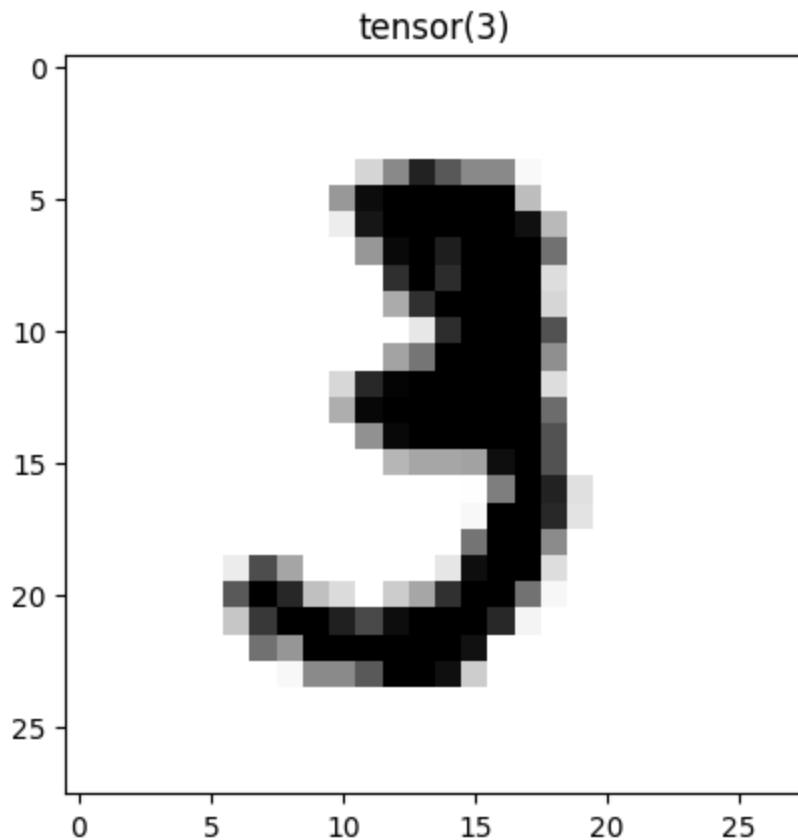
```
In [20]: train_labels
```

```
Out[20]: tensor([5, 0, 4, ..., 5, 6, 8])
```

```
In [21]: train_labels.shape
```

```
Out[21]: torch.Size([60000])
```

```
In [39]: import matplotlib.pyplot as plt  
import matplotlib.cm as cm  
plt.style.use('default') # set the same default style as jupyter notebook  
  
def visualize_mnist(image, label):  
    plt.imshow(image, cm.binary)  
    plt.title(label)  
    plt.show()  
  
visualize_mnist(train_images[10], train_labels[10])
```



- We'll feed the neural network the training data, `train_images` and `train_labels`.

- After training, we'll ask the network to produce predictions for `test_images`, and we'll verify whether these predictions match the labels from `test_labels`.
- Let's build the network.

```
In [1]: import torch
import torch.nn as nn
import torch.nn.functional as F
```

```
class SimpleNN(nn.Module):
    def __init__(self):
        super().__init__()
        self.fc1 = nn.Linear(28*28, 512)
        self.fc2 = nn.Linear(512, 10)

    def forward(self, x):
        x = self.fc1(x)
        x = F.relu(x)
        x = self.fc2(x)
        return x
```

```
In [3]: !pip install -q torchinfo
```

```
In [4]: from torchinfo import summary

model = SimpleNN()
summary(model)
```

```
Out[4]: =====
Layer (type:depth-idx)           Param #
=====
SimpleNN                         --
|Linear: 1-1                      401,920
|Linear: 1-2                      5,130
=====
Total params: 407,050
Trainable params: 407,050
Non-trainable params: 0
=====
```

```
In [ ]:
```

- As we learned, the core building block of neural networks is the *layers*, a data-processing module.
- Specifically, layers extract *representations* out of the data fed into them.
- We can obtain more useful representations as training goes.
- Here, our network consists of a sequence of two `Linear` layers, which are densely connected (also called `fully connected`) layers.

- The last layer is a 10-way output layer, which produces 10 raw scores (called logits). These are not probabilities yet; the cross-entropy loss function will apply softmax to convert them into probabilities during training.
- To make the network ready for training, we need three more things,
 - A loss function
 - An optimizer
 - Metrics to monitor during training and testing

```
In [17]: is_cuda = torch.cuda.is_available()
device = torch.device('cuda' if is_cuda else 'cpu')
print('Current device is', device)
```

Current device is cuda

```
In [18]: model = SimpleNN().to(device)
criterion = nn.CrossEntropyLoss()
optimizer = torch.optim.RMSprop(model.parameters(), lr=1e-3)

def accuracy(logits, labels):
    preds = torch.argmax(logits, dim=1)
    return (preds == labels).float().mean()
```

- Before training, we'll define the data feeding pipeline using `Dataset` and `Dataloader`.
 - Instead of loading all arrays and preprocessing them upfront, we stream mini-batches from a `Dataset` with on-the-fly transforms.
 - Scaling to [0,1] is handled by `transforms.ToTensor()` (casts to `float32` and divides by 255).
 - Flattening from (1,28,28) -> (784) can be done either in the model or as a transform.

```
In [19]: from torchvision import datasets, transforms
from torch.utils.data import DataLoader

# ToTensor(): uint8[0,255] -> float32[0,1], (H,W) -> (1,H,W)
# Lambda: (1,28,28) -> (784,)
train_tfms = transforms.Compose([
    transforms.ToTensor(),
    transforms.Lambda(t: t.view(-1)) # flatten
])
test_tfms = transforms.Compose([
    transforms.ToTensor(),
    transforms.Lambda(t: t.view(-1))
])

train_ds = datasets.MNIST(root='./data', train=True, download=True, transform=train_tfms)
test_ds = datasets.MNIST(root='./data', train=False, download=True, transform=test_tfms)
```

```
train_loader = DataLoader(train_ds, batch_size=64, shuffle=True, num_workers=0, pin_memory=True)
test_loader = DataLoader(test_ds, batch_size=256, shuffle=False, num_workers=0, pin_memory=True)
```

- Now, we are ready to train the network.

In [20]:

```
model.train()
for epoch in range(5):
    running_loss = 0.0
    running_acc = 0.0
    n_batches = 0

    for data, target in train_loader:
        data = data.to(device)
        target = target.to(device)

        optimizer.zero_grad()
        logits = model(data)
        loss = criterion(logits, target)
        loss.backward()
        optimizer.step()

        running_loss += loss.item()
        running_acc += accuracy(logits, target).item()
        n_batches += 1

    train_loss = running_loss / n_batches
    train_acc = running_acc / n_batches
    print("Epoch : {}\tLoss : {:.3f}\tAcc : {:.3f}".format(epoch, train_loss, train_acc))
```

```
Epoch : 0      Loss : 0.213554 Acc : 0.937883
Epoch : 1      Loss : 0.088478 Acc : 0.973264
Epoch : 2      Loss : 0.057449 Acc : 0.982276
Epoch : 3      Loss : 0.039944 Acc : 0.987623
Epoch : 4      Loss : 0.028961 Acc : 0.991005
```

- Two quantities are displayed during training, *loss* and *accuracy*.
- Note that these quantities do not guarantee the *test* performance.
- Let's check that the model performs well on the test set, too.

In [21]:

```
model.eval()
correct = 0
for data, target in test_loader:
    data, target = data.to(device), target.to(device)
    output = model(data)
    prediction = output.data.max(1)[1]
    correct += prediction.eq(target.data).sum()
print('Test set Accuracy : {:.2f}%'.format(100. * correct / len(test_loader.dataset)))
```

Test set Accuracy : 97.97%

- The test accuracy is a bit lower than the training accuracy.
- This gap between training accuracy and test accuracy is an example of *overfitting*.

- We saw how we can build and train a neural network to classify handwritten digits.

Data representations for neural networks

- In the previous example, the data was stored in multidimensional Numpy arrays, also called *tensors*.
- In general, all current machine learning systems use tensors as their basic data structure.
- A tensor is a container for data.
- The number of *axes* (also called *dimensions*) of a tensor is called its *rank*.

Scalars (0D tensors)

- A tensor that contains only one number is called a *scalar*.
- For example, a `float32` or `float64` number in Numpy is a scalar tensor.

```
In [22]: import numpy as np  
x = np.array(12)  
x
```

```
Out[22]: array(12)
```

```
In [23]: x.ndim
```

```
Out[23]: 0
```

Vectors (1D tensors)

- An array of numbers is called a *vector*, or 1D tensor.

```
In [24]: x = np.array([12, 3, 6, 14, 5])  
x
```

```
Out[24]: array([12, 3, 6, 14, 5])
```

```
In [25]: x.ndim
```

```
Out[25]: 1
```

- This vector has five entries and so is called a *5-dimensional vector*.
- Don't confuse a 5D vector with a 5D tensor!

Matrices (2D tensors)

- An array of vectors is a *matrix*, or 2D tensor.
- A matrix has two axes (rows and columns).

```
In [26]: x = np.array([[5, 78, 2, 34, 0],
                   [6, 79, 3, 35, 1],
                   [7, 80, 4, 36, 2]])
x
```

```
Out[26]: array([[ 5, 78,  2, 34,  0],
                 [ 6, 79,  3, 35,  1],
                 [ 7, 80,  4, 36,  2]])
```

```
In [27]: x.ndim
```

```
Out[27]: 2
```

```
In [28]: x.shape
```

```
Out[28]: (3, 5)
```

3D tensors and higher-dimensional tensors

- An array of matrices? --> a 3D tensor
- It can be visually interpreted as a cube of numbers.

```
In [29]: x = np.array([ [ [5, 78, 2, 34, 0],
                     [6, 79, 3, 35, 1],
                     [7, 80, 4, 36, 2] ],
                     [ [5, 78, 2, 34, 0],
                     [6, 79, 3, 35, 1],
                     [7, 80, 4, 36, 2] ],
                     [ [5, 78, 2, 34, 0],
                     [6, 79, 3, 35, 1],
                     [7, 80, 4, 36, 2] ] ])
print(x.ndim)
print(x.shape)
```

```
3
(3, 3, 5)
```

- By packing 3D tensors in an array, you can create a 4D tensor, and so on.
- In deep learning, you'll generally manipulate tensors that are 0D to 4D, although you may go up to 5D if you process video data.

Key attributes

- *Number of axes (rank)*: the tensor's `ndim` in Numpy

- *Shape*: the tensor's `shape` in Numpy
- *Data type*: usually called `dtype` in Python libraries
 - This is the type of the data contained in the tensor: `float32`, `float64`, `uint8`, and so on.

```
In [35]: train_data = datasets.MNIST(root='./data', train=True, download=True)
test_data = datasets.MNIST(root='./data', train=False, download=True)

train_images, train_labels = train_data.data, train_data.targets
test_images, test_labels = test_data.data, test_data.targets

train_images = train_images.numpy()

print(train_images.ndim)
```

3

```
In [36]: print(train_images.shape)

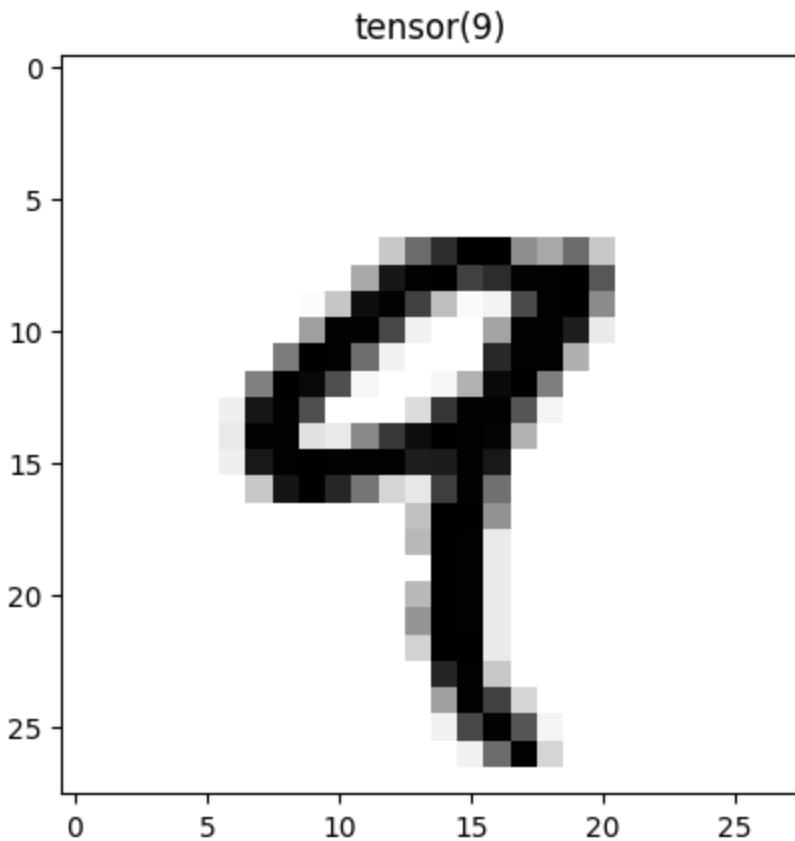
(60000, 28, 28)
```

```
In [37]: print(train_images.dtype)

uint8
```

- So, our `train_images` is a 3D tensor of 8-bit integers.
- More precisely, it is an array of 60,000 matrices of 28*28 integers.
- Each matrix is a grayscale images, with elements between 0 and 255.

```
In [40]: visualize_mnist(train_images[4], train_labels[4])
```



Manipulating tensors in Numpy

- In the previous example, we selected a specific digit alongside the first axis using the syntax `train_images[i]`.
- Selecting specific elements in a tensor is called *tensor slicing*.

```
In [41]: my_slice = train_images[10:100]
print(my_slice.shape)
```

```
(90, 28, 28)
```

```
In [42]: my_slice = train_images[10:100, :, :] # : is equivalent to selecting the entire axis
print(my_slice.shape)
```

```
my_slice = train_images[10:100, 0:28, 0:28]
print(my_slice.shape)
```

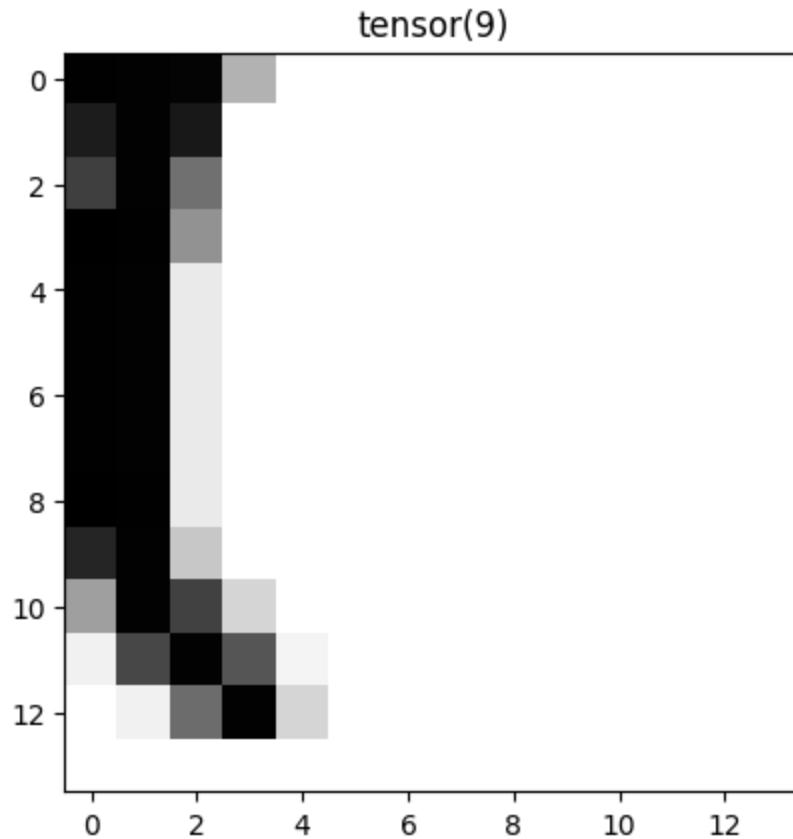
```
(90, 28, 28)
(90, 28, 28)
```

- If you want to select 14*14 pixels in the bottom-right corner of all images:

```
In [43]: my_slice = train_images[:, 14:, 14:]
print(my_slice.shape)
```

```
(60000, 14, 14)
```

```
In [44]: visualize_mnist(my_slice[4], train_labels[4])
```



- It is also possible to use negative indices.
- Negative indices indicate a position relative to the end of the current axis.
- For example, if you want to crop the images to patches of 14*14 pixels centered in the middle:

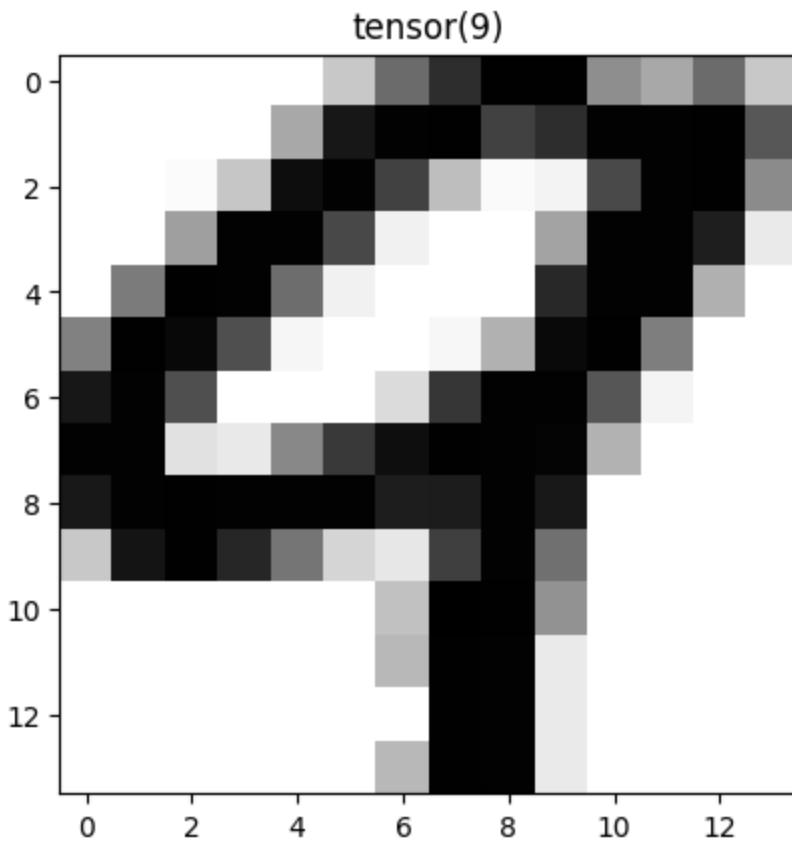
```
In [45]: np.array([0,1,2,3,4,5])[-2]
```

```
Out[45]: np.int64(4)
```

```
In [46]: my_slice = train_images[:, 7:-7, 7:-7]
print(my_slice.shape)
```

```
(60000, 14, 14)
```

```
In [47]: visualize_mnist(my_slice[4], train_labels[4])
```



The notion of data batches

- In general, the first axis in all data tensors you'll come across in deep learning will be the *samples axis* (sometimes called the *samples dimension*, *batch axis* or *batch dimension*).
- In addition, deep learning models don't process the entire dataset at once. Instead, a `DataLoader` streams the dataset in mini-batches. You simply iterate over the `DataLoader`, which yields `(inputs, labels)` tensors of the configured batch size (e.g., 128), handles shuffling, and takes care of the final, possibly smaller batch.

```
In [ ]: for xb, yb in train_loader:
    # xb: (B, 784) float32 in [0,1]
    # yb: (B,) int64 class labels
    # run forward/backward/optimizer.step()
    pass
```

Real-world examples of data tensors

- Vector data - 2D tensors of shape `(samples, features)`
- Timeseries data or sequence data - 3D tensors of shape `(samples, timesteps, features)` No description has been provided for this image
 - A dataset of stock prices

- A dataset of tweets
- Images - 4D tensors of shape `(samples, height, width, channels)` or `(samples, channels, height, width)`
 No description has been provided for this image
 - A batch of 128 color images could be stored in a tensor of shape `(128, 256, 256, 3)`.
 - There are two conventions for shapes of images tensors: the *channels-last* convention and the *channels-first* convention.
 - For example, in Tensorflow,
https://www.tensorflow.org/api_docs/python/tf/nn/convolution
- Video - 5D tensors of shape `(samples, frames, height, width, channels)` or `(samples, frames, channels, height, width)`
 - Each frame can be stored in a 3D tensor `(height, width, channels)`, a sequence of frames can be stored in a 4D tensor `(frames, height, width, channels)`, and thus a batch of different videos can be stored in a 5D tensor of shape `(samples, frames, height, width, channels)`.
 - For instance, a 60-second, 144*256 YouTube video clip samples at 4 frames per second would have 240 frames. A batch of four such video clips would be stored in a tensor of shape `(4, 240, 144, 256, 3)`. --> A total of 106,168,320 values!
 - If the `dtype` of the tensor was `float32`, then each value would be stored in 32 bits, so the tensor would represent 405MB.

The gears of neural networks: tensor operations

- All transformations learned by deep neural networks can be reduced to a handful of *tensor operations* applied to tensors of numeric data.
- In our initial example, we were building our network by stacking `Dense` layers on top of each other.
 - `keras.layers.Dense(512, activation='relu')`
- This layer can be interpreted as a function. Specifically, the function is:
 - `output = relu(dot(W, input) + b)` where `W` is a 2D tensor and `b` is a vector.
 - We have three tensor operations here:
 - A dot product (`dot`) between the input tensor and a tensor named `W`
 - An addition (`+`) between the resulting 2D tensor and a vector `b`
 - A `relu` operation of `relu(x)=max(x, 0)`

Element-wise operations

- The `relu` operation and addition are `element-wise` operations.
- If we write a naive Python implementation of `relu` using `for` loop:

```
In [ ]: def naive_relu(x):
    assert len(x.shape) == 2 # x is a 2D Numpy tensor.

    x = x.copy() # to avoid overwriting the input tensor
    for i in range(x.shape[0]):
        for j in range(x.shape[1]):
            x[i, j] = max(x[i, j], 0)
    return x
```

- We can do the same thing for addition:

```
In [ ]: def naive_add(x, y):
    assert len(x.shape) == 2
    assert x.shape == y.shape

    x = x.copy()
    for i in range(x.shape[0]):
        for j in range(x.shape[1]):
            x[i, j] += y[i, j]
    return x
```

- In practice, these operations are provided as built-in Numpy functions, which are well-optimized via BLAS (Basic Linear Algebra Subprograms) implemented in Fortran or C.
 - So, in Numpy, you can do the following element-wise operations very efficiently.
- ```
import numpy as np
z = x+y
z = np.maximum(z, 0.)
```

## Broadcasting

- What happens with addition when the shapes of the two tensors being added differ?
  - In our `naive_add` implementation, it only supports the addition of 2D tensors with identical shapes.
- When possible, and if there's no ambiguity, the smaller tensor will be *broadcasted* to match the shape of the larger tensor.
- Broadcasting consists of two steps:
  - Axes (called *broadcast axes*) are added to the smaller tensor to match the `ndim` of the larger tensor.
  - The smaller tensor is repeated alongside these new axes to match the full shape of the larger tensor.
- For example, consider `X` with shape `(32, 10)` and `y` with shape `(10,)`.
  - First, we add an empty first axis to `y`. --> The shape of `y` becomes `(1, 10)`.

- Then, we repeat `y` 32 times alongside this new axis. --> We get `Y` with shape of `(32, 10)` where `Y[i, :] == y` for `i` in `range(0, 32)`.
- Now, we can proceed to add `X` and `Y`.
- Note that the repetition operation is entirely virtual: it happens at the algorithmic level rather than at the memory level.
- Here is what a naive implementation would look like:

```
In []: def naive_add_matrix_and_vector(x, y):
 assert len(x.shape) == 2
 assert len(y.shape) == 1
 assert x.shape[1] == y.shape[0]

 x = x.copy()
 for i in range(x.shape[0]):
 for j in range(x.shape[1]):
 x[i, j] += y[j]
 return x
```

- With broadcasting, we can generally apply two-tensor element-wise operations if one tensor has shape `(a, b, ..., n, n+1, ..., m)` and the other has shape `(n, n+1, ..., m)`.
- The broadcasting will then automatically happen for axes `a` through `n-1`.
- Example: the element-wise `maximum` operation to two tensors of different shapes via broadcasting

```
In []: import numpy as np

x = np.random.random((64, 3, 32, 10))
y = np.random.random((10,))
z = np.maximum(x, y)
```

```
In []: print(y)
[0.80652492 0.85132886 0.13638803 0.55177141 0.06940284 0.90881542
 0.5270101 0.78581118 0.83440423 0.25812784]
```

```
In []:
```

## Tensor dot

- The dot operation, also called a *tensor product* is the most common, most useful tensor operation.
  - An element-wise product is done with the `*` operator in most libraries including Numpy, Keras, and Tensorflow.
  - In both Numpy and Keras, the dot operation uses the standard `dot` operator.
- ```
import numpy as np
z = np.dot(x, y)
```

- The dot product of two vectors x and y

```
In [ ]: def naive_vector_dot(x, y):
    assert len(x.shape) == 1
    assert len(y.shape) == 1
    assert x.shape[0] == y.shape[0]

    z = 0.
    for i in range(x.shape[0]):
        z += x[i] * y[i]
    return z
```

- The dot product between a matrix x and a vector y

```
In [ ]: import numpy as np

def naive_matrix_vector_dot(x, y):
    assert len(x.shape) == 2
    assert len(y.shape) == 1
    assert x.shape[1] == y.shape[0] # Note that if x is a shape of (m,n),
                                # then y should be a shape of (n,).

    z = np.zeros(x.shape[0])
    for i in range(x.shape[0]):
        for j in range(x.shape[1]):
            z[i] += x[i, j] * y[j]
    return z
```

```
In [ ]: def naive_matrix_vector_dot(x, y):
    z = np.zeros(x.shape[0])
    for i in range(x.shape[0]):
        z[i] = naive_vector_dot(x[i, :], y)
    return z
```

- Note that if one of two tensors has an `ndim` greater than 1, `dot` operations is no longer symmetric: `dot(x, y) != dot(y, x)`.
 - The dot product of two matrices x and y , `dot(x, y)`
 - It can be computed if and only if `x.shape[1] == y.shape[0]`.
 - The result is a matrix with shape `(x.shape[0], y.shape[1])`, where the coefficients are the vector products between the rows of x and the columns of y .
- 

```
In [ ]: def naive_matrix_dot(x, y):
    assert len(x.shape) == 2
    assert len(y.shape) == 2
    assert x.shape[1] == y.shape[0]

    z = np.zeros((x.shape[0], y.shape[1]))
    for i in range(x.shape[0]):
        for j in range(y.shape[1]):
```

```

    row_x = x[i, :]
    column_y = y[:, j]
    z[i, j] = naive_vector_dot(row_x, column_y)
return z

```

- More generally, we can take the dot product between higher-dimensional tensors.

$(a, b, c, d) \cdot (d,) \rightarrow (a, b, c)$
 $(a, b, c, d) \cdot (d, e) \rightarrow (a, b, c, e)$

Tensor reshaping

- In our first neural network example, we used *reshaping* operation when we preprocessed the digits data before feeding it into the network.

```
train_images = train_images.reshape((60000, 28*28))
```

- Reshaping means rearranging its rows and columns to match a target shape.

```
In [ ]: x = np.array([
    [0, 1],
    [2, 3],
    [4, 5]])
print(x.shape)
```

(3, 2)

```
In [ ]: x = x.reshape((6, 1))
print(x)
```

[
[0]
[1]
[2]
[3]
[4]
[5]]

```
In [ ]: x = x.reshape((2, 3))
print(x)
```

[
[[0 1 2]
[3 4 5]]

```
In [ ]: x = x.reshape((3, 4))
print(x)
```

ValueError Traceback (most recent call last)
<ipython-input-54-4a57eba2433c> in <cell line: 1>()
----> 1 x = x.reshape((3, 4))
 2 print(x)

ValueError: cannot reshape array of size 6 into shape (3,4)

- A special case of reshaping is *transposition*.
- *Transposing* a matrix means exchanging its rows and its columns, so that `x[i, :]` --> `x[:, i]`.

```
In [ ]: x = np.zeros((300, 20))
x = np.transpose(x)
print(x.shape)

(20, 300)
```

Geometric interpretation of tensor operations

- The contents of the tensors manipulated by tensor operations --> coordinates of points in some geometric space.
- Therefore, all tensor operations have a geometric interpretation.
- Geometric interpretation of the sum of two vectors

- In general, elementary geometric operations such as translation, rotation, scaling, skewing, and so on can be expressed as tensor operations.
 - Translation

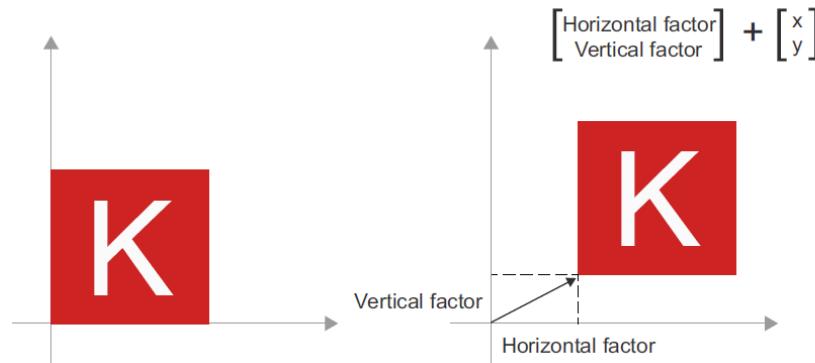


Figure 2.9 2D translation as a vector addition

- Rotation

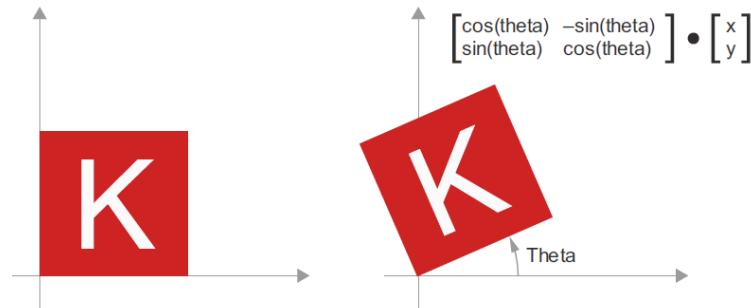


Figure 2.10 2D rotation (counterclockwise) as a dot product

- Scaling

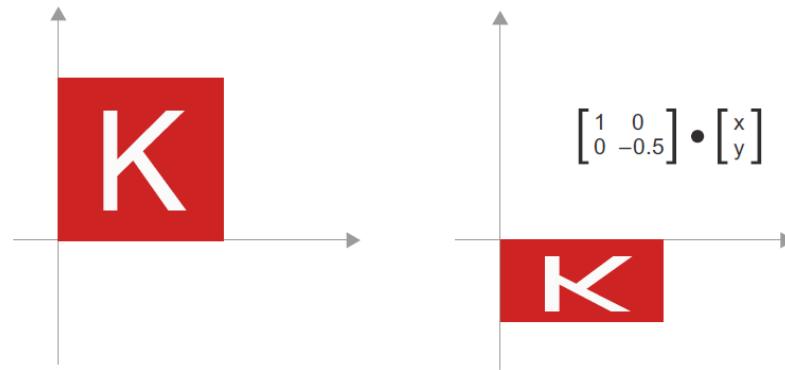


Figure 2.11
2D scaling as a
dot product

- Linear transform
 - A dot product with an arbitrary matrix, e.g., scaling and rotation
- Affine transform

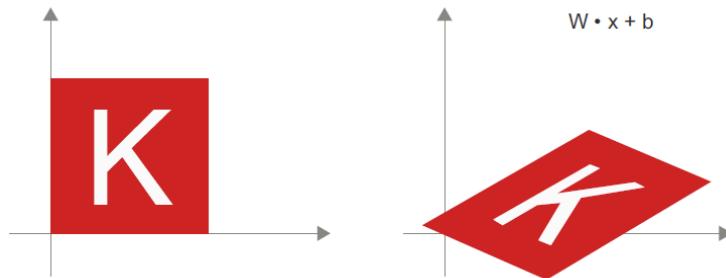


Figure 2.12 Affine
transform in the plane

- Dense layer with *relu* activation
 - Applying affine transformations repeatedly, then?
 - $\text{affine2}(\text{affine1}(x)) = W_2(W_1x+b_1)+b_2 = (W_2W_1)x+(W_2b_1b_2)$
 - A multilayer NN made entirely of *Dense* layers without activations would be equivalent to a single *Dense* layer.

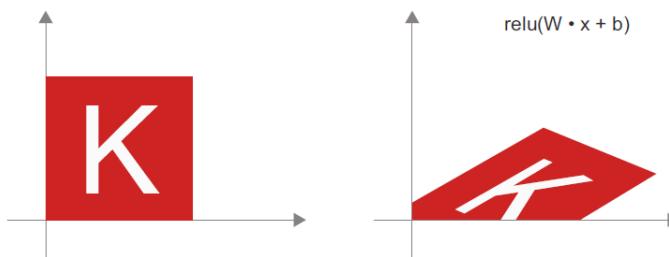


Figure 2.13 Affine
transform followed by
relu activation

```
In [ ]: import numpy as np

x = np.array([2, 0])
rotation = np.array([
    [0, -1],
    [1, 0]
])
z = np.dot(rotation, x) # Careful about the order!
print(z)
```

[0 2]

```
In [ ]: np.dot(x, rotation)
```

```
Out[ ]: array([ 0, -2])
```

```
In [ ]: rotation * x
```

```
Out[ ]: array([[0, 0],
               [2, 0]])
```

```
In [ ]: x * rotation
```

```
Out[ ]: array([[0, 0],
               [2, 0]])
```

A geometric interpretation of deep learning

- We just learned that neural networks consist entirely of chains of tensor operations and that all of these tensor operations are just geometric transformations of the input data.
- We can interpret a neural network = very complex geometric transformation.
- Image two sheets of colored paper, one red and one blue, and put one on top of the other.
 - Now crumple them together into a small ball.
 - That crumpled paper ball is your input data, and each sheet of paper is a class of data in a classification problem.
 - What a NN is meant to do is figure out a transformation of the paper ball that would uncrumple it, so as to make the two classes clearly separable again.
 - Uncrumpling paper balls is what machine learning is about: finding neat representations for complex, highly folded data *manifolds* in high-dimensional spaces (a manifold is a continuous surface, like our crumpled sheet of paper).
 - Deep learning takes the approach of incrementally decomposing a complicated geometric transformation into a long chain of elementary ones, which is pretty much the strategy a human would follow to uncrumple a paper ball.
 - Each layer in a deep network applies a transformation that disentangles the data a little, and a deep stack of layers makes tractable an extremely complicated disentanglement process.

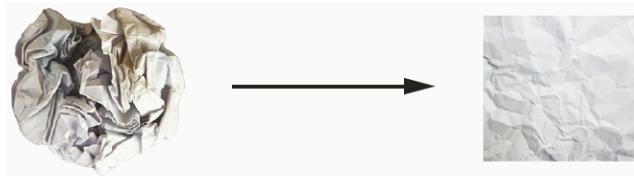


Figure 2.14 Uncrumpling a complicated manifold of data

Reading assignments

- Section 2.1 and 2.3 in "Dive into deep learning"