

# Optimizers

# Review. Backpropagation

- Output activation and loss function



| Task                      | # of output nodes | Output activation   | Loss function   |
|---------------------------|-------------------|---|---|
| Binary classification     | 1                 | Sigmoid<br>$\hat{y} = 1/(1 + \exp(-z))$                               | Binary cross-entropy<br>$L(y, \hat{y}) = -y \log \hat{y} - (1 - y) \log(1 - \hat{y})$ |
| Multiclass classification | $k$               | Softmax<br>$\hat{y}_i = \frac{\exp\{z_i\}}{\sum_{i=1}^k \exp\{z_i\}}$ | Cross-entropy<br>$L(y, \hat{y}) = - \sum_{i=1}^k y_i \log(\hat{y}_i)$                 |
| Regression                | $k$               | None<br>$\hat{y}_i = z_i$   | Mean squared error<br>$L(y, \hat{y}) = \frac{1}{2} \sum_{i=1}^k (\hat{y}_i - y_i)^2$  |

# Review. Backpropagation

- Refer to [Chap. 2 in Nielsen's Neural Networks and Deep Learning](#)

- Notations:  $w_{jk}^l, b_j^l, a_j^l, z_j^l, \delta_j^l = \frac{\delta C}{\delta z_j^l}$

- The Hadamard product  $\odot$
- Four fundamental equations behind backpropagation
  - (BP1)  $\delta^L = \nabla_a C \odot \sigma'(z^L)$
  - (BP2)  $\delta^l = (w^{l+1})^T \delta^{l+1} \odot \sigma'(z^l)$
  - (BP3)  $\frac{\partial C}{\partial b_j^l} = \delta_j^l$
  - (BP4)  $\frac{\partial C}{\partial w_{jk}^l} = a_k^{l-1} \delta_j^l$



# Review. Backpropagation

- (BP1)  $\delta^L = \nabla_a C \odot \sigma'(z^L)$  for multiclass classification

$$\frac{\delta C}{\delta z_j^L} = \sum_{i=1}^k \frac{\delta C}{\delta \hat{y}_i} \frac{\delta \hat{y}_i}{\delta z_j} = \hat{y}_j - y_j$$
$$\frac{\delta C}{\delta \hat{y}_i} = -\frac{y_i}{\hat{y}_i} \quad \frac{\delta \hat{y}_i}{\delta z_j} = \hat{y}_i (I_{ij} - \hat{y}_j)$$


# Review. Backpropagation

- Refer to [Chap. 2 in Nielsen's Neural Networks and Deep Learning](#)

1. **Input  $x$ :** Set the corresponding activation  $a^1$  for the input layer.

2. **Feedforward:** For each  $l = 2, 3, \dots, L$  compute

$$z^l = w^l a^{l-1} + b^l \text{ and } a^l = \sigma(z^l).$$

3. **Output error  $\delta^L$ :** Compute the vector  $\delta^L = \nabla_a C \odot \sigma'(z^L)$ .

4. **Backpropagate the error:** For each  $l = L - 1, L - 2, \dots, 2$  compute  $\delta^l = ((w^{l+1})^T \delta^{l+1}) \odot \sigma'(z^l)$ .

5. **Output:** The gradient of the cost function is given by

$$\frac{\partial C}{\partial w_{jk}^l} = a_k^{l-1} \delta_j^l \text{ and } \frac{\partial C}{\partial b_j^l} = \delta_j^l.$$

# Basic Gradient Descent Algorithms

- Batch gradient descent

---

**Algorithm 1** Batch Gradient Descent at Iteration  $k$ 

---

**Require:** Learning rate  $\epsilon_k$

**Require:** Initial Parameter  $\theta$

- 1: **while** stopping criteria not met **do**
  - 2:     Compute gradient estimate over  $N$  examples:
  - 3:      $\hat{\mathbf{g}} \leftarrow +\frac{1}{N} \nabla_{\theta} \sum_i L(f(\mathbf{x}^{(i)}; \theta), \mathbf{y}^{(i)})$
  - 4:     Apply Update:  $\theta \leftarrow \theta - \epsilon \hat{\mathbf{g}}$
  - 5: **end while**
-

# Basic Gradient Descent Algorithms

- Stochastic gradient descent

---

**Algorithm 8.1** Stochastic gradient descent (SGD) update

---

**Require:** Learning rate schedule  $\epsilon_1, \epsilon_2, \dots$

**Require:** Initial parameter  $\theta$

$k \leftarrow 1$

**while** stopping criterion not met **do**

    Sample a minibatch of  $m$  examples from the training set  $\{\mathbf{x}^{(1)}, \dots, \mathbf{x}^{(m)}\}$  with corresponding targets  $\mathbf{y}^{(i)}$ .

    Compute gradient estimate:  $\hat{\mathbf{g}} \leftarrow \frac{1}{m} \nabla_{\theta} \sum_i L(f(\mathbf{x}^{(i)}; \theta), \mathbf{y}^{(i)})$

    Apply update:  $\theta \leftarrow \theta - \epsilon_k \hat{\mathbf{g}}$


$k \leftarrow k + 1$

**end while**

---


# Basic Gradient Descent Algorithms

- Batch gradient descent VS Stochastic gradient descent



# Basic Gradient Descent Algorithms

- How to determine the learning rate?
  - If it is too large, the objective value may explode.
  - If it is too small, a lot of iterations are needed.



# Basic Gradient Descent Algorithms

- Momentum: accumulates an exponentially decaying moving average of past gradients and continues to move in that direction


previous direction  
(accumulation of past gradients)

$$v \leftarrow \alpha v - \epsilon \nabla_{\theta} \left( \frac{1}{m} \sum_{i=1}^m L(f(x^{(i)}; \theta), y^{(i)}) \right)$$

current direction

current gradient

Figure Credit: Prof. Kang (SKKU)



# Basic Gradient Descent Algorithms

- Stochastic gradient descent with momentum

---

**Algorithm 8.2** Stochastic gradient descent (SGD) with momentum

---

**Require:** Learning rate  $\epsilon$ , momentum parameter  $\alpha$

**Require:** Initial parameter  $\boldsymbol{\theta}$ , initial velocity  $\boldsymbol{v}$

**while** stopping criterion not met **do**

    Sample a minibatch of  $m$  examples from the training set  $\{\mathbf{x}^{(1)}, \dots, \mathbf{x}^{(m)}\}$  with corresponding targets  $\mathbf{y}^{(i)}$ .

    Compute gradient estimate:  $\mathbf{g} \leftarrow \frac{1}{m} \nabla_{\boldsymbol{\theta}} \sum_i L(f(\mathbf{x}^{(i)}; \boldsymbol{\theta}), \mathbf{y}^{(i)})$ .

    Compute velocity update:  $\boldsymbol{v} \leftarrow \alpha \boldsymbol{v} - \epsilon \mathbf{g}$ .

    Apply update:  $\boldsymbol{\theta} \leftarrow \boldsymbol{\theta} + \boldsymbol{v}$ .


**end while**

---



# Basic Gradient Descent Algorithms

- Stochastic gradient descent with momentum

Local Minima





Saddle points




# Basic Gradient Descent Algorithms

- Nesterov Momentum: the gradient is evaluated after the current velocity is applied.





# Basic Gradient Descent Algorithms

## Momentum



## Nesterov Momentum

Figure Credit: Prof. Kang (SKKU)



# Basic Gradient Descent Algorithms

- Stochastic gradient descent with Nesterov's momentum

---

**Algorithm 8.3** Stochastic gradient descent (SGD) with Nesterov momentum

---

**Require:** Learning rate  $\epsilon$ , momentum parameter  $\alpha$

**Require:** Initial parameter  $\boldsymbol{\theta}$ , initial velocity  $\boldsymbol{v}$

**while** stopping criterion not met **do**

    Sample a minibatch of  $m$  examples from the training set  $\{\mathbf{x}^{(1)}, \dots, \mathbf{x}^{(m)}\}$  with corresponding labels  $\mathbf{y}^{(i)}$ .

    Apply interim update:  $\tilde{\boldsymbol{\theta}} \leftarrow \boldsymbol{\theta} + \alpha \boldsymbol{v}$ .

    Compute gradient (at interim point):  $\mathbf{g} \leftarrow \frac{1}{m} \nabla_{\tilde{\boldsymbol{\theta}}} \sum_i L(f(\mathbf{x}^{(i)}; \tilde{\boldsymbol{\theta}}), \mathbf{y}^{(i)})$ .

    Compute velocity update:  $\boldsymbol{v} \leftarrow \alpha \boldsymbol{v} - \epsilon \mathbf{g}$ .


    Apply update:  $\boldsymbol{\theta} \leftarrow \boldsymbol{\theta} + \boldsymbol{v}$ .

**end while**

---

# Basic Gradient Descent Algorithms

## Why Momentum Really Works



<https://distill.pub/2017/momentum/>

# Algorithms with Adaptive Learning Rates

- The learning rate significantly affects the performance.
- How to determine the learning rate?
- Adaptive learning rate
  - To use a separate learning rate for each parameter and automatically adapt these learning rates throughout the course of learning

# Algorithms with Adaptive Learning Rates

- AdaGrad (2011)
  - Large gradients → small update, small gradients → large update
  - Code snippet

```
# Assume the gradient dx and parameter vector x
cache += dx**2
x += - learning_rate * dx / (np.sqrt(cache) + eps)
```

- Cache keeps growing during training → learning rate shrinks too quickly

# Algorithms with Adaptive Learning Rates

- AdaGrad (2011)

---

**Algorithm 8.4** The AdaGrad algorithm

---

**Require:** Global learning rate  $\epsilon$

**Require:** Initial parameter  $\theta$

**Require:** Small constant  $\delta$ , perhaps  $10^{-7}$ , for numerical stability

Initialize gradient accumulation variable  $r = \mathbf{0}$

**while** stopping criterion not met **do**

    Sample a minibatch of  $m$  examples from the training set  $\{\mathbf{x}^{(1)}, \dots, \mathbf{x}^{(m)}\}$  with corresponding targets  $\mathbf{y}^{(i)}$ .

    Compute gradient:  $\mathbf{g} \leftarrow \frac{1}{m} \nabla_{\theta} \sum_i L(f(\mathbf{x}^{(i)}; \theta), \mathbf{y}^{(i)})$

    Accumulate squared gradient:  $\mathbf{r} \leftarrow \mathbf{r} + \mathbf{g} \odot \mathbf{g}$

    Compute update:  $\Delta\theta \leftarrow -\frac{\epsilon}{\delta + \sqrt{r}} \odot \mathbf{g}$ . (Division and square root applied element-wise)

    Apply update:  $\theta \leftarrow \theta + \Delta\theta$

**end while**

# Algorithms with Adaptive Learning Rates

- RMSProp (2012)
  - Extension of AdaGrad which solves AdaGrad's aggressive, monotonically decreasing learning rate problem
  - The gradient accumulation → exponentially weighted moving average

$$r \leftarrow r + g \odot g \quad \blacktriangleright \quad r \leftarrow \rho r + (1 - \rho)g \odot g$$

```
cache = decay_rate * cache + (1 - decay_rate) * dx**2
x += - learning_rate * dx / (np.sqrt(cache) + eps)
```

# Algorithms with Adaptive Learning Rates

- RMSProp (2012)

---

**Algorithm 8.5** The RMSProp algorithm

---

**Require:** Global learning rate  $\epsilon$ , decay rate  $\rho$ .

**Require:** Initial parameter  $\theta$

**Require:** Small constant  $\delta$ , usually  $10^{-6}$ , used to stabilize division by small numbers.

Initialize accumulation variables  $r = 0$

**while** stopping criterion not met **do**

    Sample a minibatch of  $m$  examples from the training set  $\{\mathbf{x}^{(1)}, \dots, \mathbf{x}^{(m)}\}$  with corresponding targets  $\mathbf{y}^{(i)}$ .

    Compute gradient:  $\mathbf{g} \leftarrow \frac{1}{m} \nabla_{\theta} \sum_i L(f(\mathbf{x}^{(i)}; \theta), \mathbf{y}^{(i)})$

    Accumulate squared gradient:  $\mathbf{r} \leftarrow \rho \mathbf{r} + (1 - \rho) \mathbf{g} \odot \mathbf{g}$

    Compute parameter update:  $\Delta \theta = -\frac{\epsilon}{\sqrt{\delta + \mathbf{r}}} \odot \mathbf{g}$ . ( $\frac{1}{\sqrt{\delta + \mathbf{r}}}$  applied element-wise)

    Apply update:  $\theta \leftarrow \theta + \Delta \theta$

**end while**

# Algorithms with Adaptive Learning Rates

- Adam (2015)
  - Adaptive moment estimation
  - It uses the first moment and the second moment.
  - A variant on the combination of RMSProp and momentum

```
m = beta1*m + (1-beta1)*dx
v = beta2*v + (1-beta2)*(dx**2)
x += - learning_rate * m / (np.sqrt(v) + eps)
```

# Algorithms with Adaptive Learning Rates

- Adam (2015)

---

**Algorithm 8.7** The Adam algorithm

---

Require: Step size  $\epsilon$  (Suggested default: 0.001)

Require: Exponential decay rates for moment estimates,  $\rho_1$  and  $\rho_2$  in  $[0, 1]$ .  
(Suggested defaults: 0.9 and 0.999 respectively)

Require: Small constant  $\delta$  used for numerical stabilization. (Suggested default:  
 $10^{-8}$ )

Require: Initial parameters  $\theta$

Initialize 1st and 2nd moment variables  $s = 0, r = 0$

Initialize time step  $t = 0$

**while** stopping criterion not met **do**

    Sample a minibatch of  $m$  examples from the training set  $\{x^{(1)}, \dots, x^{(m)}\}$  with  
    corresponding targets  $y^{(i)}$ .

    Compute gradient:  $\mathbf{g} \leftarrow \frac{1}{m} \nabla_{\theta} \sum_i L(f(x^{(i)}; \theta), y^{(i)})$

$t \leftarrow t + 1$

    Update biased first moment estimate:  $s \leftarrow \rho_1 s + (1 - \rho_1) \mathbf{g}$

    Update biased second moment estimate:  $r \leftarrow \rho_2 r + (1 - \rho_2) \mathbf{g} \odot \mathbf{g}$

    Correct bias in first moment:  $\hat{s} \leftarrow \frac{s}{1 - \rho_1^t}$


    Correct bias in second moment:  $\hat{r} \leftarrow \frac{r}{1 - \rho_2^t}$

    Compute update:  $\Delta\theta = -\epsilon \frac{\hat{s}}{\sqrt{\hat{r}} + \delta}$  (operations applied element-wise)

    Apply update:  $\theta \leftarrow \theta + \Delta\theta$


**end while**

# And Many More




# Learning Rate Scheduling


- Learning rate decaying
  - Constant learning rate often prevents convergence.



**Step decay**




**Exponential decay**



**Test accuracy**

# Animations for Optimization Algorithms



# Reading assignments

- “Understanding deep learning”
  - Chapter 6
- “Dive into deep learning”
  - Section 12