

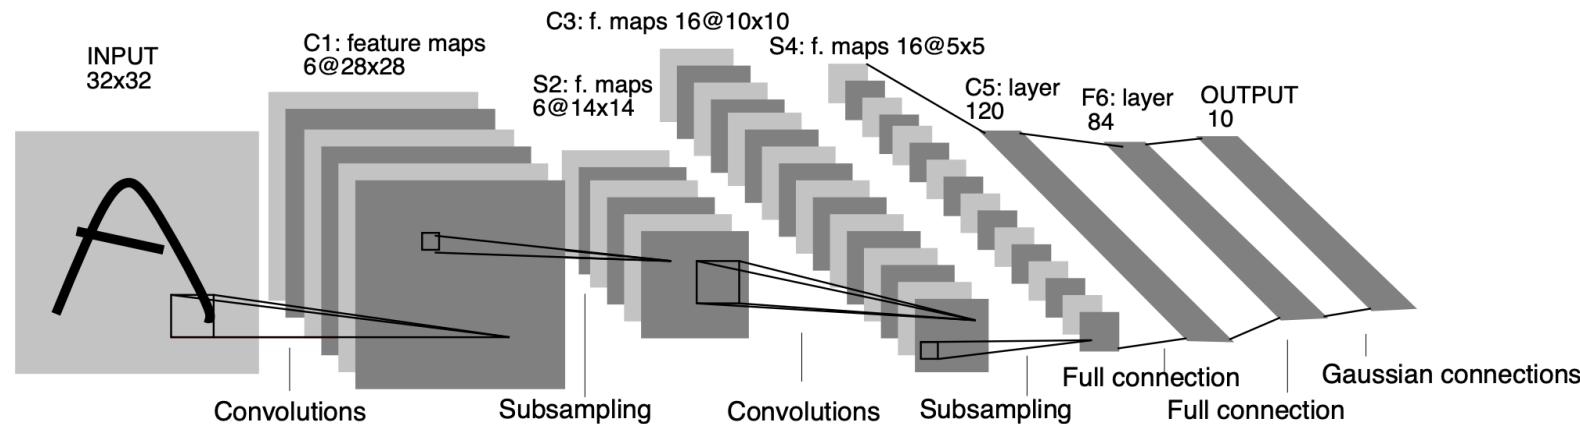
Convolution and Pooling

Computer Vision Tasks

- Classification
- Object detection
- Semantic or instance segmentation
- Others
 - Tracking in videos, camera pose estimation, body pose estimation, 3D reconstruction, denoising, super-resolution, auto-captioning, synthesis, etc.

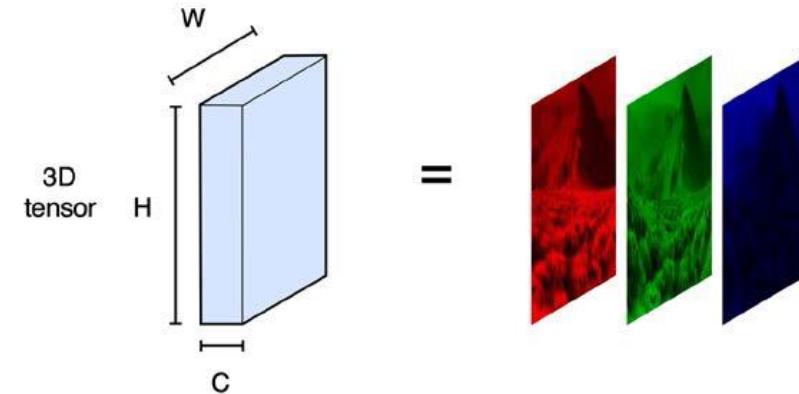
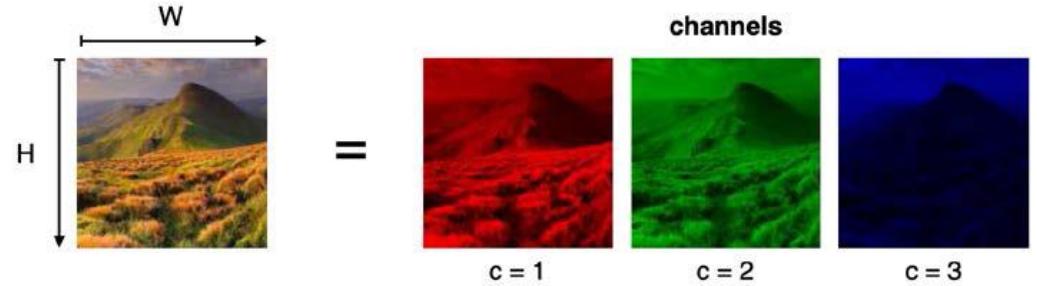
Convolutional Neural Networks

- A specialized kind of neural network for processing data that has a known grid-like topology such as time-series data, image data, video data, etc.
- Ex. LeNet-5 (LeCun et al., 1998)



Recap. Tensors

- 3D tensor = 3-dimensional array
 - RGB color image
 - (height, width, channel)
 - Grayscale image = 2D tensor
- 4D tensor = 4-dimensional array
 - Color video
 - (time, height, width, channel)



Motivation

- A linear layer taking a 256*256 RGB image as input, and producing an image of same size would require $(256 \times 256 \times 3)^2 \cong 3.87 \times 10^{10}$.
- Some input signals have some “invariance in translation”.
 - A function f of x is invariant to a transformation T if $f(T(x)) = f(x)$.
 - A transformation meaningful at a certain location can be used everywhere.
- A convolution layer embodies this idea.
 - It applies the same linear transformation locally, everywhere.

Convolution

- Convolution is a spatial filtering.



$$\begin{matrix} & \xrightarrow{*^{1/8}} \\ \xleftarrow{*} & \end{matrix}$$

0	1	0
1	4	1
0	1	0

0	-1	0
-1	4	-1
0	-1	0

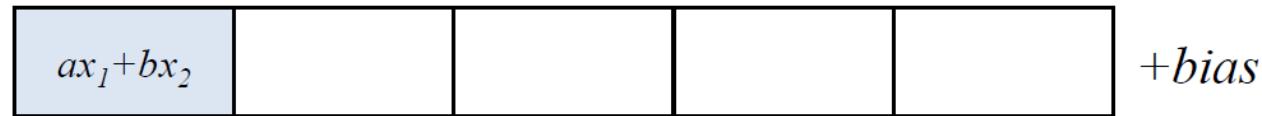
1	0	-1
2	0	-2
1	0	-1



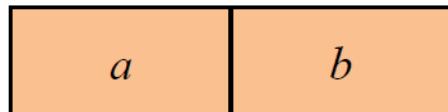
Convolution with 1D Array Input

- **kernel = [a,b] ← parameters**
(a.k.a. filter)
- input size m=6, kernel size k=2, stride(kernel step size) s=1, output size n= $(m-k)/s+1=5$

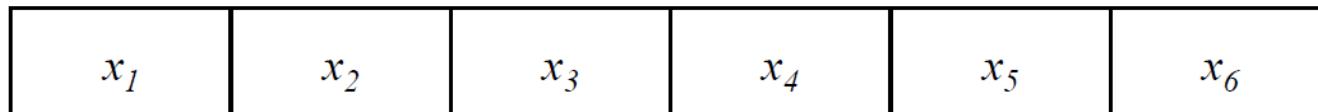
output



kernel



input



Convolution with 1D Array Input

- **kernel = [a,b]** ← parameters
- input size m=6, kernel size k=2, stride(kernel step size) s=1, output size n= $(m-k)/s+1=5$

output

ax_1+bx_2	ax_2+bx_3	ax_3+bx_4	ax_4+bx_5	ax_5+bx_6	+bias
-------------	-------------	-------------	-------------	-------------	-------

kernel



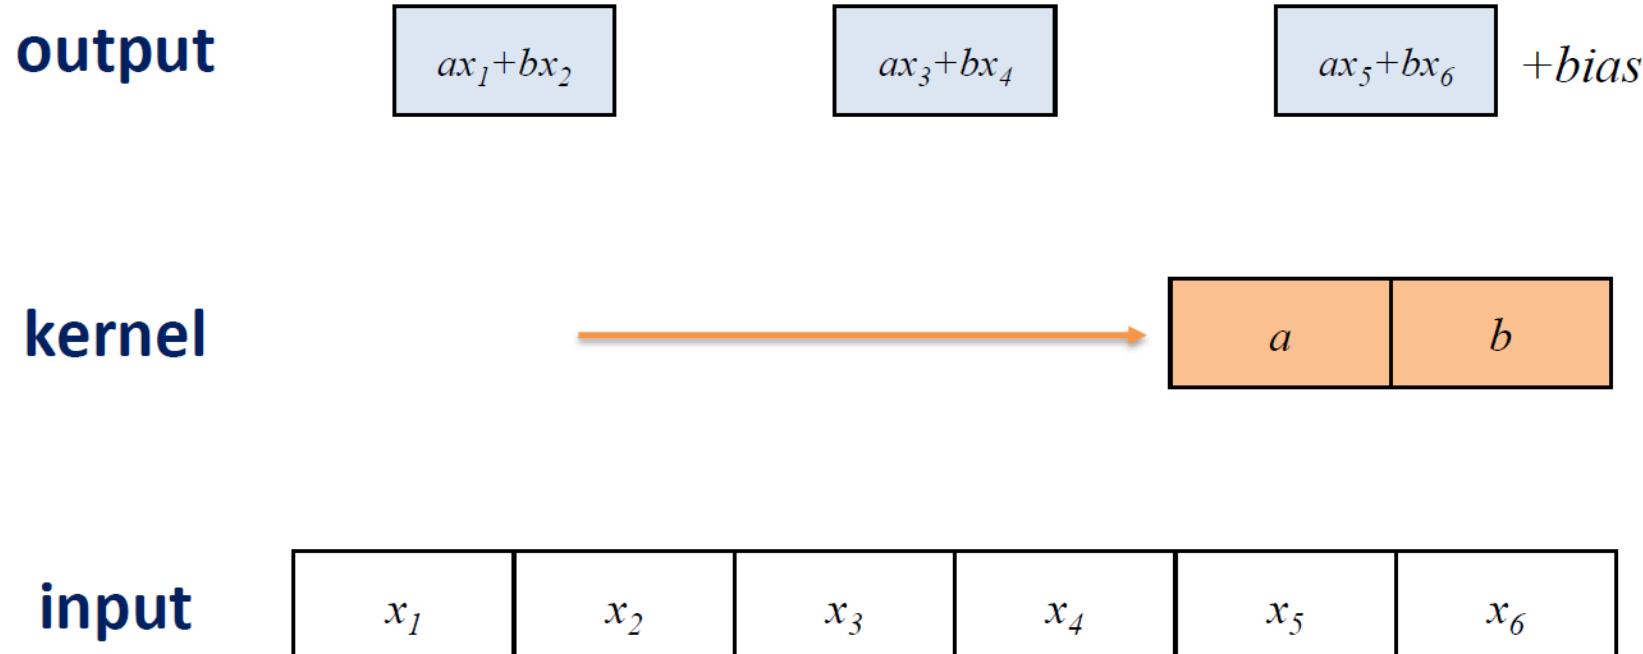
a	b
---	---

input

x_1	x_2	x_3	x_4	x_5	x_6
-------	-------	-------	-------	-------	-------

Convolution with 1D Array Input

- **kernel = [a,b]** ← parameters
- input size m=6, kernel size k=2, **stride(kernel step size) s=2**, output size n= $(m-k)/s+1=3$



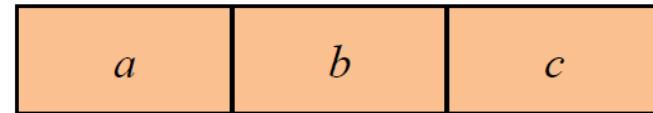
Convolution with 1D Array Input

- **kernel** = [a,b,c] ← parameters
- input size m=6, **kernel size k=3**, stride(kernel step size) s=1, output size n= $(m-k)/s+1=4$

output



kernel



input



Convolution with 1D Array Input

- **kernel = [a,b,c]** ← parameters
- input size m=6, **kernel size k=3**, stride(kernel step size) s=1, output size n= $(m-k)/s+1=4$

output

$$[ax_1+bx_2+cx_3 | ax_2+bx_3+cx_4 | ax_3+bx_4+cx_5 | ax_4+bx_5+cx_6] + bias$$

kernel



a	b	c
---	---	---

input

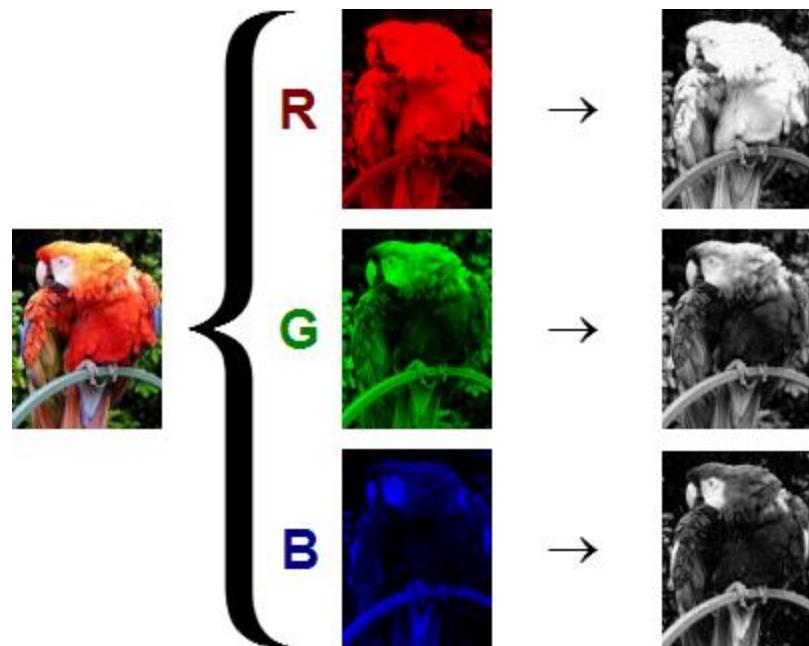
x_1	x_2	x_3	x_4	x_5	x_6
-------	-------	-------	-------	-------	-------

How Computers See Images

- Image is just an array of numbers.

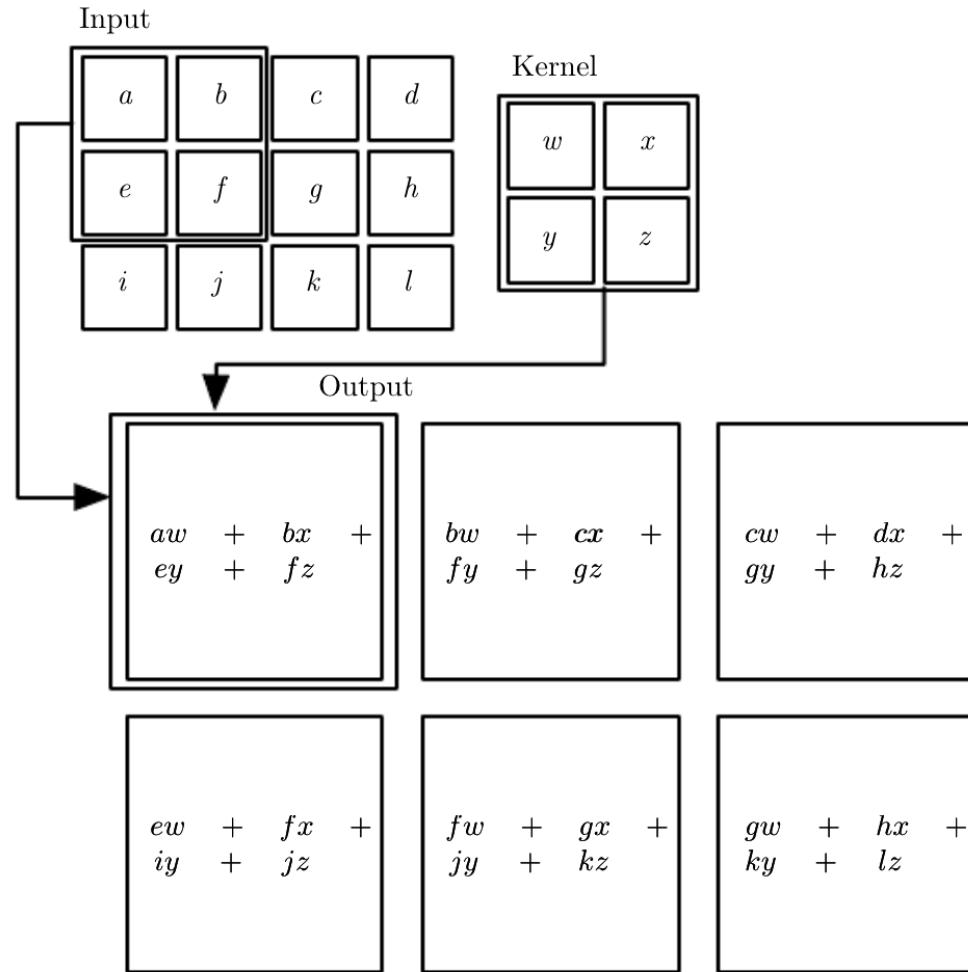


```
[[ 9  1  29  70 114  76  0  8  4  5  5  0 111 162  9  8  62  62]
 [ 3  0  33  61 102 106  34  0  0  0  0  49 182 150  1 12  65  62]
 [ 1  0  40  54 123  90  72  77  52  51  49 121 205  98  0 15  67  59]
 [ 3  1  41  57  74  54  96 181 220 170  90 149 208  56  0 16  69  59]
 [ 6  1  32  36  47  81  85  90 176 206 140 171 186  22  3 15  72  63]
 [ 4  1  31  39  66  71  71  97 147 214 203 190 198  22  6 17  73  65]
 [ 2  3  15  30  52  57  68 123 161 197 207 200 179  8  8 18  73  66]
 [ 2  2  17  37  34  40  78 103 148 187 205 225 165  1  8 19  76  68]
 [ 2  3  20  44  37  34  35  26  78 156 214 145 200  38  2 21  78  69]
 [ 2  2  20  34  21  43  70  21  43 139 205  93 211  70  0 23  78  72]
 [ 3  4  16  24  14  21 192 175 120 130 226 212 236  75  0 25  78  72]
 [ 6  5  13  21  28  28  97 216 184  90 196 255 255  84  4 24  79  74]
 [ 6  5  15  25  30  39  63 185 140  66 113 252 251  74  4 28  79  75]
 [ 5  5  16  32  38  57  69  85  93 120 128 251 255 154  19  26  80  76]
 [ 6  5  20  42  55  62  66  76  86 104 148 242 254 241  83  26  80  77]
 [ 2  3  20  38  55  64  69  80  78 109 195 247 252 255 172  40  78  77]
 [ 10  8  23  34  44  64  88 104 119 173 234 247 253 254 227  66  74  74]
 [ 32  6  24  37  45  63  85 114 154 196 226 245 251 252 250 112  66  71]]
```



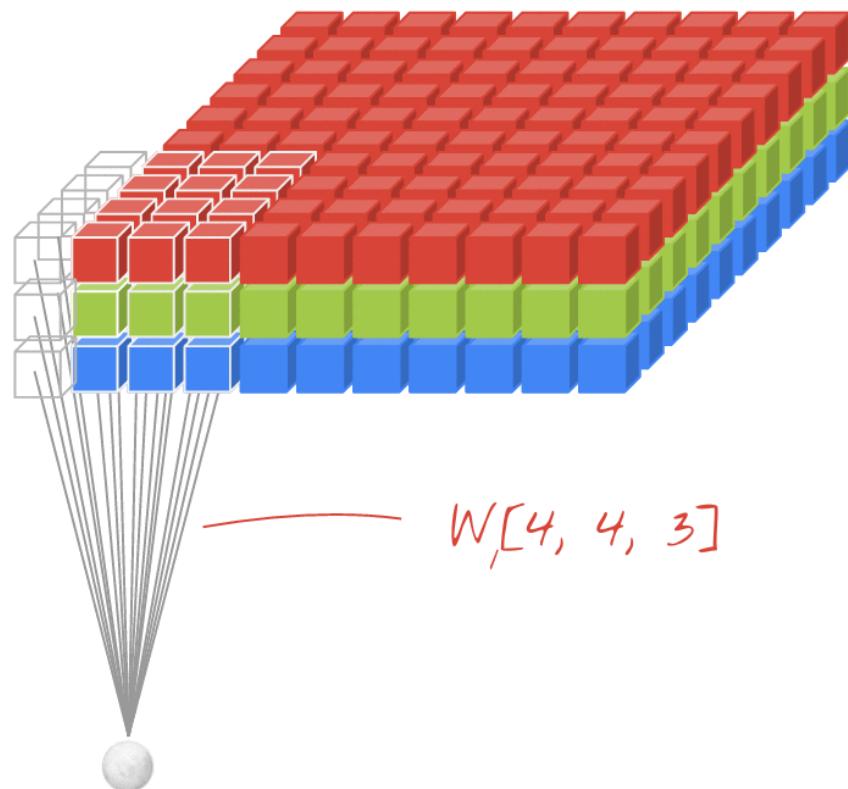
<https://savan77.github.io/blog/how-computers-see-image.html>

Convolution with 2D Tensor



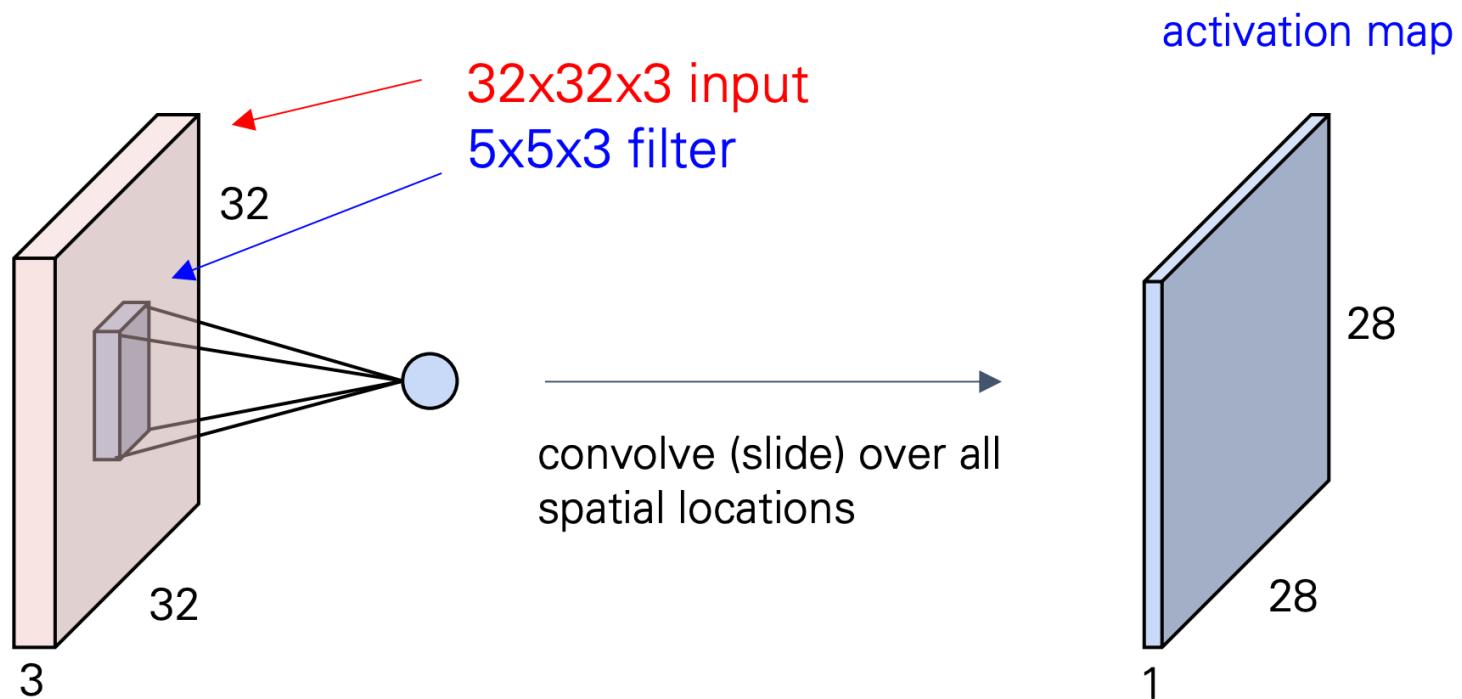
What are parameters?
What are hyperparameters?

Convolution with 3D Tensor



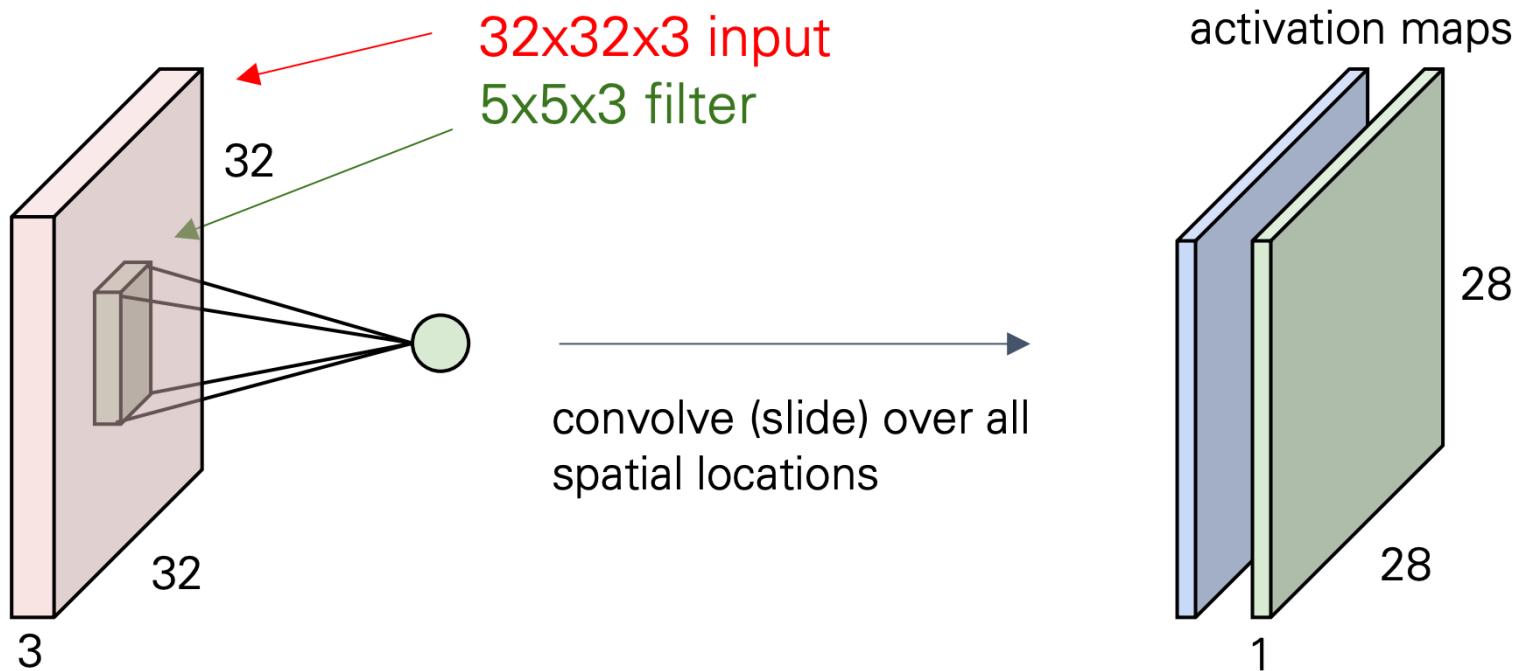
Convolution with 3D Tensor

- Generally, we use multiple kernels for single convolutional layer.



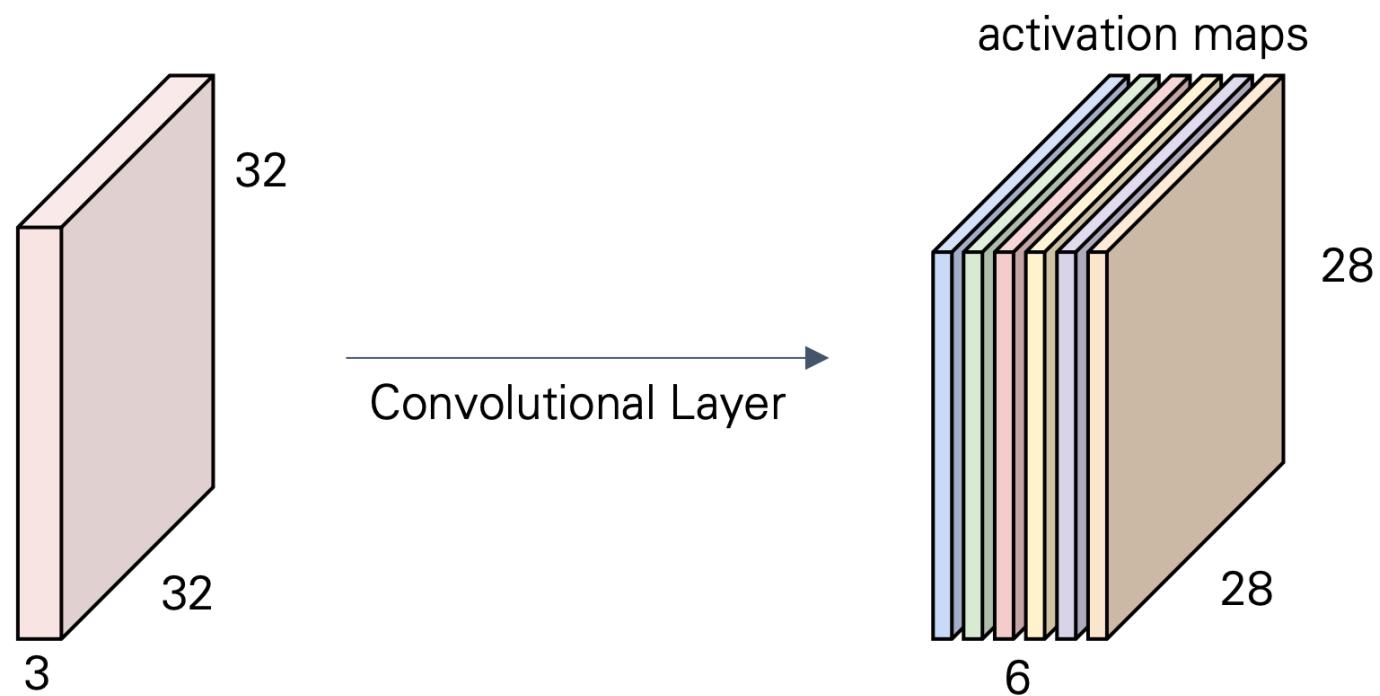
Convolution with 3D Tensor

- Generally, we use multiple kernels for single convolutional layer.



Convolution with 3D Tensor

- Generally, we use multiple kernels for single convolutional layer.



Parameters in Convolution

- The **padding** specifies the size of a zeroed frame added around the input.
- The **stride** specifies a step size when moving the kernel across the signal.
- The **dilation** modulates the expansion of the filter without adding weights.

Parameters in Convolution: Padding

- There is border effects in convolution.
- Valid convolution

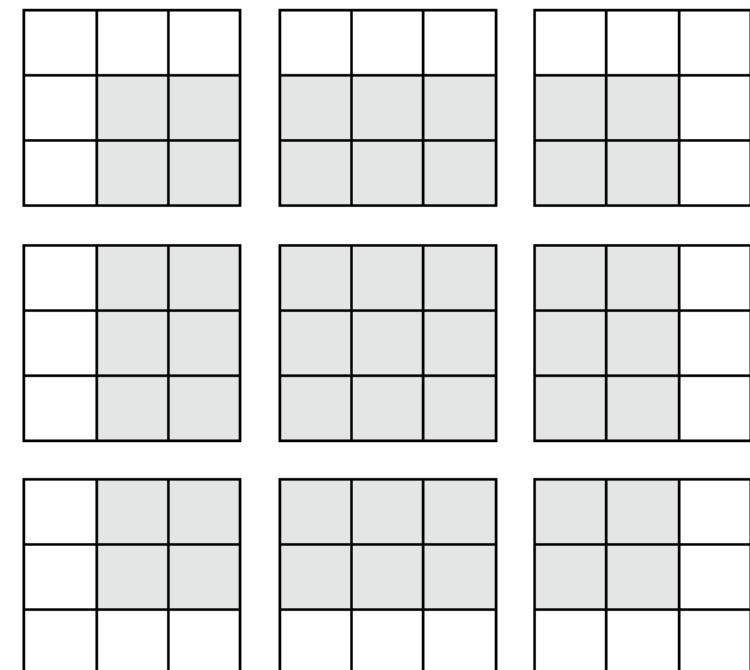
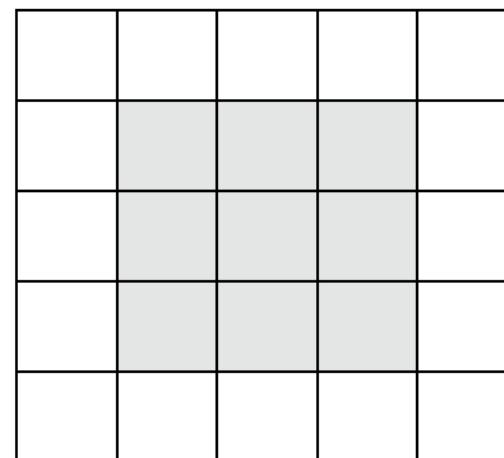


Figure 8.5 Valid locations of 3×3 patches in a 5×5 input feature map

Parameters in Convolution: Padding

- Padding to an input

Typically set to zero → Zero padding

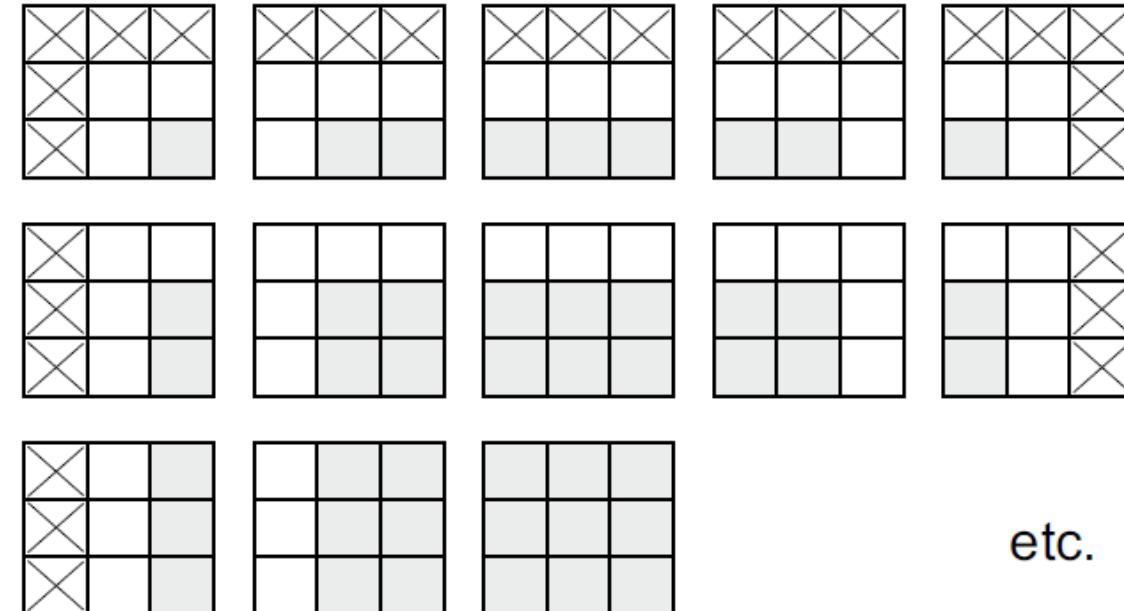
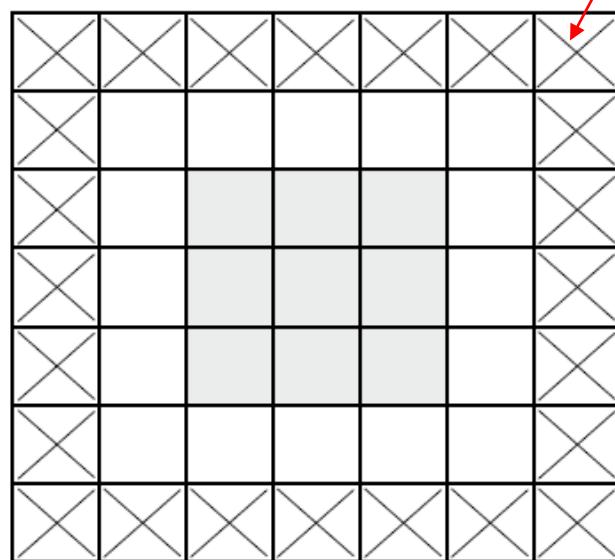


Figure 8.6 Padding a 5×5 input in order to be able to extract 25 3×3 patches

Parameters in Convolution: Stride

- Strided convolution: convolutions with a stride higher than 1

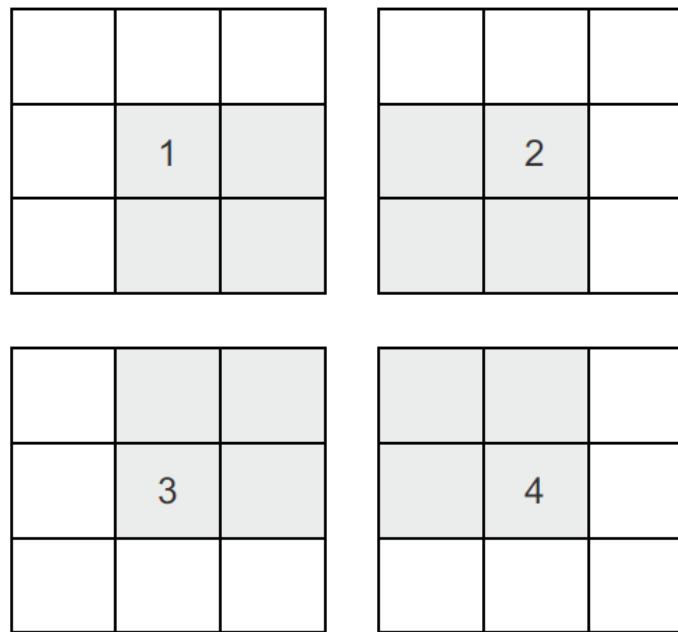
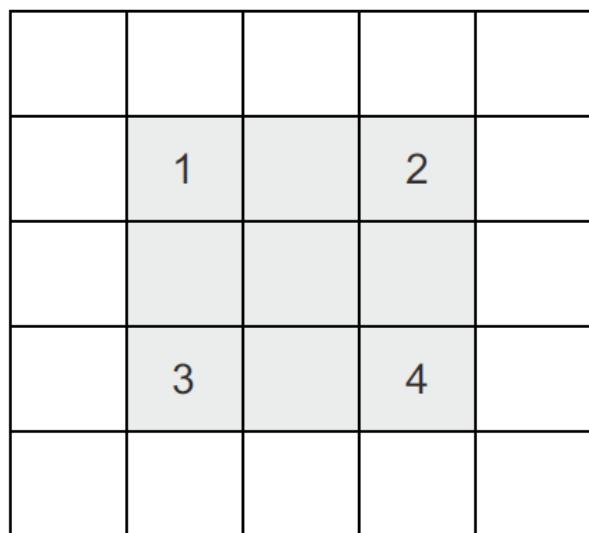
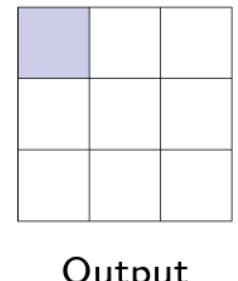
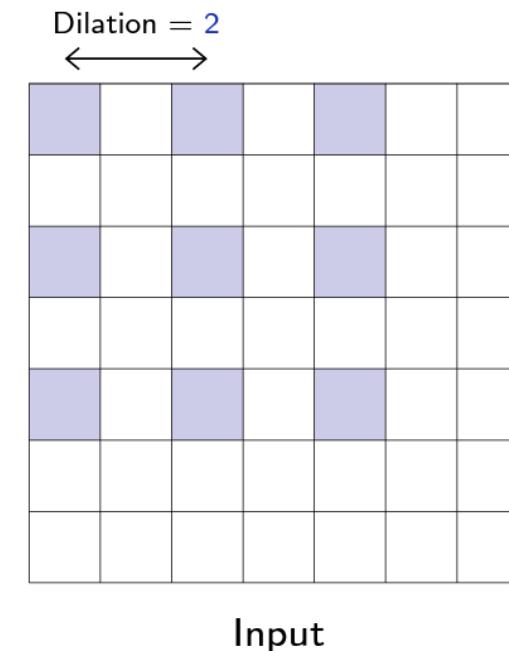
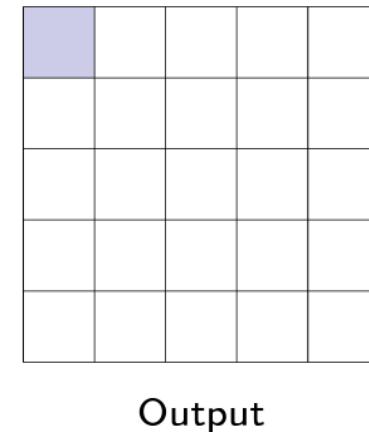
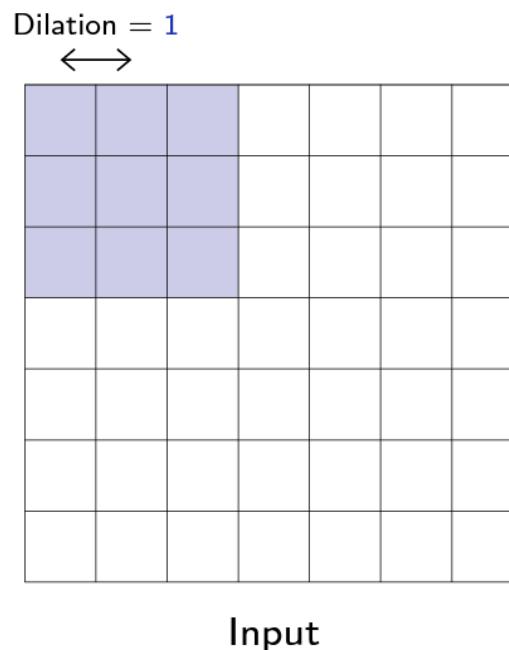


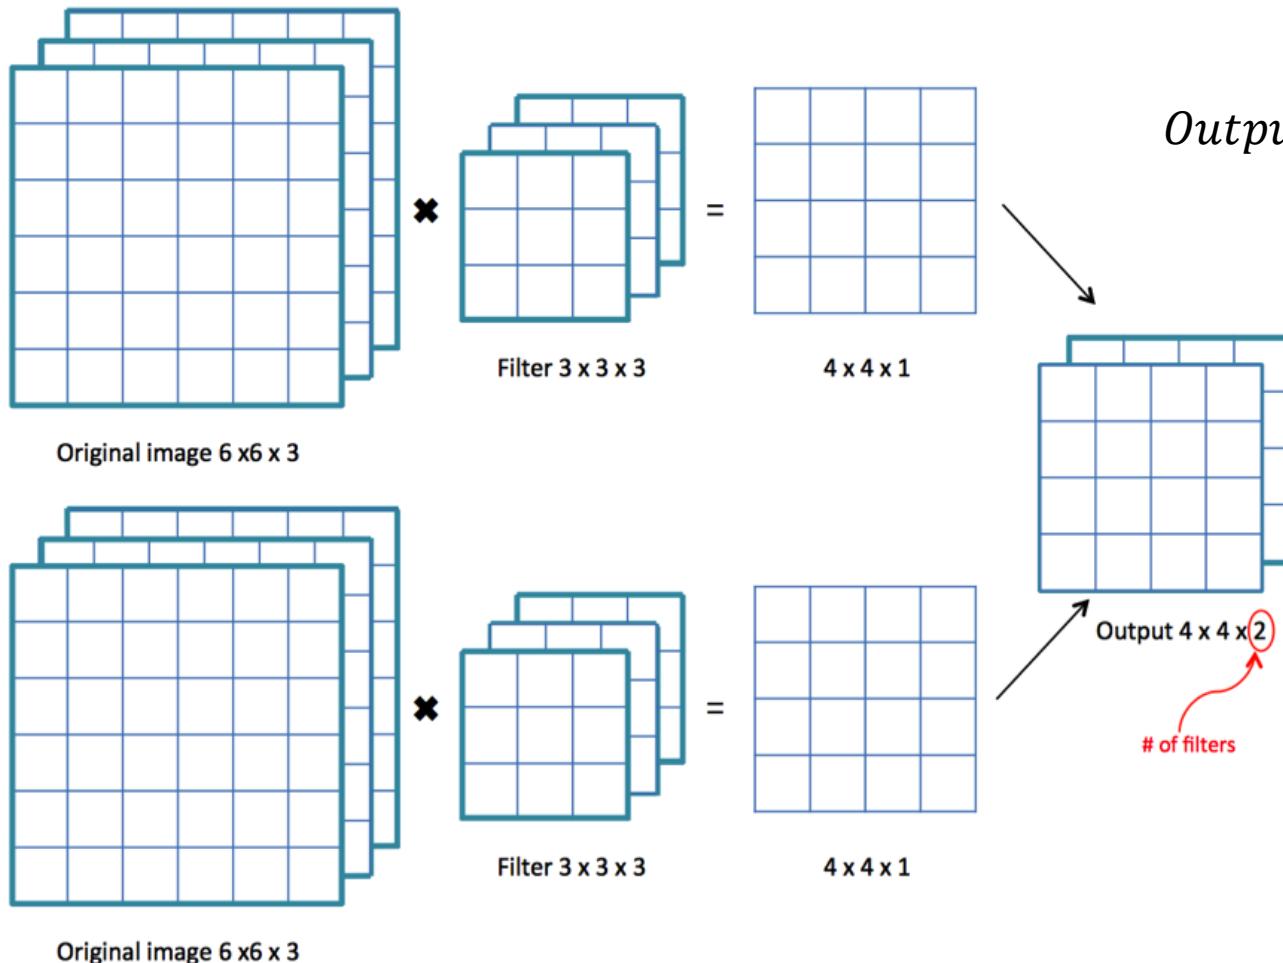
Figure 8.7 3×3 convolution patches with 2×2 strides

Parameters in Convolution: Dilation

- The expansion of a filter by adding rows and columns of zeros between coefficient.



How to Compute Output Shape



$$\text{Output Size} = \left(\frac{H + 2P - F}{S} + 1 \right) \times \left(\frac{W + 2P - F}{S} + 1 \right)$$

Properties of Convolution

- A convolution preserves the signal support structure.
- Sparse interactions
 - Inputs and outputs are not fully connected but have local connectivity
- Parameter sharing
 - The same kernel is used repeatedly.
- Equivariance to transition
 - $\text{convolution}(\text{shift}(\text{input})) = \text{shift}(\text{convolution}(\text{input}))$

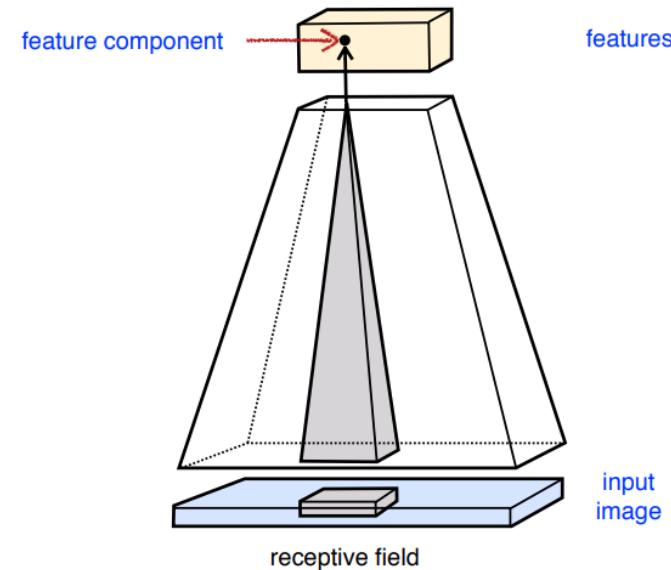
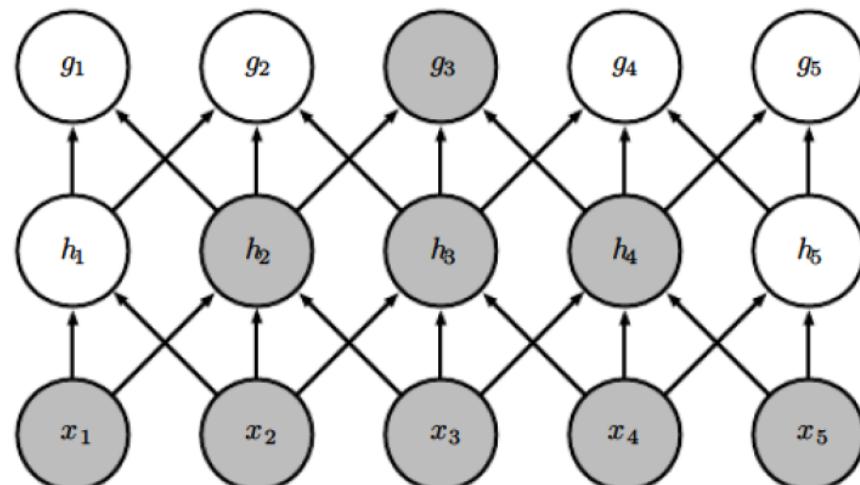
- **Receptive field: Spatial locality**

- Each element of the feature map process only for its receptive field (a local region of the input)
- higher kernel size $k \rightarrow$ larger receptive field
- Higher-level layers \rightarrow larger receptive field

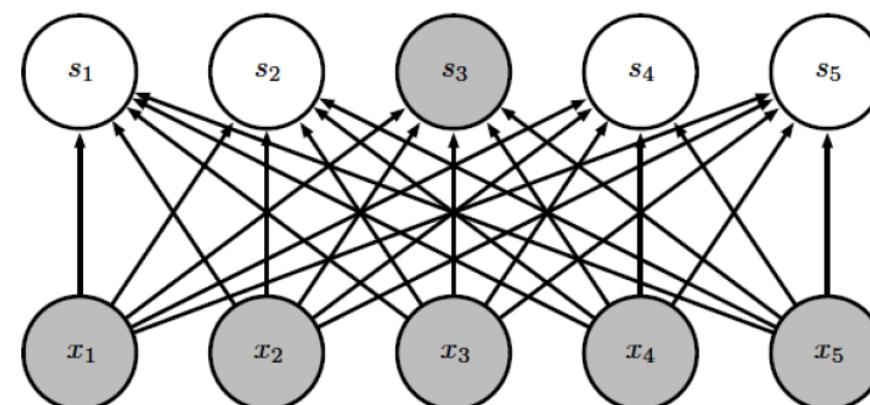
Example: a simple CNN with $k=3$

Receptive field of h_3 in the input layer = $\{x_2, x_3, x_4\}$

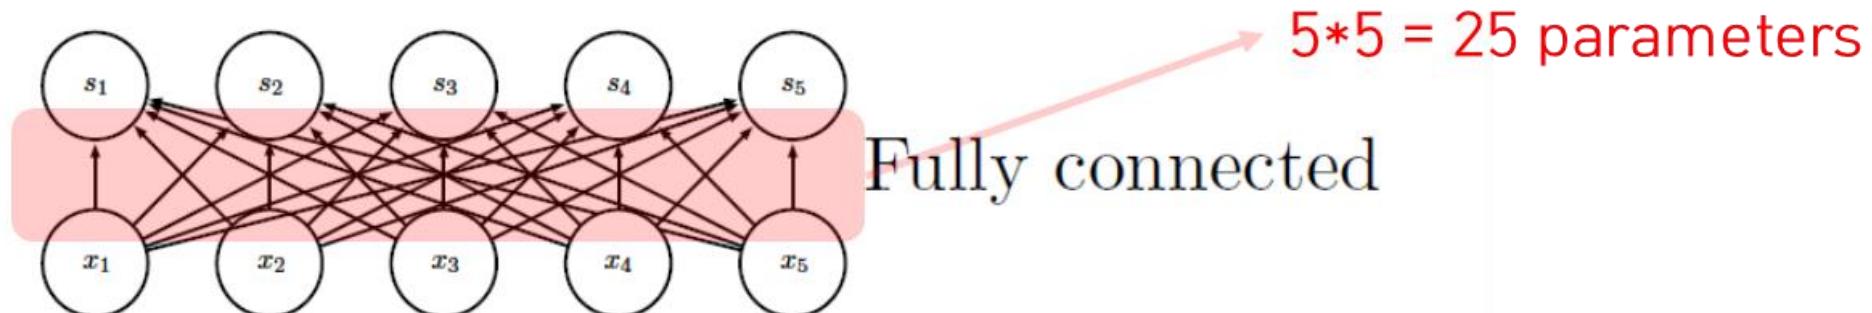
Receptive field of g_3 in the input layer = $\{x_1, x_2, x_3, x_4, x_5\}$



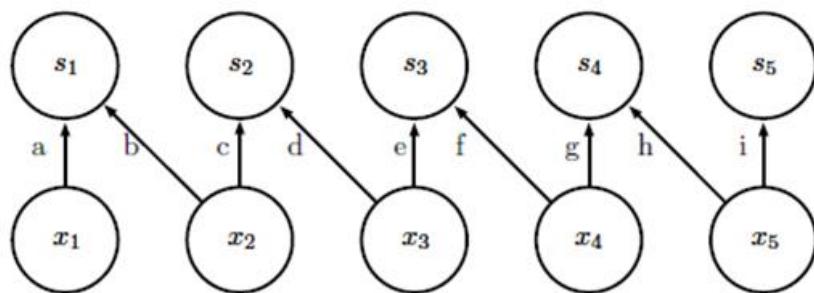
Example: a simple FNN (fully connected)
Receptive field of a hidden unit?



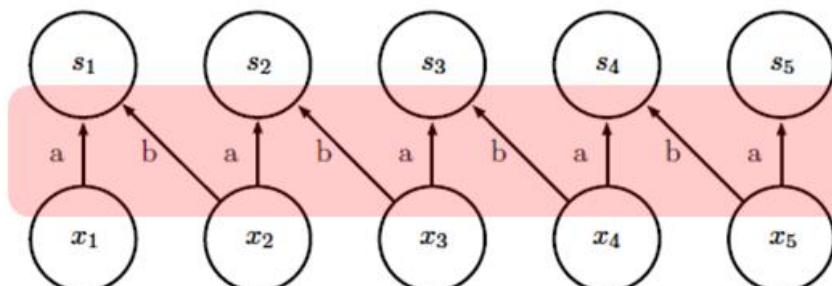
Properties of Convolution



Fully connected



Local connection:
like convolution,
but no sharing



Convolution

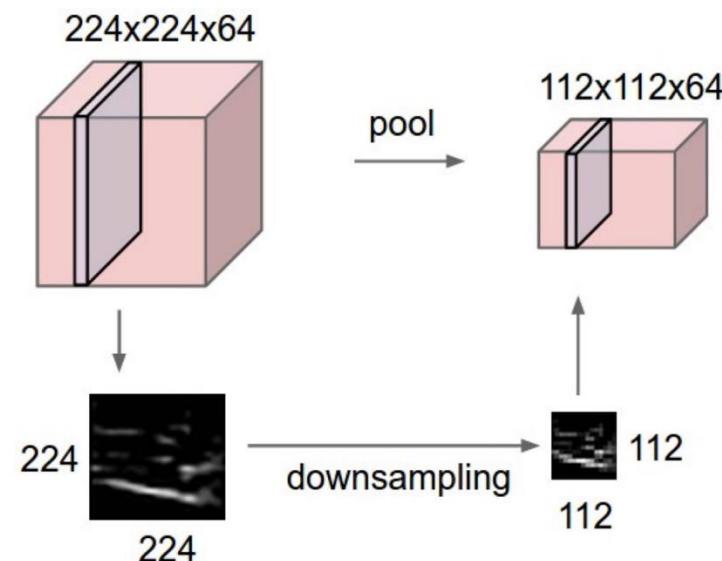
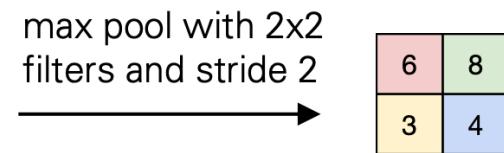
2 parameters

Pooling

- Pooling function
 - replaces the output of the layer at a certain location with a summary statistic of the nearby outputs.
 - makes the representation smaller and more manageable.
 - operates over each activation map independently.
 - Max pooling, average pooling, etc.

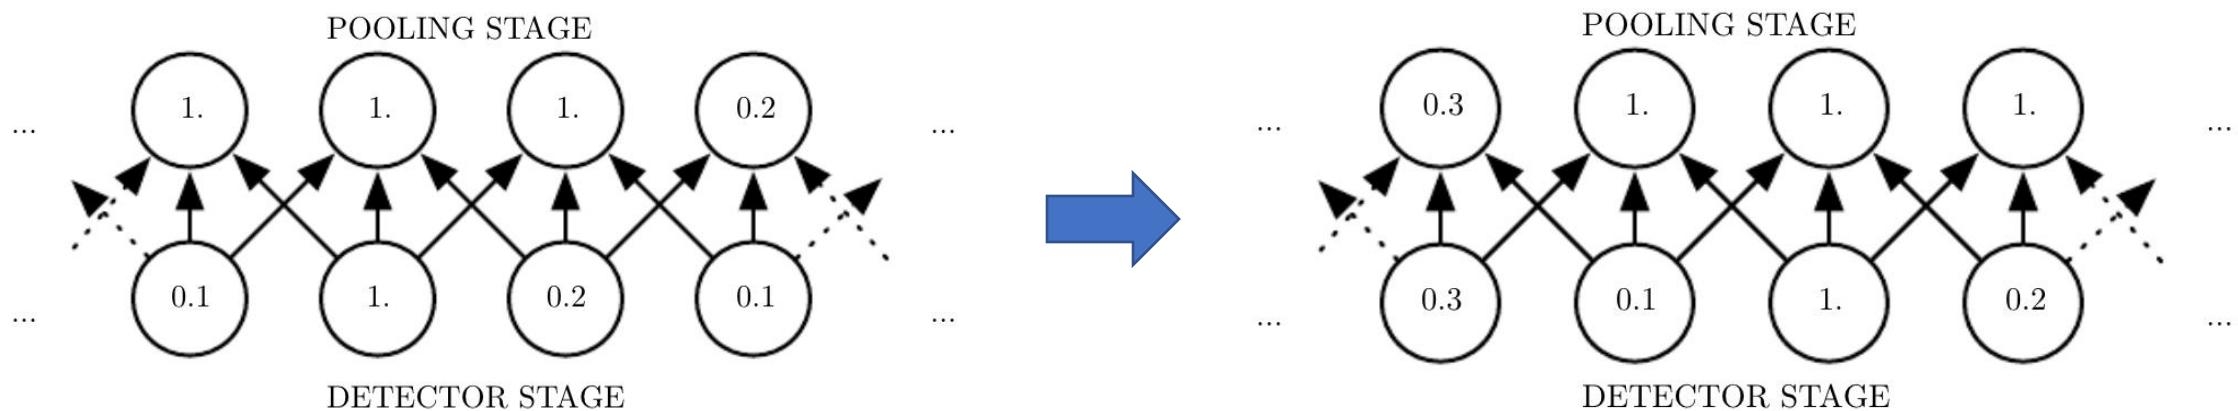
Single depth slice

1	1	2	4
5	6	7	8
3	2	1	0
1	2	3	4



Pooling

- Invariant to small translations of the input
 - It means that if we translate the input by a small amount, the values of most of the pooled outputs do not change.



Reading assignments

- “Dive into deep learning”
 - Chapter 7
- “Understanding deep learning”
 - Chapter 10