

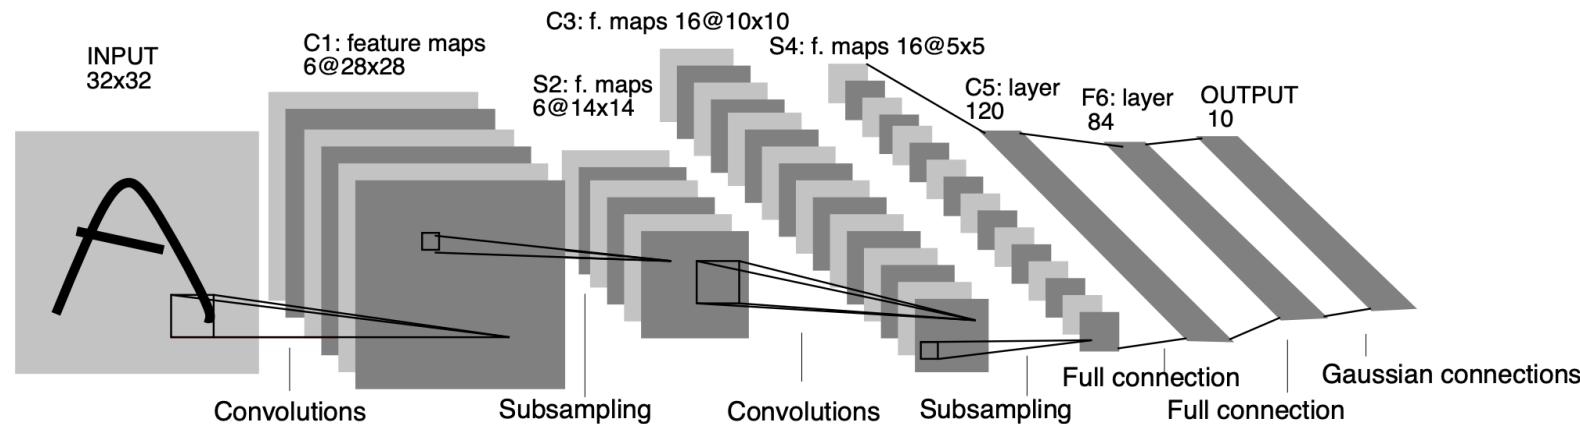
# Convolution and Pooling

# Computer Vision Tasks

- Classification
- Object detection
- Semantic or instance segmentation
- Others
  - Tracking in videos, camera pose estimation, body pose estimation, 3D reconstruction, denoising, super-resolution, auto-captioning, synthesis, etc.

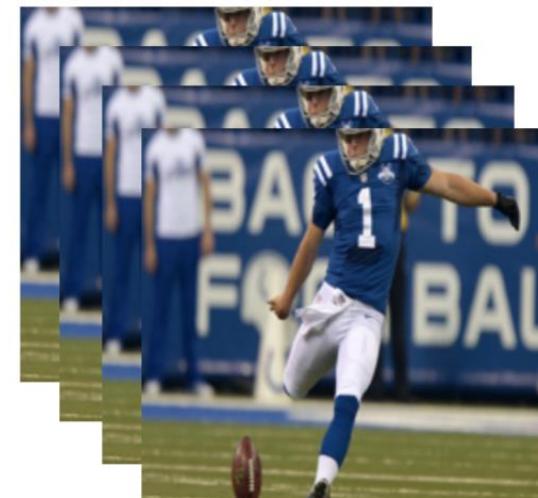
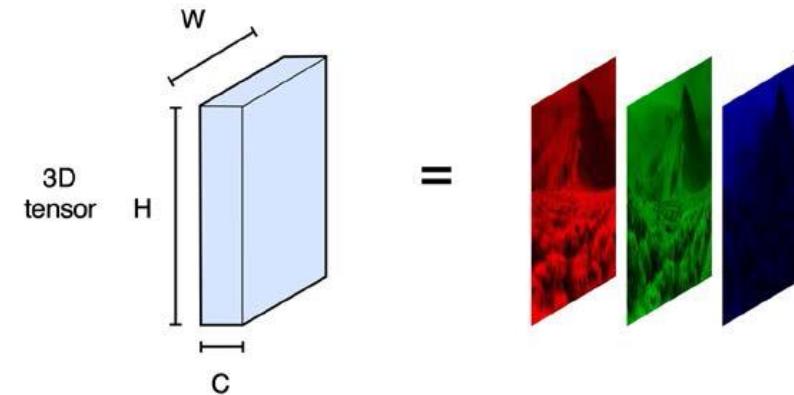
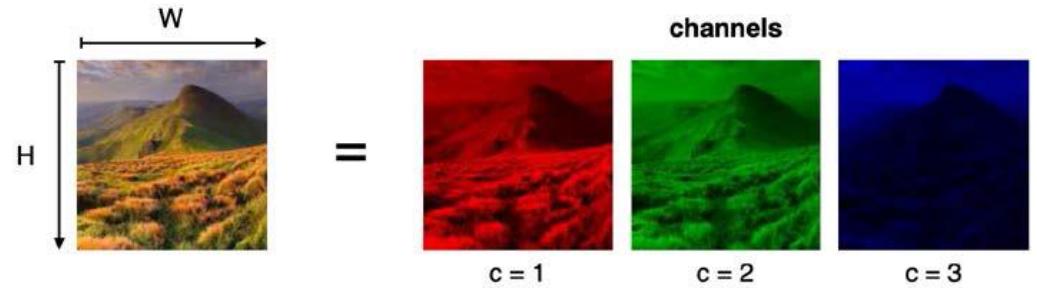
# Convolutional Neural Networks

- A **specialized** kind of neural network for processing data that has a known **grid-like topology** such as time-series data, image data, video data, etc.
- Ex. LeNet-5 (LeCun et al., 1998)



# Recap. Tensors

- 3D tensor = 3-dimensional array
  - RGB color image
  - (height, width, channel)
  - Grayscale image = 2D tensor
- 4D tensor = 4-dimensional array
  - Color video
  - (time, height, width, channel)



# Motivation

- A linear layer taking a 256\*256 RGB image as input, and producing an image of same size would require  $(256 \times 256 \times 3)^2 \cong 3.87 \times 10^{10}$ .
- Some input signals have some “invariance in translation”.
  - A function  $f$  of  $x$  is invariant to a transformation  $T$  if  $f(T(x)) = f(x)$ .
  - A transformation meaningful at a certain location can be used everywhere.
- A convolution layer embodies this idea.
  - It applies the same linear transformation locally, everywhere.

# Convolution

- Convolution is a spatial filtering.



$$\begin{matrix} & \xrightarrow{*^{1/8}} \\ \xleftarrow{*} & \end{matrix}$$

0	1	0
1	4	1
0	1	0

0	-1	0
-1	4	-1
0	-1	0

1	0	-1
2	0	-2
1	0	-1



# Convolution with 1D Array Input

- **kernel = [a,b]** <sup>learnable  
(a.k.a. filter)</sup> ← parameters
- input size m=6, kernel size k=2, stride(kernel step size) s=1, output size n=  $(m-k)/s+1=5$

$$\frac{6-2}{1} + 1 = 5$$

**output**

$ax_1+bx_2$	$ax_2+bx_3$	$ax_3+bx_4$	$ax_4+bx_5$	$ax_5+bx_6$
-------------	-------------	-------------	-------------	-------------

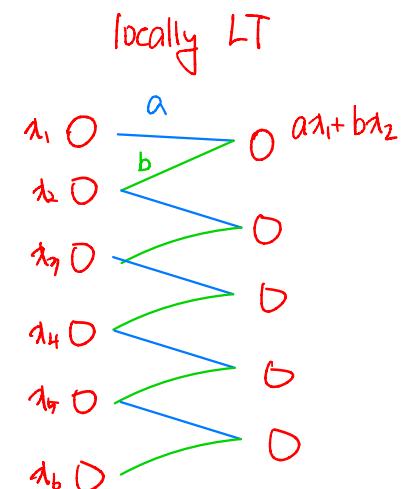
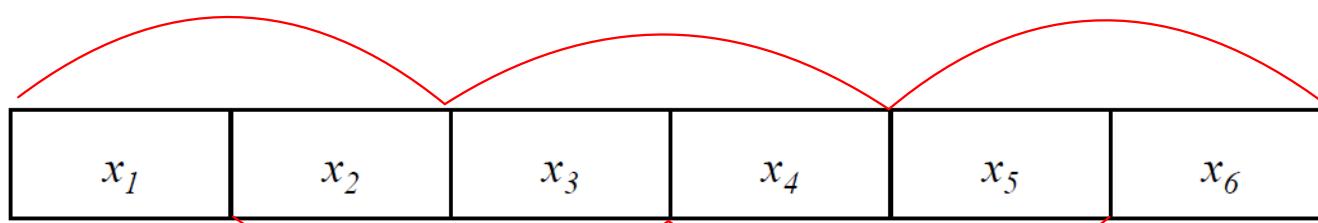
+bias

**kernel**

a	b
---	---

⇒ Shared all connections

**input**  
6 node



# Convolution with 1D Array Input

- **kernel = [a,b]** ← parameters
- input size m=6, kernel size k=2, stride(kernel step size) s=1, output size n=  $(m-k)/s+1=5$

**output**

$ax_1+bx_2$	$ax_2+bx_3$	$ax_3+bx_4$	$ax_4+bx_5$	$ax_5+bx_6$	+bias
-------------	-------------	-------------	-------------	-------------	-------

**kernel**

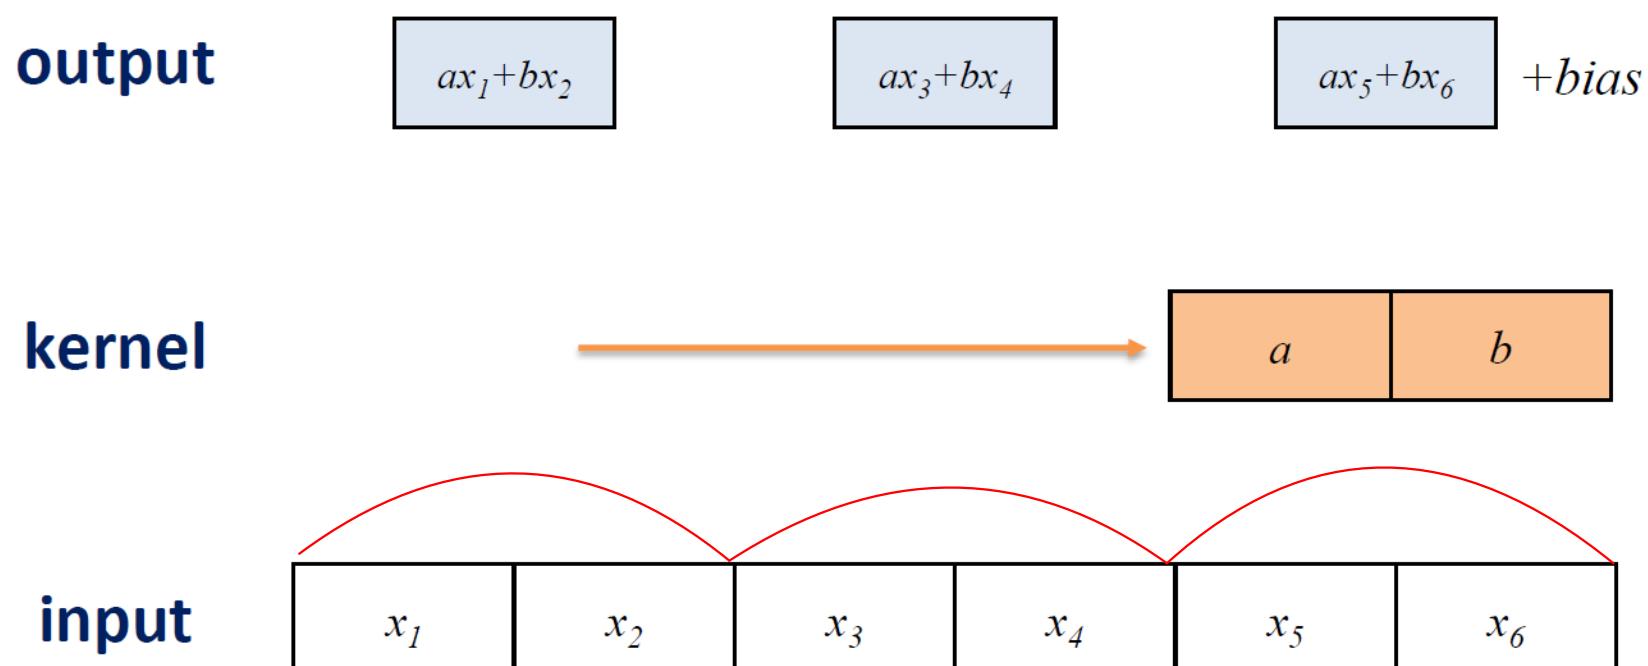


**input**

$x_1$	$x_2$	$x_3$	$x_4$	$x_5$	$x_6$
-------	-------	-------	-------	-------	-------

# Convolution with 1D Array Input

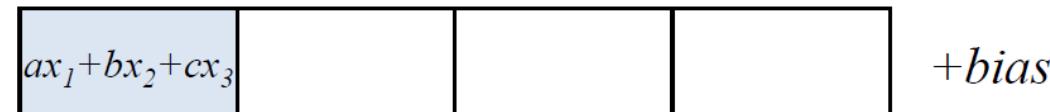
- **kernel = [a,b]** ← parameters
- input size m=6, kernel size k=2, **stride(kernel step size) s=2**, output size n=  $(m-k)/s+1=3$



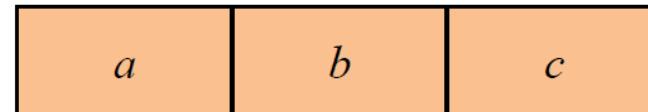
# Convolution with 1D Array Input

- **kernel** = [a,b,c] ← parameters
- input size m=6, **kernel size k=3**, stride(kernel step size) s=1, output size n=  $(m-k)/s+1=4$

**output**



**kernel**



**input**



# Convolution with 1D Array Input

- **kernel = [a,b,c]** ← parameters
- input size m=6, **kernel size k=3**, stride(kernel step size) s=1, output size n=  $(m-k)/s+1=4$

**output**

$$[ax_1+bx_2+cx_3 | ax_2+bx_3+cx_4 | ax_3+bx_4+cx_5 | ax_4+bx_5+cx_6] + bias$$

**kernel**



a	b	c
---	---	---

**input**

$x_1$	$x_2$	$x_3$	$x_4$	$x_5$	$x_6$
-------	-------	-------	-------	-------	-------

# How Computers See Images

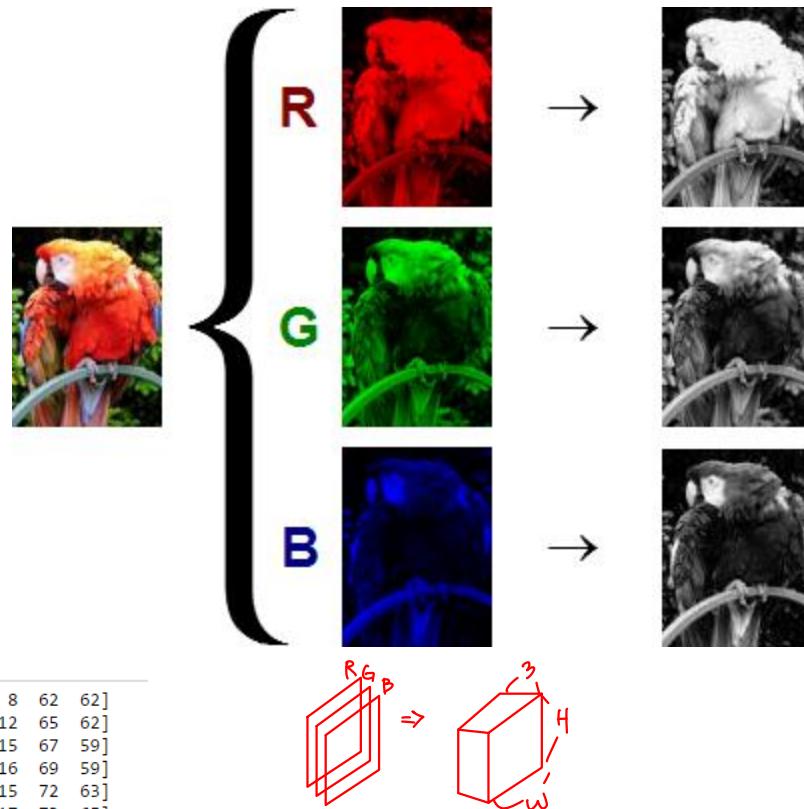
- Image is just an array of numbers.



2D tensor

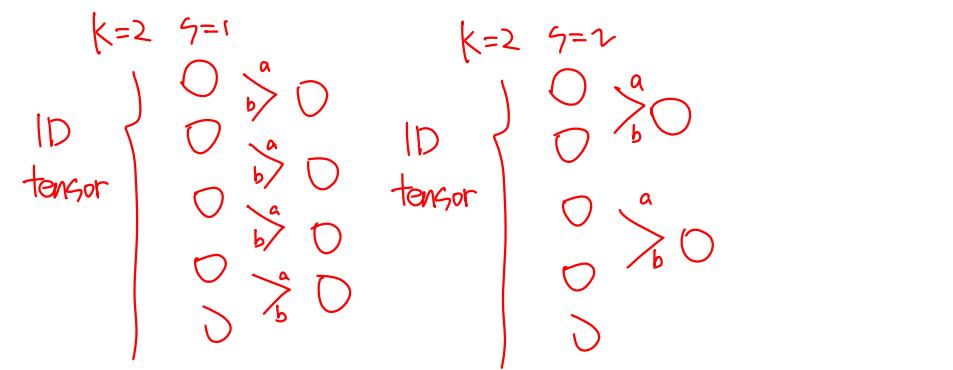
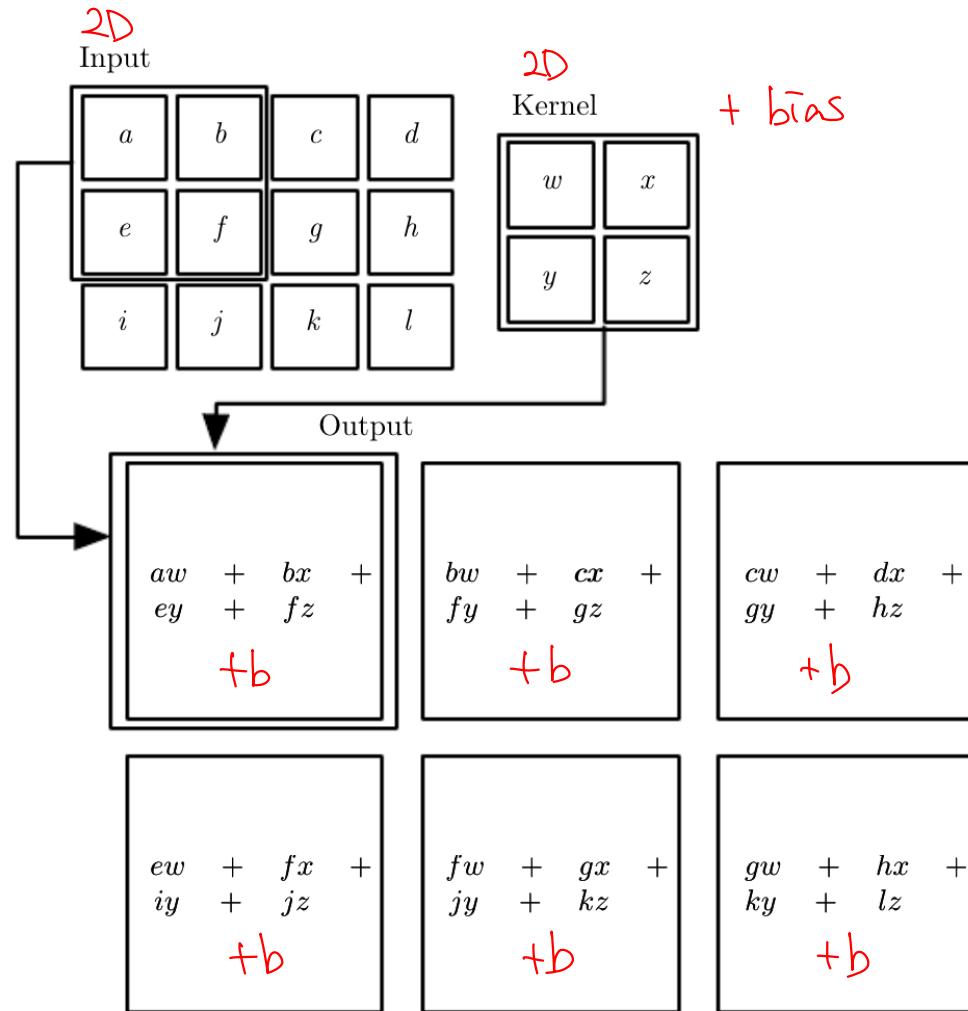
```
[[ 9  1  29  70 114  76  0  8  4  5  5  0 111 162  9  8  62  62]
 [ 3  0  33  61 102 106  34  0  0  0  0  49 182 150  1 12  65  62]
 [ 1  0  40  54 123  90  72  77  52  51  49 121 205  98  0 15  67  59]
 [ 3  1  41  57  74  54  96 181 220 170  90 149 208  56  0 16  69  59]
 [ 6  1  32  36  47  81  85  90 176 206 140 171 186  22  3 15  72  63]
 [ 4  1  31  39  66  71  71  97 147 214 203 190 198  22  6 17  73  65]
 [ 2  3  15  30  52  57  68 123 161 197 207 200 179  8  8 18  73  66]
 [ 2  2  17  37  34  40  78 103 148 187 205 225 165  1  8 19  76  68]
 [ 2  3  20  44  37  34  35  26  78 156 214 145 200  38  2 21  78  69]
 [ 2  2  20  34  21  43  70  21  43 139 205 93 211  70  0 23  78  72]
 [ 3  4  16  24  14  21 102 175 120 130 226 212 236  75  0 25  78  72]
 [ 6  5  13  21  28  28  97 216 184  90 196 255 255  84  4 24  79  74]
 [ 6  5  15  25  30  39  63 105 140  66 113 252 251  74  4 28  79  75]
 [ 5  5  16  32  38  57  69  85  93 120 128 251 255 154  19  26  80  76]
 [ 6  5  20  42  55  62  66  76  86 104 148 242 254 241  83  26  80  77]
 [ 2  3  20  38  55  64  69  80  78 109 195 247 252 255 172  40  78  77]
 [ 10  8  23  34  44  64  88 104 119 173 234 247 253 254 227  66  74  74]
 [ 32  6  24  37  45  63  85 114 154 196 226 245 251 252 250 112  66  71]]
```

3 channels



<https://savan77.github.io/blog/how-computers-see-image.html>

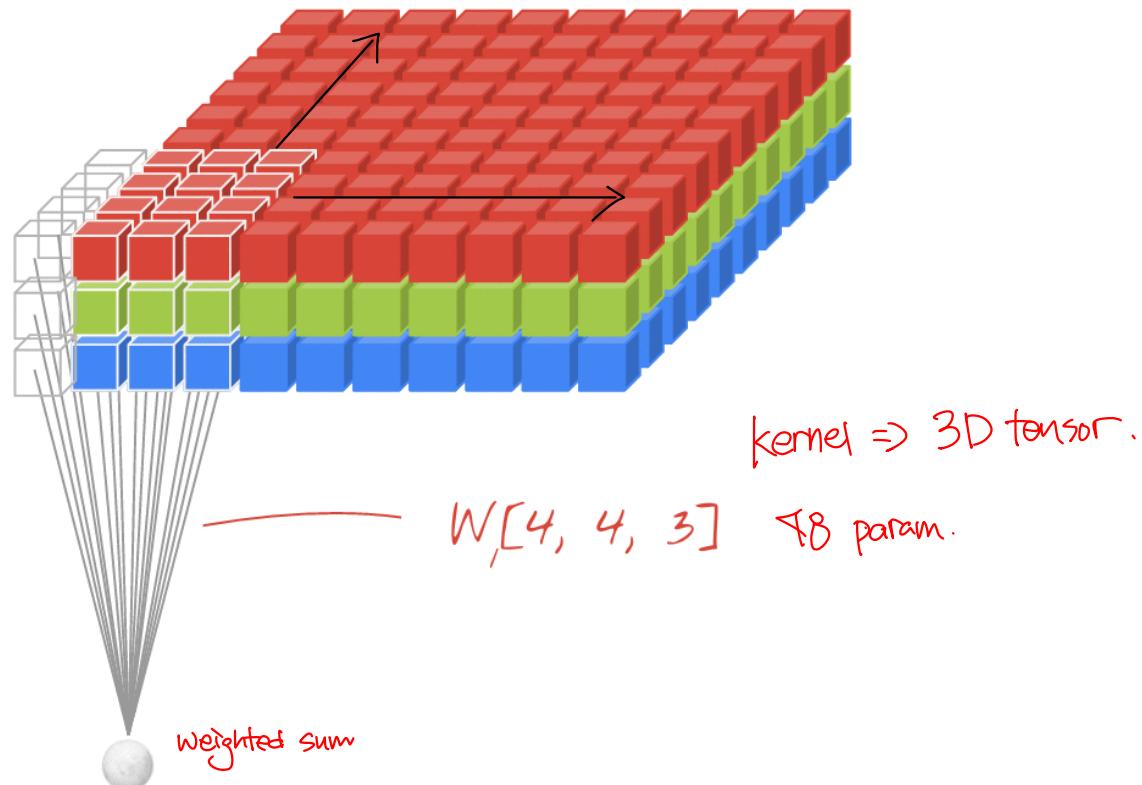
# Convolution with 2D Tensor



kernel & bias  
What are parameters? ( $w \rightarrow y \Leftarrow b$ )  
What are hyperparameters?  
kernel size  
stride  
# output node  
indirectly specifying!

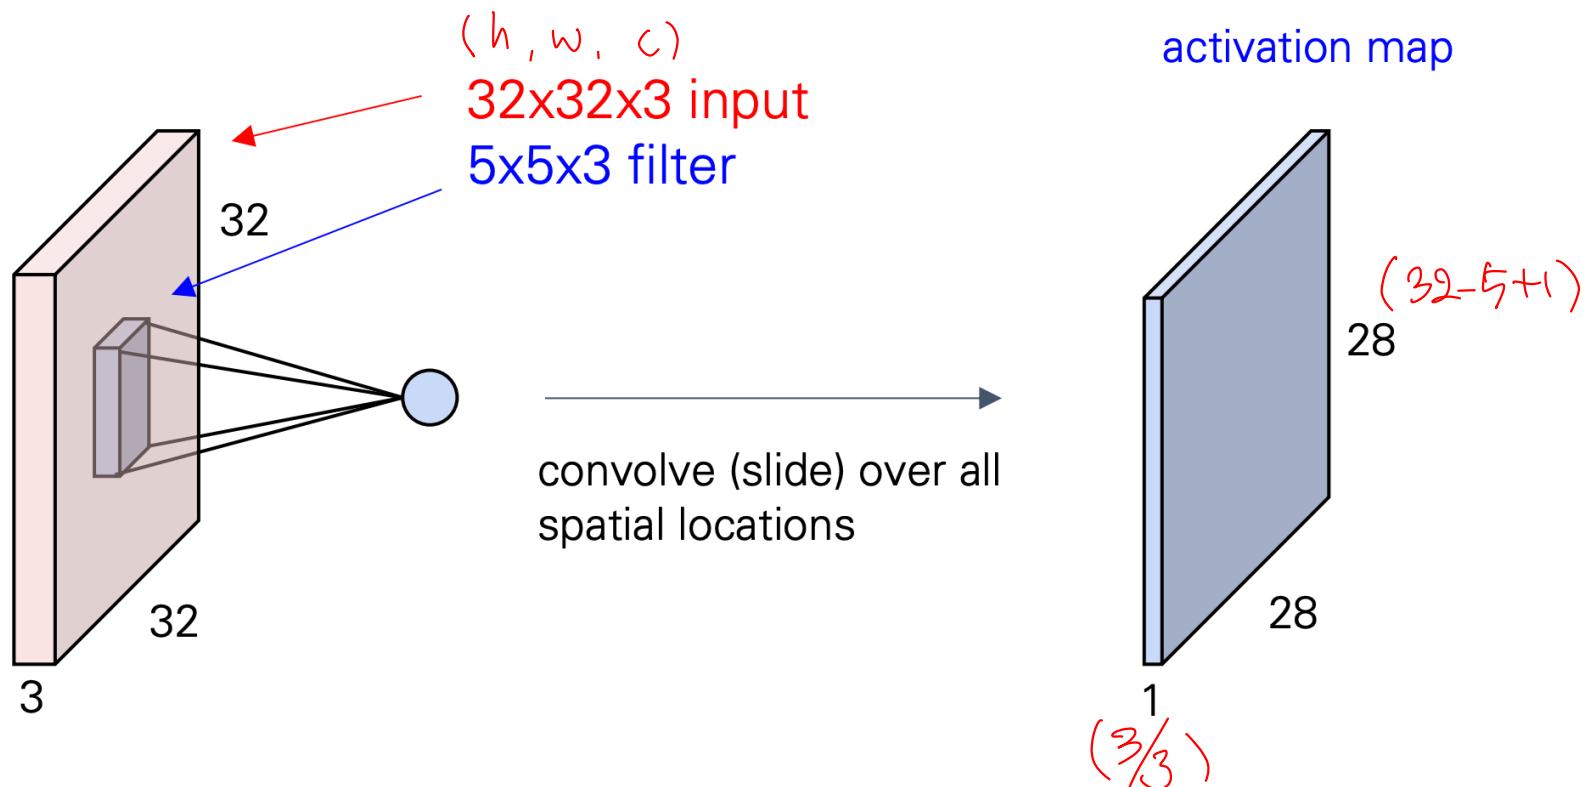
# Convolution with 3D Tensor

$$(H, W) \Rightarrow d_f = V.$$



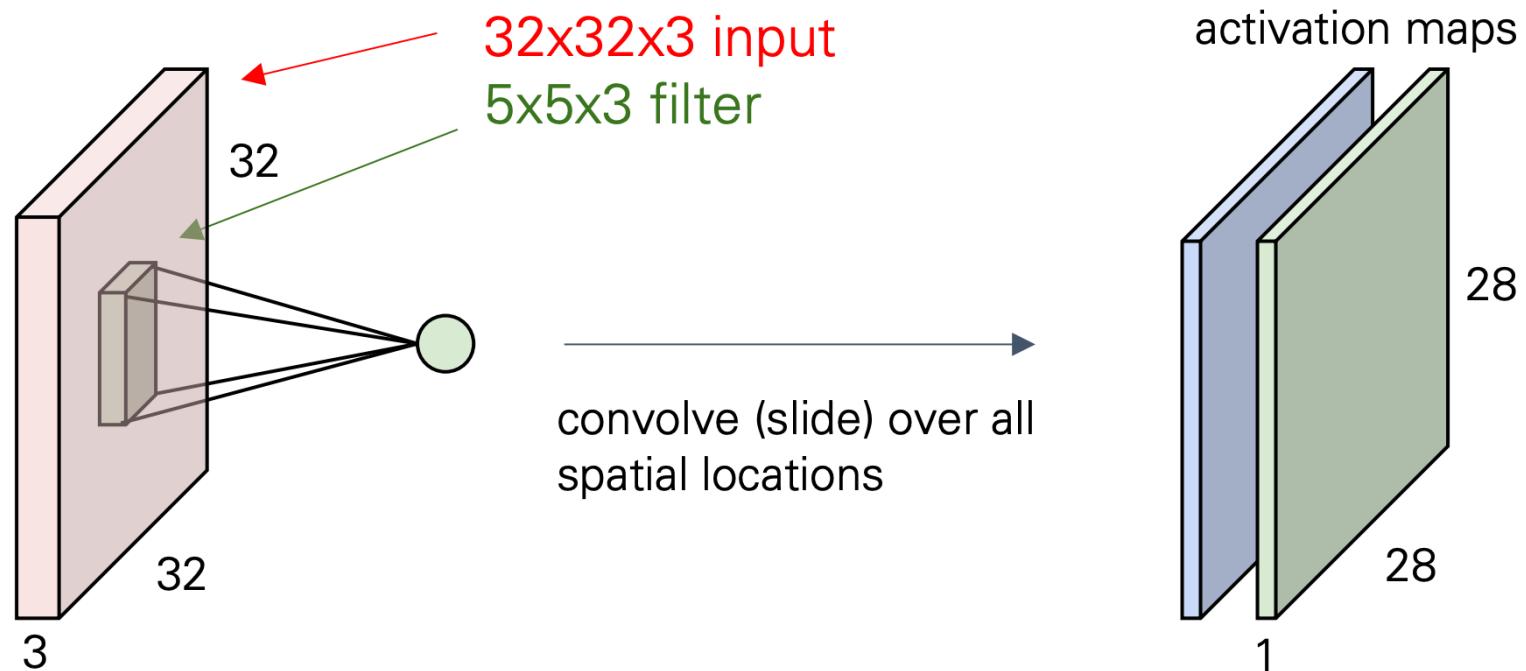
# Convolution with 3D Tensor

- Generally, we use multiple kernels for single convolutional layer.



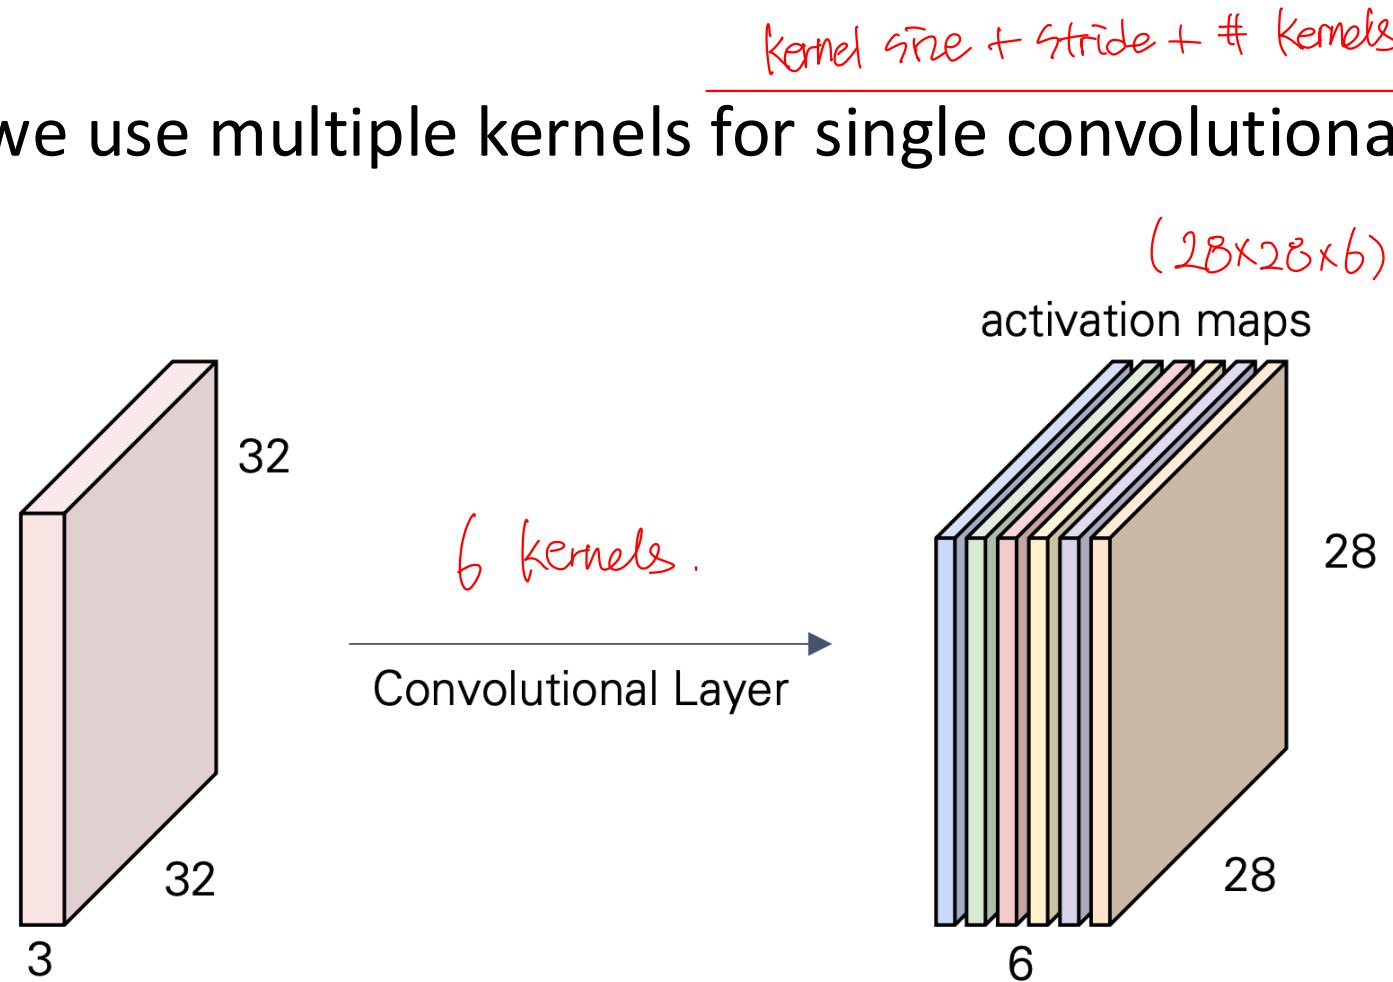
# Convolution with 3D Tensor

- Generally, we use multiple kernels for single convolutional layer.



# Convolution with 3D Tensor

- Generally, we use multiple kernels for single convolutional layer.



# Parameters in Convolution

- The **padding** specifies the size of a zeroed frame added around the input.
- The **stride** specifies a step size when moving the kernel across the signal.
- The **dilation** modulates the expansion of the filter without adding weights.  
optional `in skip`

# Parameters in Convolution: ~~A~~ Padding

- There is border effects in convolution.
- Valid convolution

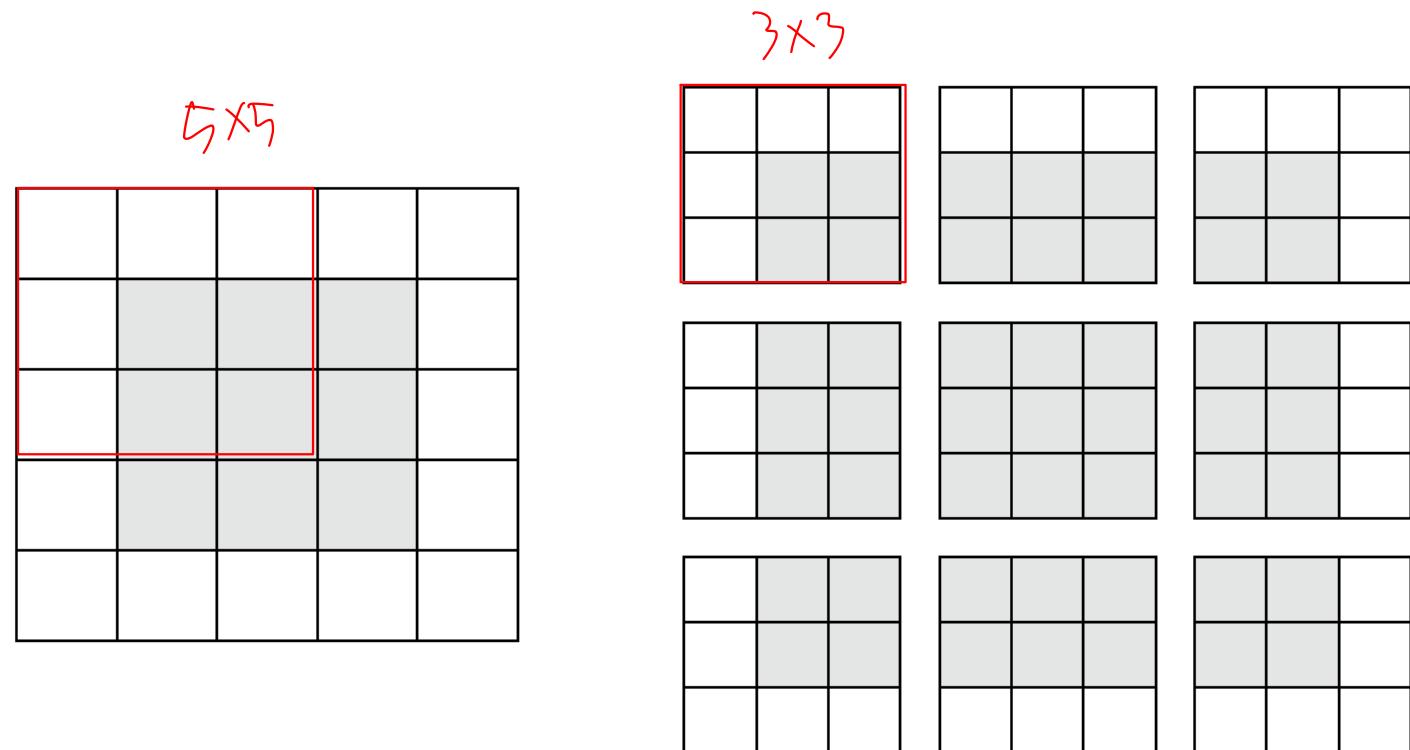


Figure 8.5 Valid locations of  $3 \times 3$  patches in a  $5 \times 5$  input feature map

# Parameters in Convolution: Padding

- Padding to an input

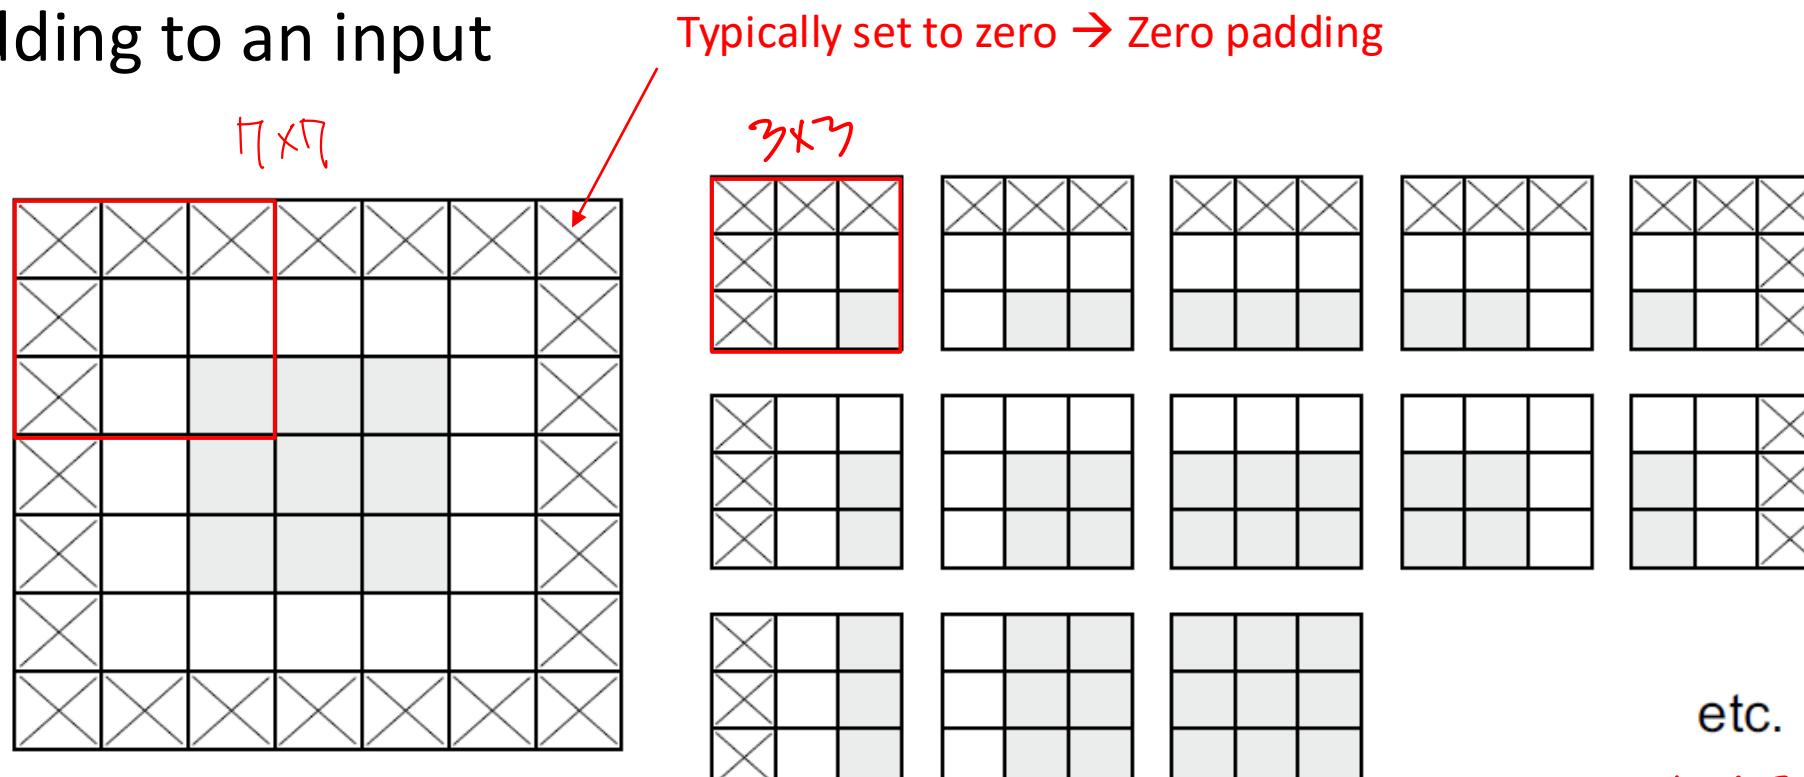


Figure 8.6 Padding a  $5 \times 5$  input in order to be able to extract 25  $3 \times 3$  patches

$$\begin{aligned} & (\underline{5}-\underline{3}+1) \times (\underline{5}-\underline{3}+1) \\ & = 2^5 \text{ patches} \end{aligned}$$

# Parameters in Convolution: Stride

- Strided convolution: convolutions with a stride higher than 1

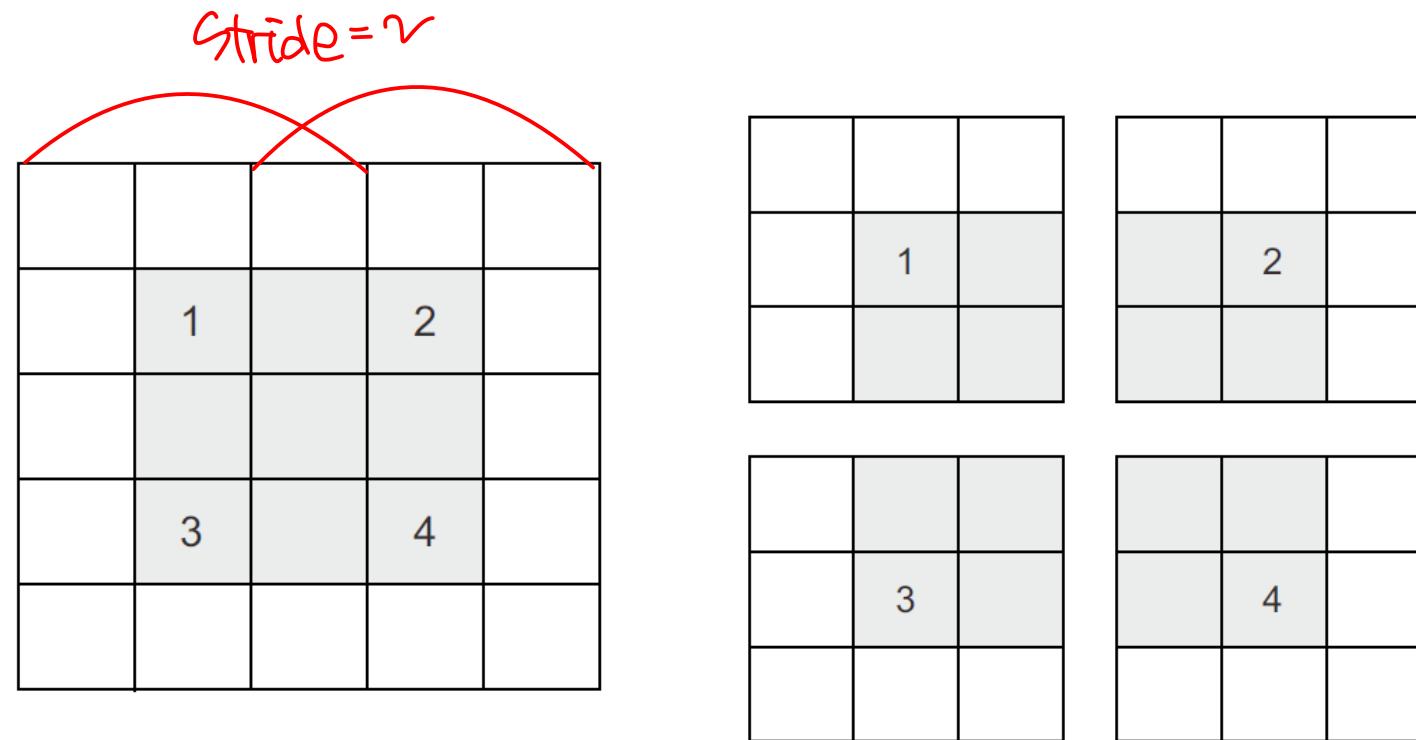
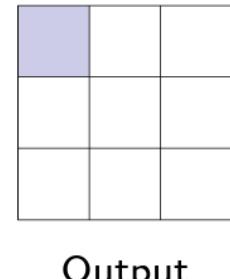
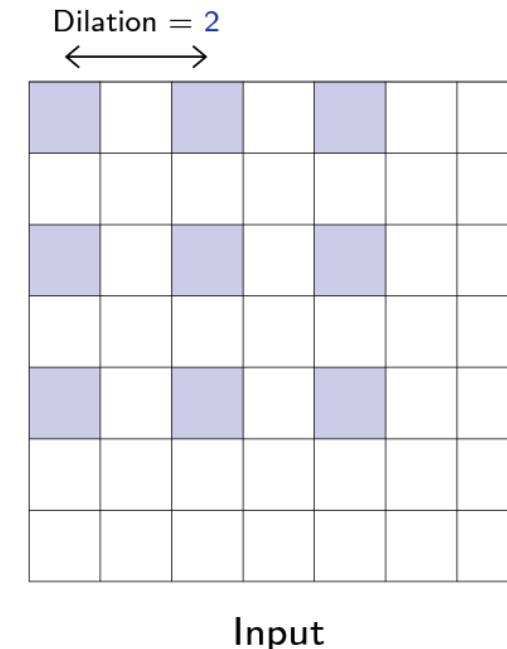
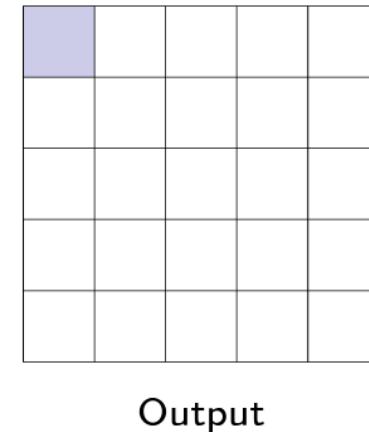
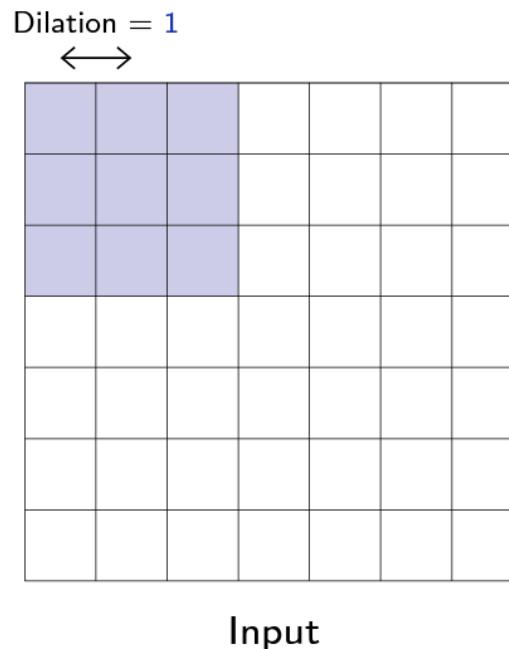


Figure 8.7  $3 \times 3$  convolution patches with  $2 \times 2$  strides

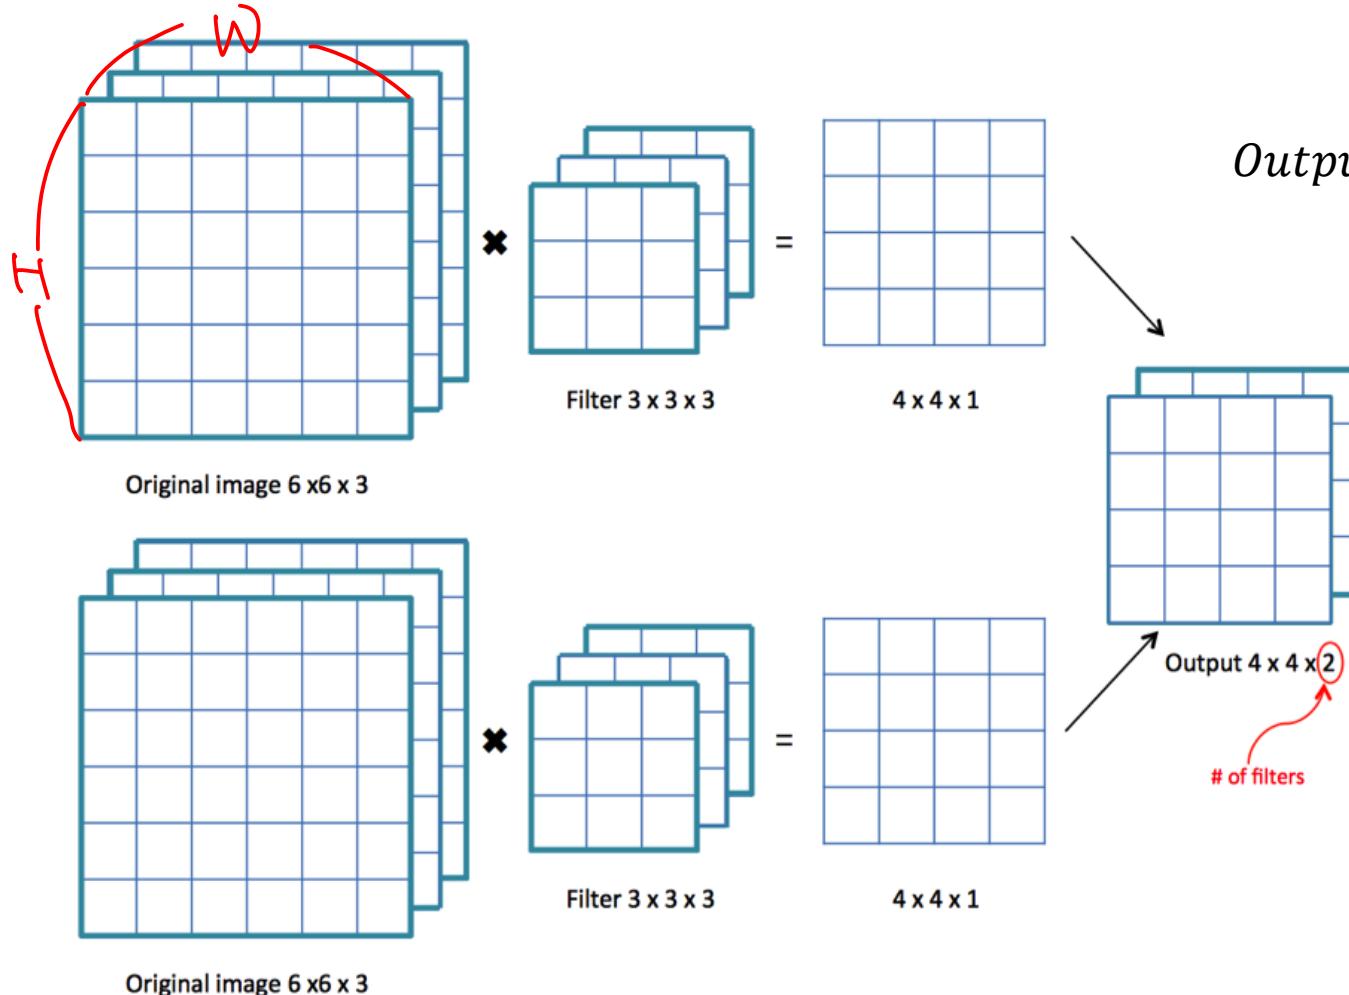
*Skip*

# Parameters in Convolution: Dilation

- The expansion of a filter by adding rows and columns of zeros between coefficient.



# How to Compute Output Shape



$$\text{Output Size} = \left( \frac{H + \underset{\text{padding}}{2P} - \underset{\text{filter size}}{F}}{\underset{\text{Stride}}{S}} + 1 \right) \times \left( \frac{W + \underset{\text{padding}}{2P} - \underset{\text{filter size}}{F}}{\underset{\text{Stride}}{S}} + 1 \right)$$

$$\left( \frac{6+0-3}{1} + 1 \right) \left( \frac{6+0-3}{1} + 1 \right) = 16$$

# Properties of Convolution

- A convolution preserves the signal support structure. ?
- Sparse interactions
  - Inputs and outputs are not fully connected but have local connectivity
- Parameter sharing
  - The same kernel is used repeatedly.
- Equivariance to transition
  - $\text{convolution}(\text{shift}(\text{input})) = \text{shift}(\text{convolution}(\text{input}))$

?

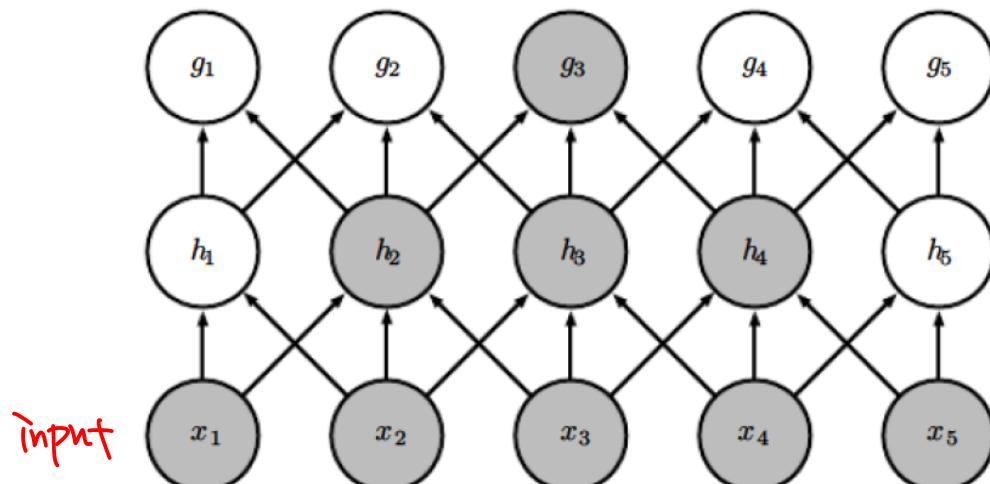
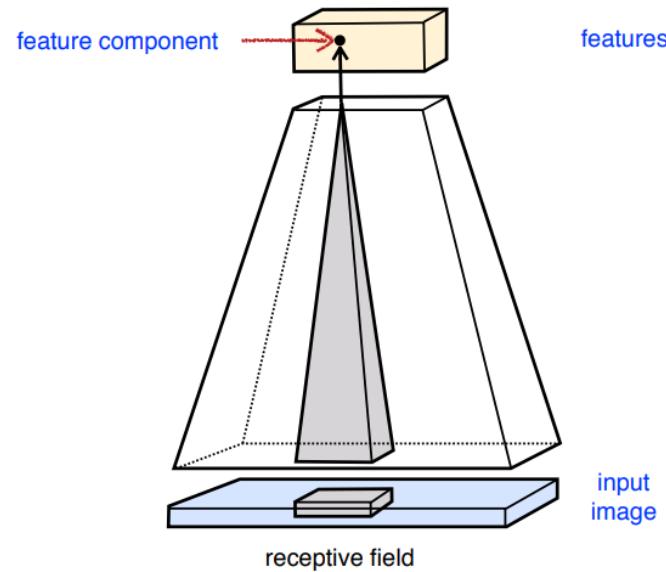
- **Receptive field: Spatial locality**

- Each element of the feature map process only for its receptive field (a local region of the input)
- higher kernel size  $k \rightarrow$  larger receptive field
- Higher-level layers  $\rightarrow$  larger receptive field

**Example:** a simple CNN with  $k=3$

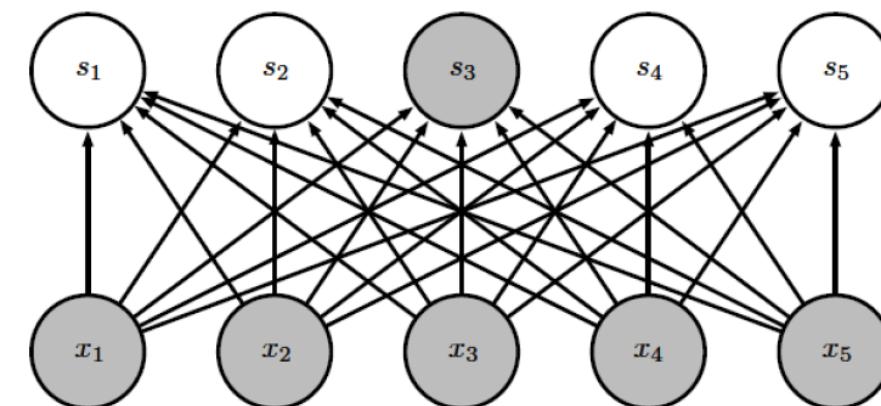
Receptive field of  $h_3$  in the input layer = { $x_2, x_3, x_4$ }

Receptive field of  $g_3$  in the input layer = { $x_1, x_2, x_3, x_4, x_5$ }

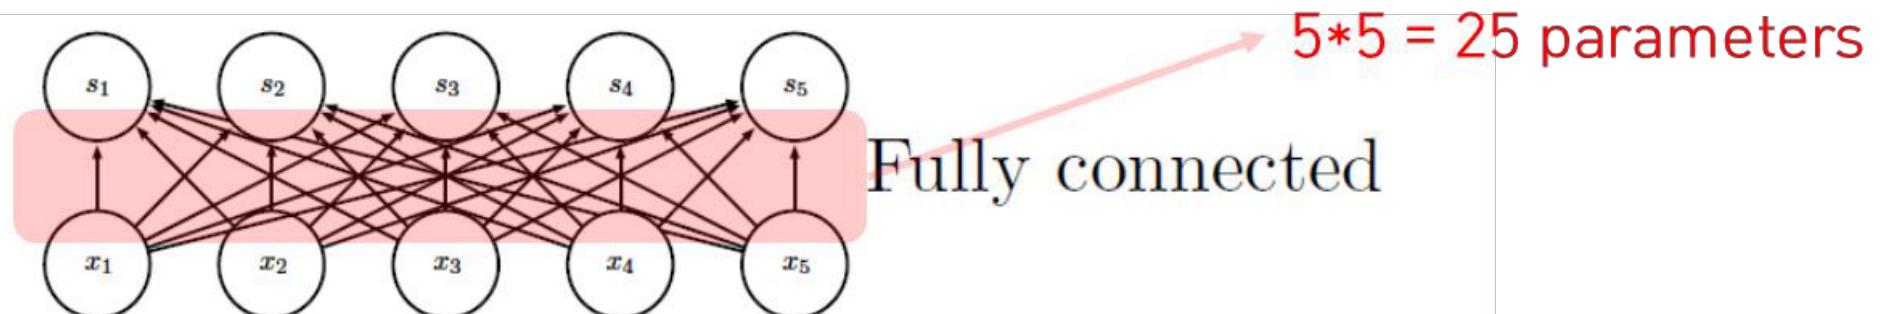


**Example:** a simple FNN (fully connected)

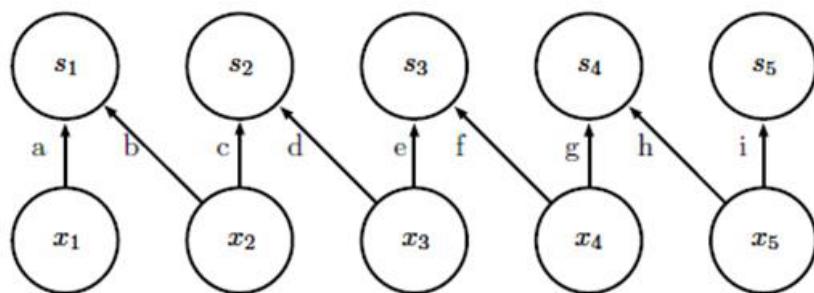
Receptive field of a hidden unit?  $\{x_1, x_2, x_3, x_4, x_5\}$



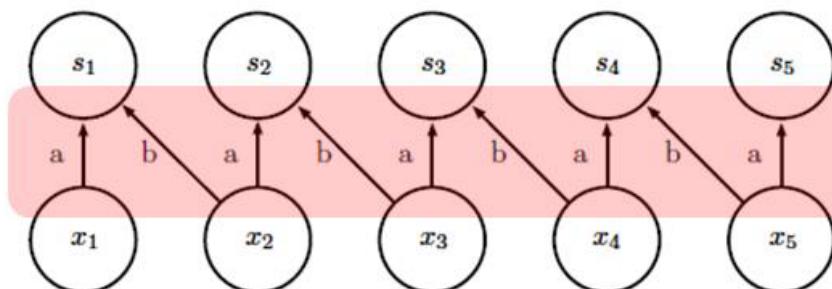
# Properties of Convolution



Fully connected



Local connection:  
like convolution,  
but no sharing



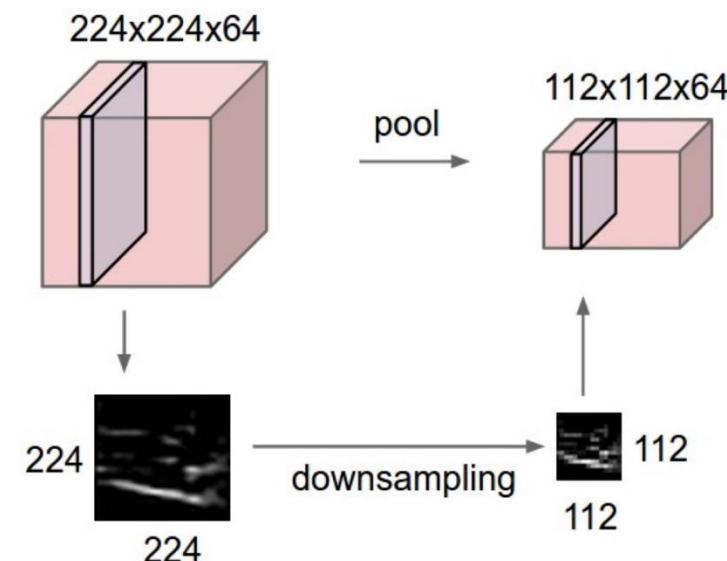
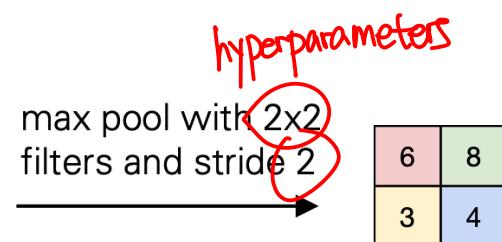
Convolution

# Pooling

- Pooling function
  - replaces the output of the layer at a certain location with a summary statistic of the nearby outputs.
  - makes the representation smaller and more manageable.
  - operates over each activation map independently. ?
  - Max pooling, average pooling, etc.

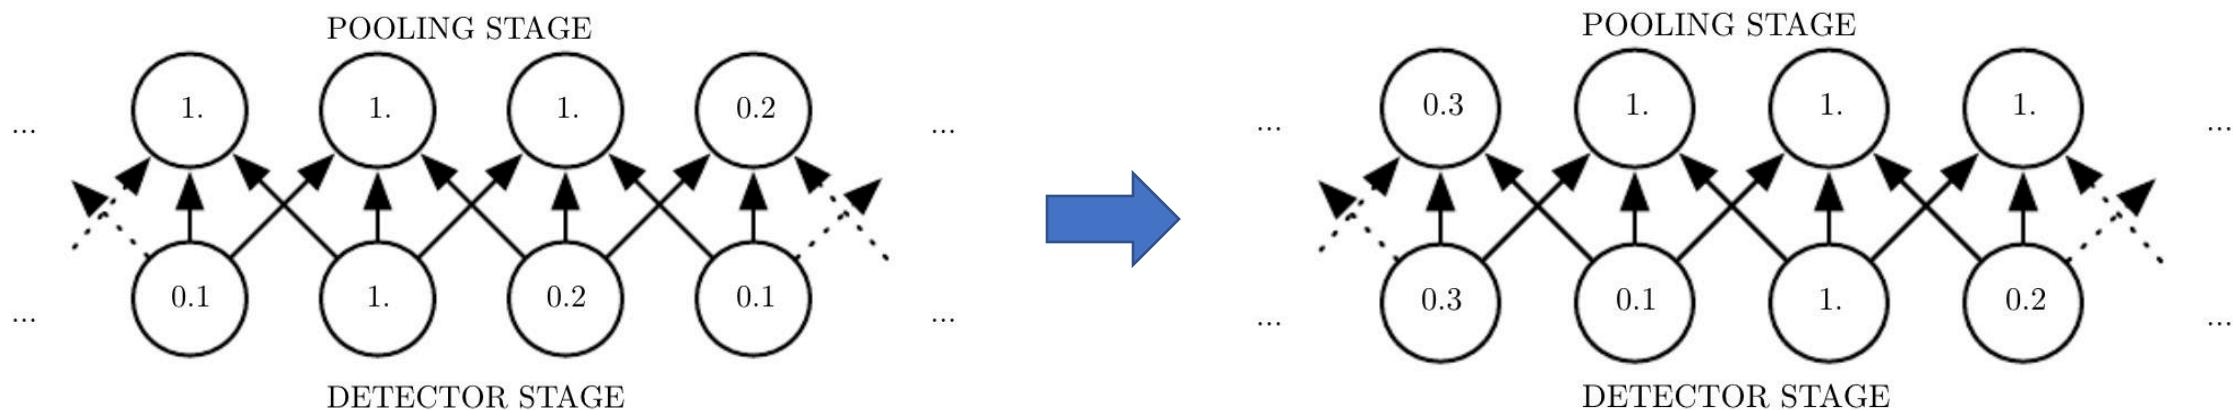
Single depth slice

1	1	2	4
5	6	7	8
3	2	1	0
1	2	3	4



# Pooling

- Invariant to small translations of the input
  - It means that if we translate the input by a small amount, the values of most of the pooled outputs do not change.



# Reading assignments

- “Dive into deep learning”
  - Chapter 7
- “Understanding deep learning”
  - Chapter 10

MLP

⇒ compute gradient

“chain rule” terms

why  $z_j \rightarrow 1$ ?

$$\frac{\partial L}{\partial w_{ji}} \rightarrow 0$$

Vanishing gradient

GyMOD activation

→ drawback

→ snap  
relu

perceptron

backprop

grad computation

overfit / regularization

techniques dropout

# Reading assignments

- “Dive into deep learning”
  - Chapter 7
- “Understanding deep learning”
  - Chapter 10

compute  
# trainable params

modern CNN (VGGNet ResNet)  
architecture