

- Deep learning models that can process text (sequences of words or sequences of characters), time-series, etc.
- Two fundamental approaches for sequence processing: *recurrent neural networks* and *1D convnets*
- Applications
 - Document/time-series classification
 - Sequence-to-sequence learning such as machine translation
 - Sentiment analysis

6.0 Sequential data and its representation

- Sequential data

 No description has been provided for this image

- Sequence representation

- Bag of words

 No description has been provided for this image

- How can we deal with an order in sequence?
 - Concatenate one-hot vectors

 No description has been provided for this image

- To develop a model for sequence data, we need

- to deal with variable length sequences
 - to maintain sequence order
 - to keep track of long-term dependencies
 - to share parameters across the sequence


6.1 Working with text data

- Text is one of the most common forms of sequence data: a sequence of characters or a sequence of words.
- *Vectorizing* text is the process of transforming text into numeric tensors.
 - Segment text into words, and transform each word into a vector.
 - Segment text into characters, and transform each character into a vector.
 - Extract n-grams of words or characters, and transform each n-gram into a vector.
 - *N-grams* are overlapping groups of multiple consecutive words or characters.
 - Example. "The cat sat on the mat."

- A bag-of-2-grams = {"The", "The cat", "cat", "cat sat", "sat", "sat on", "on", "on the", "the", "the mat", "mat"}
- A bag-of-3-grams = {"The", "The cat", "cat", "cat sat", "The cat sat", "sat", "sat on", "on", "cat sat on", "on the", "the", "sat on the", "the mat", "mat", "on the mat"}
- The different units into which you can break down text (words, characters, or n-grams) are called *tokens*, and breaking text into such tokens is called *tokenization*.
- All text-vectorization processes consist of applying some tokenization scheme and then associating numeric vectors with the generated tokens.
 - These vectors, packed into sequence tensors, are fed into deep neural networks.

Tokenization

- One of the most important steps in text preprocessing
- The process of splitting a phrase, sentence, paragraph, one or multiple text documents into smaller units
 - Each of these smaller units is called a token.
 - Tokens can be anything: a word, a subword, or even a character.



- Word-based tokenization
 - The most commonly used delimiter is space.
 - "Is it weird I don't like coffee?" --> ["Is", "it", "weird", "I", "don't", "like", "coffee?"]
 - (if punctuation into account) --> ["Is", "it", "weird", "I", "don", "", "t", "like", "coffee", "?"]
 - Simple, but a big vocabulary (a huge embedding matrix)
 - What if we limit the number of words that can be added to the vocabulary?
- Character-based tokenization
 - Split the raw text into individual characters
 - This results in a very small vocabulary.
 - Drawbacks
 - A character usually don't carry any meaning or information.
 - Reducing the vocabulary size has a trade-off-with the sequence length.
- Subword-based tokenization
 - A solution between word and character-based tokenization

- Word-level: large vocab size, OOV tokens, misspelled words
 - Char-level: long sequences, less meaningful individual tokens
- Principle
 - Do not split the frequently used words into smaller subwords.
 - Split the rare words into smaller meaningful subwords.
 - E.g., "boys" --> ["boy", "s"]
- It is even possible for a model to process a word which it has never seen before as the decomposition can lead to known subwords.

- Byte-Pair Encoding (BPE)
 - BPE was originally a data compression algorithm proposed in 1994.
 - BPE is a powerful tool for subword tokenization.
 - The core idea
 - Iteratively find the most frequent pair of adjacent characters (or bytes) in the data and merge them into a single, new token.
 - Through this process, sequence of characters that frequently appear together are gradually grouped into meaningful subword units.
 - Process
 - Pre-tokenization --> measure frequency --> merge
 - Suppose that our corpus uses these five words.
 - ["hug", "pug", "pun", "bun", "hugs"]
 - The base vocab = ["b", "g", "h", "n", "p", "s", "u"]
 - Base vocab will contain all the ASCII characters or Unicode characters.
 - Some characters will be converted to the unknown token.
 - E.g., emojis
 - Byte-level BPE can solve this issue.
 - With base vocab, we add new tokens until the desired vocab size is reached by learning merges.
 - merge = rules to merge two elements of the existing vocab together into a new one
 - During the tokenizer training, the BPE algorithm searches for the most frequent pair of existing tokens.
 - That most frequent pair is the one that will be merged, and we rinse and repeat for the next step.

- BPE example
 - 1. Assume that the words have the following frequencies.
 - ("hug", 10), ("pug", 5), ("pun", 12), ("bun", 4), ("hugs", 5)
 - 2. Split each word into characters (tokens in our base vocab)
 - ("h" "u" "g", 10), ("p" "u" "g", 5), ("p" "u" "n", 12), ("b" "u" "n", 4), ("h" "u" "g" "s", 5)
 - 3. Merge the most frequent pair: ("u", "g") -> "ug"
 - Vocab: ["b", "g", "h", "n", "p", "s", "u", "ug"]
 - Corpus: ("h" "ug", 10), ("p" "ug", 5), ("p" "u" "n", 12), ("b" "u" "n", 4), ("h" "ug" "s", 5)
 - 4. Merge the most frequent pair: ("u", "n") -> "un"
 - Vocab: ["b", "g", "h", "n", "p", "s", "u", "ug", "un"]
 - Corpus: ("h" "ug", 10), ("p" "ug", 5), ("p" "un", 12), ("b" "un", 4), ("h" "ug" "s", 5)
 - 5. Merge the most frequent pair: ("h", "ug") -> "hug"
 - Vocab: ["b", "g", "h", "n", "p", "s", "u", "ug", "un", "hug"]
 - Corpus: ("hug", 10), ("p" "ug", 5), ("p" "un", 12), ("b" "un", 4), ("hug" "s", 5)

- Continue until we reach the desired vocab size.
- Given the following vocab:
 - Vocab: ["b", "g", "h", "n", "p", "s", "u", "ug", "un", "hug"]
 - Corpus: ("hug", 10), ("p" "ug", 5), ("p" "un", 12), ("b" "un", 4), ("hug" "s", 5)
- New inputs are tokenized as follows:
 - "bug" -> ["b", "ug"]
 - "mug" -> ["[UNK]", "ug"]
 - "thug" -> ["[UNK]", "hug"]
 - "unhug" -> ?

```
In [ ]: # =====
# Exercise: BPE-based Tokenizer (OpenAI's Tiktoken)
# - The method used by models like GPT-4 and GPT-3.5-turbo.
# - Feature: Sophisticated Regex Pre-tokenization + Byte-Level BPE.
# =====
import tiktoken

print("=*25 + " Exercise 1: BPE-based Tokenizer (tiktoken) " + "=*25)

# 1. Load the tokenizer
# 'cl100k_base' is the encoding name used by OpenAI for models like GPT-4.
tokenizer_bpe = tiktoken.get_encoding("cl100k_base")

# 2. Define a sample text with English, numbers, and Korean
text = "Hello, world! 2024 is the year of AI. 안녕하세요!"

# 3. Encode (Text -> Token IDs)
encoded_ids_bpe = tokenizer_bpe.encode(text)

print(f"Original Text: {text}")
print(f"BPE Encoding Result (Token IDs): {encoded_ids_bpe}")
print(f"Total BPE Tokens: {len(encoded_ids_bpe)}\n")

# 4. Decode (Token IDs -> Text)
decoded_text_bpe = tokenizer_bpe.decode(encoded_ids_bpe)

print(f"BPE Decoding Result: {decoded_text_bpe}\n")

# 5. Inspect the individual tokens
print("Inspecting individual BPE tokens:")
tokens_bpe = [tokenizer_bpe.decode([token_id]) for token_id in encoded_ids_bpe]
print(tokens_bpe)

print("\n[Analysis Point 1] Note how the leading space is part of the token itself (e.g., ' world')
print("[Analysis Point 2] Note how '안녕하세요' (Korean) is broken down into multiple, less intuit
```

```
===== Exercise 1: BPE-based Tokenizer (tiktoken) =====
Original Text: Hello, world! 2024 is the year of AI. 안녕하세요!
BPE Encoding Result (Token IDs): [9906, 11, 1917, 0, 220, 2366, 19, 374, 279, 1060, 315, 15592, 1
3, 96270, 75265, 243, 92245, 0]
Total BPE Tokens: 18
```

BPE Decoding Result: Hello, world! 2024 is the year of AI. 안녕하세요!

Inspecting individual BPE tokens:

```
['Hello', ',', ' world', '!', ' ', '202', '4', ' is', ' the', ' year', ' of', ' AI', '.', ' 안',
'?', '?', '하세요', '!']
```

[Analysis Point 1] Note how the leading space is part of the token itself (e.g., ' world').

[Analysis Point 2] Note how '안녕하세요' (Korean) is broken down into multiple, less intuitive byte-level tokens. This is inefficient but prevents 'unknown token' errors.

Vectorization

- There are multiple ways to associate a vector with a token.
 - *one-hot encoding* of tokens
 - *token embedding* (typically used exclusively for words, and called *word embedding*)

|  No description has been provided for this image

One-hot encoding

- The most common, most basic way to turn a token into a vector
- It consists of associating a unique integer index with every token and then turning this integer index i into a binary vector of size N (the size of the vocabulary).
 - The vector is all zeros except for the i th entry, which is 1.

```
In [ ]: # Word-Level one-hot encoding

import numpy as np

samples = ['The cat sat on the mat.', 'The dog ate my homework.']

# build vocabulary dictionary
token_index = {}
for sample in samples:
    for word in sample.split():
        if word not in token_index:
            token_index[word] = len(token_index)+1 # assign a unique index to each word

max_length = 10 # only consider the first max_length words
results = np.zeros(shape=(len(samples),
                         max_length,
                         max(token_index.values())+1))

for i, sample in enumerate(samples):
    for j, word in list(enumerate(sample.split()))[:max_length]:
```

```
index = token_index.get(word)
results[i, j, index] = 1.
```

In []: token_index

```
Out[ ]: {'The': 1,
         'cat': 2,
         'sat': 3,
         'on': 4,
         'the': 5,
         'mat.': 6,
         'dog': 7,
         'ate': 8,
         'my': 9,
         'homework.': 10}
```

In []: results

```
Out[ ]: array([[[0., 1., 0., 0., 0., 0., 0., 0., 0., 0., 0., 0., 0., 0.],  
[0., 0., 1., 0., 0., 0., 0., 0., 0., 0., 0., 0., 0., 0.],  
[0., 0., 0., 1., 0., 0., 0., 0., 0., 0., 0., 0., 0., 0.],  
[0., 0., 0., 0., 1., 0., 0., 0., 0., 0., 0., 0., 0., 0.],  
[0., 0., 0., 0., 0., 1., 0., 0., 0., 0., 0., 0., 0., 0.],  
[0., 0., 0., 0., 0., 0., 1., 0., 0., 0., 0., 0., 0., 0.],  
[0., 0., 0., 0., 0., 0., 0., 1., 0., 0., 0., 0., 0., 0.],  
[0., 0., 0., 0., 0., 0., 0., 0., 1., 0., 0., 0., 0., 0.],  
[0., 0., 0., 0., 0., 0., 0., 0., 0., 1., 0., 0., 0., 0.],  
[0., 0., 0., 0., 0., 0., 0., 0., 0., 0., 1., 0., 0., 0.],  
[0., 0., 0., 0., 0., 0., 0., 0., 0., 0., 0., 1., 0., 0.],  
[0., 0., 0., 0., 0., 0., 0., 0., 0., 0., 0., 0., 1., 0.],  
[0., 0., 0., 0., 0., 0., 0., 0., 0., 0., 0., 0., 0., 1.],  
[0., 0., 0., 0., 0., 0., 0., 0., 0., 0., 0., 0., 0., 0.],  
[0., 0., 0., 0., 0., 0., 0., 0., 0., 0., 0., 0., 0., 0.],  
[0., 0., 0., 0., 0., 0., 0., 0., 0., 0., 0., 0., 0., 0.],  
[0., 0., 0., 0., 0., 0., 0., 0., 0., 0., 0., 0., 0., 0.],  
[0., 0., 0., 0., 0., 0., 0., 0., 0., 0., 0., 0., 0., 0.]]])
```

```
In [ ]: # Character-Level one-hot encoding
```

```
import string

samples = ['The cat sat on the mat.', 'The dog ate my homework.']
characters = string.printable
token_index = dict(zip(characters, range(1, len(characters) + 1)))

max_length = 50
results = np.zeros((len(samples), max_length, max(token_index.values())+1))
for i, sample in enumerate(samples):
    for j, character in enumerate(sample):
        index = token_index.get(character)
        results[i, j, index] = 1.
```

In []: token index

```
Out[ ]: {'0': 1,
         '1': 2,
         '2': 3,
         '3': 4,
         '4': 5,
         '5': 6,
         '6': 7,
         '7': 8,
         '8': 9,
         '9': 10,
         'a': 11,
         'b': 12,
         'c': 13,
         'd': 14,
         'e': 15,
         'f': 16,
         'g': 17,
         'h': 18,
         'i': 19,
         'j': 20,
         'k': 21,
         'l': 22,
         'm': 23,
         'n': 24,
         'o': 25,
         'p': 26,
         'q': 27,
         'r': 28,
         's': 29,
         't': 30,
         'u': 31,
         'v': 32,
         'w': 33,
         'x': 34,
         'y': 35,
         'z': 36,
         'A': 37,
         'B': 38,
         'C': 39,
         'D': 40,
         'E': 41,
         'F': 42,
         'G': 43,
         'H': 44,
         'I': 45,
         'J': 46,
         'K': 47,
         'L': 48,
         'M': 49,
         'N': 50,
         'O': 51,
         'P': 52,
         'Q': 53,
         'R': 54,
         'S': 55,
         'T': 56,
         'U': 57,
         'V': 58,
         'W': 59,
         'X': 60,
         'Y': 61,
```

'Z': 62,
'!': 63,
'''': 64,
'#': 65,
'\$': 66,
'%': 67,
'&': 68,
'''': 69,
'('': 70,
')': 71,
'*': 72,
'+': 73,
',': 74,
'-': 75,
'.': 76,
'/': 77,
::': 78,
';': 79,
'<': 80,
'='': 81,
'>': 82,
'?': 83,
('@': 84,
['': 85,
'\\': 86,
']': 87,
'^': 88,
'_': 89,
'`': 90,
'{': 91,
'|': 92,
'~': 94,
' ': 95,
'\t': 96,
'\n': 97,
'\r': 98,
'\x0b': 99,
'\x0c': 100}

In []: results.shape

Out[]: (2, 50, 101)

```
In [ ]: results[0,0,:]
```

```
In [ ]: np.where(results[0,0,:]==1)
```

Out[]: (array([56]),)

Word embeddings

- The use of dense *word vectors*, also called *word embeddings*

- One-hot encoding = binary, sparse, high-dimensional

 No description has been provided for this image

- How to obtain word embeddings?

- Learn word embeddings jointly with the main task
- Load the precomputed word embeddings (*pretrained word embeddings*)

- **Learning word embeddings**

- In embedding space, the geometric relationship (distance, direction, etc.) between word vectors should reflect the semantic relationships between these words.

 No description has been provided for this image

- Is there some ideal word-embedding space that would perfectly map human language and could be used for any natural-language-processing task?

- The `Embedding` layer in PyTorch

```
import torch
import torch.nn as nn
```

```
embedding_layer = nn.Embedding(num_embeddings=1000, embedding_dim=64)
```

- The `Embedding` layer takes as input a 2D tensor of integers `(samples, sequence_length)`, and returns a 3D floating-point tensor of shape `(samples, sequence_length, embedding_dimensionality)`.
- Each integer in the input corresponds to an index in the embedding matrix, and the layer looks up the corresponding learned vector during the forward pass.

In []:

```
import torch
import torch.nn as nn

embedding_layer = nn.Embedding(num_embeddings=1000, embedding_dim=64)

# Example input: batch of 2 samples, each with 4 token indices
input_indices = torch.tensor([[1, 2, 3, 4],
                             [4, 3, 2, 1]])

output = embedding_layer(input_indices)
print("Output shape:", output.shape)
```

Output shape: torch.Size([2, 4, 64])

- An embedding layer is mathematically equivalent to a linear layer without bias applied to one-hot vectors.

In []:

```
import torch
import torch.nn as nn

# Vocabulary of size 5, embedding dim 3
embedding = nn.Embedding(5, 3)
linear = nn.Linear(5, 3, bias=False)
```

```

# Copy weights so they're identical
linear.weight.data = embedding.weight.data.clone()

# Input: index 2 -> one-hot equivalent is [0,0,1,0,0]
index = torch.tensor([2])
one_hot = torch.nn.functional.one_hot(index, num_classes=5).float()

# Compare outputs
print("Embedding:", embedding(index))
print("Linear (one-hot):", one_hot @ linear.weight)

```

Embedding: tensor([-0.5881, -0.8058, 0.9569]), grad_fn=<EmbeddingBackward0>
Linear (one-hot): tensor([-0.5881, -0.8058, 0.9569]), grad_fn=<MMBackward0>

- Then, why use `Embedding` instead of `Linear`
 - Efficiency: One-hot encoding would create extremely large, mostly zero vectors. `Embedding` performs direct row indexing, which is O(1) and memory-efficient.
 - Interpretability: `Embedding` directly stores the token vectors and can be pretrained, visualized, or reused.
 - Simplicity: It's conceptually clear, "Each token has its own learnable vectors."

Revisited) Binary Classification with IMDB

In [108...]

```

import re, random
import torch
from collections import Counter
from datasets import load_dataset

ds = load_dataset("imdb")

id2label = {0: "neg", 1: "pos"}
train_list = [(id2label[int(r["label"])], r["text"]) for r in ds["train"]]
test_list = [(id2label[int(r["label"])], r["text"]) for r in ds["test"]]

```

In []:

```

# Simple tokenizer (English)
def simple_tokenize(s: str):
    # alphanumeric word tokens, lowercased
    return re.findall(r"\b\w+\b", s.lower())

# Build vocab from training set
# Keep top-N frequent tokens
MAX_VOCAB = 10000
specials = ["<unk>", "<pad>"]
counter = Counter()
for _, txt in train_list:
    counter.update(simple_tokenize(txt))

most_common = [w for w, _ in counter.most_common(MAX_VOCAB - len(specials))]
itos = specials + most_common
stoi = {w: i for i, w in enumerate(itos)}
UNK_IDX = stoi["<unk>"]
PAD_IDX = stoi["<pad>"]

```

In []:

```

label_to_int = {"neg": 0, "pos": 1}

# set MAX sequence length
MAX_LEN = 300

```

```

def text_pipeline(x: str):
    ids = [stoi.get(tok, UNK_IDX) for tok in simple_tokenize(x)]
    if len(ids) > MAX_LEN:
        ids = ids[:MAX_LEN]
    return torch.tensor(ids, dtype=torch.long)

def label_pipeline(y: str):
    return torch.tensor(label_to_int[y], dtype=torch.float32)

```

```

In [ ]: random.seed(42)
random.shuffle(train_list)
split_idx = int(len(train_list) * 0.9)
train_data = train_list[:split_idx]
valid_data = train_list[split_idx:]
test_data = test_list

```

```

In [ ]: print(f"Vocab size: {len(stoi)}, Train/Valid/Test: {len(train_data)}/{len(valid_data)}/{len(test_data)}")
Vocab size: 10000, Train/Valid/Test: 22500/2500/25000

```

```

In [ ]: import os, random, math, time
import torch
from torch import nn
from torch.utils.data import DataLoader
from torch.nn.utils.rnn import pad_sequence
import torch.optim as optim

SEED = 42
random.seed(SEED)
os.environ["PYTHONHASHSEED"] = str(SEED)
torch.manual_seed(SEED)
torch.cuda.manual_seed_all(SEED)

DEVICE = torch.device("cuda" if torch.cuda.is_available() else "cpu")
DEVICE

```

```
Out[ ]: device(type='cuda')
```

```

In [ ]: from torch.utils.data import DataLoader
from torch.nn.utils.rnn import pad_sequence

def collate_batch(batch):
    text_list, label_list = [], []
    for (label, text) in batch:
        text_list.append(text_pipeline(text))
        label_list.append(label_pipeline(label))
    # pad to the same length
    text_padded = pad_sequence(text_list, batch_first=True, padding_value=PAD_IDX)
    labels = torch.stack(label_list)
    return text_padded.to(DEVICE), labels.to(DEVICE)

BATCH_SIZE = 128
train_loader = DataLoader(train_data, batch_size=BATCH_SIZE, shuffle=True, collate_fn=collate_batch)
valid_loader = DataLoader(valid_data, batch_size=BATCH_SIZE, shuffle=False, collate_fn=collate_batch)
test_loader = DataLoader(test_data, batch_size=BATCH_SIZE, shuffle=False, collate_fn=collate_batch)
len(train_loader), len(valid_loader), len(test_loader)

```

```
Out[ ]: (176, 20, 196)
```

```
In [ ]: class SentimentMLP(nn.Module):
    def __init__(self, vocab_size, embed_dim=64, hidden_dims=(32, 32), pad_idx=None):
        super().__init__()
        self.embedding = nn.Embedding(vocab_size, embed_dim, padding_idx=pad_idx if pad_idx is not None else -1)
        layers = []
        in_dim = embed_dim
        for h in hidden_dims:
            layers += [nn.Linear(in_dim, h), nn.ReLU()]
            in_dim = h
        layers += [nn.Linear(in_dim, 1)] # output logit
        self.mlp = nn.Sequential(*layers)

    def forward(self, x):
        # x: (B, T)
        emb = self.embedding(x) # (B, T, E)
        # Create mask for non-pad tokens
        if hasattr(self.embedding, "padding_idx") and self.embedding.padding_idx is not None:
            pad_idx = self.embedding.padding_idx
        else:
            pad_idx = 0
        mask = (x != pad_idx).unsqueeze(-1) # (B, T, 1)
        emb = emb * mask # zero out pad embeddings
        lengths = mask.sum(dim=1).clamp(min=1) # (B, 1)
        mean_pooled = emb.sum(dim=1) / lengths # (B, E)
        logits = self.mlp(mean_pooled).squeeze(1) # (B,)
        return logits

model = SentimentMLP(vocab_size=len(stoi), embed_dim=64, hidden_dims=(32, 32), pad_idx=PAD_IDX).to(device)
```

```
Out[ ]: SentimentMLP(
  (embedding): Embedding(10000, 64, padding_idx=1)
  (mlp): Sequential(
    (0): Linear(in_features=64, out_features=32, bias=True)
    (1): ReLU()
    (2): Linear(in_features=32, out_features=32, bias=True)
    (3): ReLU()
    (4): Linear(in_features=32, out_features=1, bias=True)
  )
)
```

```
In [ ]: def binary_accuracy_from_logits(logits, targets):
    probs = torch.sigmoid(logits)
    preds = (probs >= 0.5).float()
    correct = (preds == targets).sum().item()
    return correct / targets.numel()

def run_epoch(dataloader, model, criterion, optimizer=None):
    is_train = optimizer is not None
    model.train() if is_train else model.eval()
    total_loss, total_acc, total_count = 0.0, 0.0, 0
    for xb, yb in dataloader:
        if is_train:
            optimizer.zero_grad(set_to_none=True)
        logits = model(xb)
        loss = criterion(logits, yb)
        if is_train:
            loss.backward()
            optimizer.step()
        bs = yb.size(0)
        total_loss += loss.item() * bs
```

```

        total_acc += binary_accuracy_from_logits(logits.detach(), yb) * bs
        total_count += bs
    return total_loss / total_count, total_acc / total_count

criterion = nn.BCEWithLogitsLoss()
optimizer = optim.Adam(model.parameters(), lr=1e-3)

EPOCHS = 20
history = {"train_loss": [], "train_acc": [], "valid_loss": [], "valid_acc": []}
best_val_acc, best_state = 0.0, None

for epoch in range(1, EPOCHS + 1):
    tr_loss, tr_acc = run_epoch(train_loader, model, criterion, optimizer)
    va_loss, va_acc = run_epoch(valid_loader, model, criterion, optimizer=None)

    history["train_loss"].append(tr_loss)
    history["train_acc"].append(tr_acc)
    history["valid_loss"].append(va_loss)
    history["valid_acc"].append(va_acc)

    print(f"Epoch {epoch:02d} | train_loss={tr_loss:.4f} acc={tr_acc:.4f} | valid_loss={va_loss:.4f} acc={va_acc:.4f}")
    if va_acc > best_val_acc:
        best_val_acc = va_acc
        best_state = {k: v.detach().cpu().clone() for k, v in model.state_dict().items()}

if best_state is not None:
    model.load_state_dict(best_state)

```

| | | | | |
|----------|-------------------|------------|-------------------|------------|
| Epoch 01 | train_loss=0.6262 | acc=0.6577 | valid_loss=0.4695 | acc=0.7880 |
| Epoch 02 | train_loss=0.3948 | acc=0.8227 | valid_loss=0.3757 | acc=0.8400 |
| Epoch 03 | train_loss=0.3111 | acc=0.8693 | valid_loss=0.3277 | acc=0.8636 |
| Epoch 04 | train_loss=0.2654 | acc=0.8932 | valid_loss=0.3172 | acc=0.8676 |
| Epoch 05 | train_loss=0.2344 | acc=0.9088 | valid_loss=0.3293 | acc=0.8672 |
| Epoch 06 | train_loss=0.2094 | acc=0.9199 | valid_loss=0.3058 | acc=0.8816 |
| Epoch 07 | train_loss=0.1877 | acc=0.9288 | valid_loss=0.3125 | acc=0.8844 |
| Epoch 08 | train_loss=0.1696 | acc=0.9385 | valid_loss=0.3205 | acc=0.8824 |
| Epoch 09 | train_loss=0.1520 | acc=0.9449 | valid_loss=0.3273 | acc=0.8828 |
| Epoch 10 | train_loss=0.1382 | acc=0.9524 | valid_loss=0.3384 | acc=0.8824 |
| Epoch 11 | train_loss=0.1224 | acc=0.9593 | valid_loss=0.3579 | acc=0.8820 |
| Epoch 12 | train_loss=0.1081 | acc=0.9662 | valid_loss=0.3786 | acc=0.8784 |
| Epoch 13 | train_loss=0.0968 | acc=0.9705 | valid_loss=0.3995 | acc=0.8752 |
| Epoch 14 | train_loss=0.0855 | acc=0.9766 | valid_loss=0.4289 | acc=0.8752 |
| Epoch 15 | train_loss=0.0742 | acc=0.9814 | valid_loss=0.4495 | acc=0.8728 |
| Epoch 16 | train_loss=0.0640 | acc=0.9864 | valid_loss=0.4805 | acc=0.8716 |
| Epoch 17 | train_loss=0.0554 | acc=0.9896 | valid_loss=0.5222 | acc=0.8692 |
| Epoch 18 | train_loss=0.0483 | acc=0.9917 | valid_loss=0.5369 | acc=0.8684 |
| Epoch 19 | train_loss=0.0416 | acc=0.9940 | valid_loss=0.5820 | acc=0.8672 |
| Epoch 20 | train_loss=0.0356 | acc=0.9949 | valid_loss=0.6148 | acc=0.8644 |

In []: `import matplotlib.pyplot as plt`


```

plt.figure()
plt.plot(history["train_loss"], label="train_loss")
plt.plot(history["valid_loss"], label="valid_loss")
plt.xlabel("epoch"); plt.ylabel("loss"); plt.legend(); plt.title("Loss Curves")
plt.show()


plt.figure()
plt.plot(history["train_acc"], label="train_acc")
plt.plot(history["valid_acc"], label="valid_acc")
plt.xlabel("epoch"); plt.ylabel("accuracy"); plt.legend(); plt.title("Accuracy Curves")
plt.show()

```

Loss Curves



Accuracy Curves



```
In [ ]: test_loss, test_acc = run_epoch(test_loader, model, criterion, optimizer=None)
print(f"TEST | loss={test_loss:.4f} acc={test_acc:.4f}")
```

TEST | loss=0.3433 acc=0.8610

Using pretrained word embeddings

- You don't have enough data available to learn truly powerful features on your own, but you expect the features that you need to be fairly generic.
- Word embedding models
 - Word2Vec (2013)
 - Global Vectors for Word Representation (GloVe) (2014)
 - fastText (2016)
 - BERT (2018)
- From raw text to word embeddings
 - Refer to the implementation above

```
In [ ]: # downloading the GloVe word embedding and preprocessing
# download url = https://drive.google.com/open?id=1NgMR-bnt02gYTr44BVTo2fRePastYypW

# mount Google Drive
from google.colab import drive
drive.mount('/content/gdrive')

EMBED_DIM = 100
glove_path = '/content/gdrive/My Drive/Lectures/deep-learning/datasets/glove.6B/glove.6B.100d.txt'

Drive already mounted at /content/gdrive; to attempt to forcibly remount, call drive.mount("/content/gdrive", force_remount=True).
```

```
In [ ]: import numpy as np

def load_glove_txt(glove_path, embedding_dim=100):
    """Read glove.6B.100d.txt into a dict: token -> np.array(embedding_dim)."""
    vectors = {}
    with open(glove_path, "r", encoding="utf-8") as f:
        for line in f:
            parts = line.rstrip().split(" ")
            token = parts[0]
            coefs = np.asarray(parts[1:], dtype=np.float32)
            if coefs.shape[0] != embedding_dim:
                continue # skip malformed rows
            vectors[token] = coefs
    return vectors

def build_embedding_matrix(stoi, glove_dict, embedding_dim=100, pad_idx=0, oov_init="normal"):
    """
    Build a (vocab_size x embedding_dim) embedding matrix aligned to 'stoi'.
    PAD row is zeros. OOV tokens use small random init.
    """
    vocab_size = len(stoi)
    emb = np.zeros((vocab_size, embedding_dim), dtype=np.float32)

    # OOV initialization
    if oov_init == "normal":
        def new_oov(): return np.random.normal(0, 0.05, size=(embedding_dim,)).astype(np.float32)
    elif oov_init == "uniform":
        def new_oov(): return np.random.uniform(-0.05, 0.05, size=(embedding_dim,)).astype(np.float32)
    else:
        def new_oov(): return np.zeros((embedding_dim,), dtype=np.float32)

    for token, idx in stoi.items():
        if idx == pad_idx:
            continue
        if token in glove_dict:
            emb[idx] = glove_dict[token]
        else:
            emb[idx] = new_oov()
```

```

    for token, idx in stoi.items():
        if idx == pad_idx:
            emb[idx] = np.zeros((embedding_dim,), dtype=np.float32)
            continue
        vec = glove_dict.get(token)
        if vec is None:
            vec = glove_dict.get(token.lower()) # just in case
        emb[idx] = vec if vec is not None else new_oov()

    return torch.tensor(emb) # (V, D)

```

```

In [ ]: # 1) Load GloVe dict
glove = load_glove_txt(glove_path, embedding_dim=EMBED_DIM)

# 2) Create embedding matrix aligned with 'stoi'
W = build_embedding_matrix(stoi, glove, embedding_dim=EMBED_DIM, pad_idx=PAD_IDX, oov_init="norm")

# 3) Build model with embed_dim=100 and load GloVe weights
model = SentimentMLP(vocab_size=len(stoi), embed_dim=EMBED_DIM, hidden_dims=(32, 32), pad_idx=PAD_IDX)

# Copy weights to model's embedding. Match device.
with torch.no_grad():
    model.embedding.weight.data.copy_(W) # copy on CPU first
    model.embedding.weight.data[PAD_IDX].zero_() # safety: make sure PAD row is 0

# (Option A) Freeze embeddings initially (common practice), then later unfreeze for small-lr fine-tuning
for p in model.embedding.parameters():
    p.requires_grad = False

# Move to device after weight init
model = model.to(DEVICE)
model

```

```

Out[ ]: SentimentMLP(
  (embedding): Embedding(10000, 100, padding_idx=1)
  (mlp): Sequential(
    (0): Linear(in_features=100, out_features=32, bias=True)
    (1): ReLU()
    (2): Linear(in_features=32, out_features=32, bias=True)
    (3): ReLU()
    (4): Linear(in_features=32, out_features=1, bias=True)
  )
)

```

```

In [ ]: criterion = nn.BCEWithLogitsLoss()
optimizer = optim.Adam(model.parameters(), lr=1e-3)

EPOCHS = 20
history = {"train_loss": [], "train_acc": [], "valid_loss": [], "valid_acc": []}
best_val_acc, best_state = 0.0, None

for epoch in range(1, EPOCHS + 1):
    tr_loss, tr_acc = run_epoch(train_loader, model, criterion, optimizer)
    va_loss, va_acc = run_epoch(valid_loader, model, criterion, optimizer=None)

    history["train_loss"].append(tr_loss)
    history["train_acc"].append(tr_acc)
    history["valid_loss"].append(va_loss)
    history["valid_acc"].append(va_acc)

    print(f"Epoch {epoch:02d} | train_loss={tr_loss:.4f} acc={tr_acc:.4f} | valid_loss={va_loss:.4f} acc={va_acc:.4f}")

```

```

if va_acc > best_val_acc:
    best_val_acc = va_acc
    best_state = {k: v.detach().cpu().clone() for k, v in model.state_dict().items()}

if best_state is not None:
    model.load_state_dict(best_state)

Epoch 01 | train_loss=0.6530 acc=0.6231 | valid_loss=0.5802 acc=0.7100
Epoch 02 | train_loss=0.5364 acc=0.7389 | valid_loss=0.5005 acc=0.7552
Epoch 03 | train_loss=0.4979 acc=0.7589 | valid_loss=0.4893 acc=0.7628
Epoch 04 | train_loss=0.4852 acc=0.7674 | valid_loss=0.4969 acc=0.7588
Epoch 05 | train_loss=0.4775 acc=0.7743 | valid_loss=0.4688 acc=0.7800
Epoch 06 | train_loss=0.4732 acc=0.7764 | valid_loss=0.4759 acc=0.7720
Epoch 07 | train_loss=0.4681 acc=0.7795 | valid_loss=0.4628 acc=0.7844
Epoch 08 | train_loss=0.4658 acc=0.7802 | valid_loss=0.4725 acc=0.7752
Epoch 09 | train_loss=0.4614 acc=0.7836 | valid_loss=0.4578 acc=0.7840
Epoch 10 | train_loss=0.4603 acc=0.7829 | valid_loss=0.4576 acc=0.7860
Epoch 11 | train_loss=0.4594 acc=0.7849 | valid_loss=0.4553 acc=0.7836
Epoch 12 | train_loss=0.4565 acc=0.7857 | valid_loss=0.4570 acc=0.7860
Epoch 13 | train_loss=0.4596 acc=0.7850 | valid_loss=0.4532 acc=0.7860
Epoch 14 | train_loss=0.4531 acc=0.7901 | valid_loss=0.4533 acc=0.7872
Epoch 15 | train_loss=0.4535 acc=0.7919 | valid_loss=0.4510 acc=0.7880
Epoch 16 | train_loss=0.4482 acc=0.7940 | valid_loss=0.4496 acc=0.7876
Epoch 17 | train_loss=0.4492 acc=0.7924 | valid_loss=0.4483 acc=0.7908
Epoch 18 | train_loss=0.4499 acc=0.7926 | valid_loss=0.4514 acc=0.7888
Epoch 19 | train_loss=0.4477 acc=0.7928 | valid_loss=0.4468 acc=0.7912
Epoch 20 | train_loss=0.4477 acc=0.7918 | valid_loss=0.4492 acc=0.7884

```

In []: `import matplotlib.pyplot as plt`


```

plt.figure()
plt.plot(history["train_loss"], label="train_loss")
plt.plot(history["valid_loss"], label="valid_loss")
plt.xlabel("epoch"); plt.ylabel("loss"); plt.legend(); plt.title("Loss Curves")
plt.show()


plt.figure()
plt.plot(history["train_acc"], label="train_acc")
plt.plot(history["valid_acc"], label="valid_acc")
plt.xlabel("epoch"); plt.ylabel("accuracy"); plt.legend(); plt.title("Accuracy Curves")
plt.show()

```

Loss Curves



Accuracy Curves



```
In [ ]: test_loss, test_acc = run_epoch(test_loader, model, criterion, optimizer=None)
print(f"TEST | loss={test_loss:.4f} acc={test_acc:.4f}")
```

TEST | loss=0.4489 acc=0.7913

- Finetuning the pretrained embeddings

```
In [ ]: # 1) Load GloVe dict
glove = load_glove_txt(glove_path, embedding_dim=EMBED_DIM)

# 2) Create embedding matrix aligned with 'stoi'
W = build_embedding_matrix(stoi, glove, embedding_dim=EMBED_DIM, pad_idx=PAD_IDX, oov_init="norm")

# 3) Build model with embed_dim=100 and Load GloVe weights
model = SentimentMLP(vocab_size=len(stoi), embed_dim=EMBED_DIM, hidden_dims=(32, 32), pad_idx=PAD_IDX)

# Copy weights to model's embedding. Match device.
with torch.no_grad():
    model.embedding.weight.data.copy_(W) # copy on CPU first
    model.embedding.weight.data[PAD_IDX].zero_() # safety: make sure PAD row is 0

# (Option A) Freeze embeddings initially (common practice), then later unfreeze for small-lr fine-tuning
# for p in model.embedding.parameters():
#     p.requires_grad = False

# Move to device after weight init
model = model.to(DEVICE)
model
```

```
Out[ ]: SentimentMLP(
    (embedding): Embedding(10000, 100, padding_idx=1)
    (mlp): Sequential(
        (0): Linear(in_features=100, out_features=32, bias=True)
        (1): ReLU()
        (2): Linear(in_features=32, out_features=32, bias=True)
        (3): ReLU()
        (4): Linear(in_features=32, out_features=1, bias=True)
    )
)
```

```
In [ ]: criterion = nn.BCEWithLogitsLoss()
optimizer = optim.Adam(model.parameters(), lr=1e-3)

EPOCHS = 20
history = {"train_loss": [], "train_acc": [], "valid_loss": [], "valid_acc": []}
best_val_acc, best_state = 0.0, None

for epoch in range(1, EPOCHS + 1):
    tr_loss, tr_acc = run_epoch(train_loader, model, criterion, optimizer)
    va_loss, va_acc = run_epoch(valid_loader, model, criterion, optimizer=None)

    history["train_loss"].append(tr_loss)
    history["train_acc"].append(tr_acc)
    history["valid_loss"].append(va_loss)
    history["valid_acc"].append(va_acc)

    print(f"Epoch {epoch:02d} | train_loss={tr_loss:.4f} acc={tr_acc:.4f} | valid_loss={va_loss:.4f} acc={va_acc:.4f}")
    if va_acc > best_val_acc:
        best_val_acc = va_acc
        best_state = {k: v.detach().cpu().clone() for k, v in model.state_dict().items()}


if best_state is not None:
    model.load_state_dict(best_state)
```

| | | | | |
|----------|-------------------|------------|-------------------|------------|
| Epoch 01 | train_loss=0.5995 | acc=0.6773 | valid_loss=0.4076 | acc=0.8256 |
| Epoch 02 | train_loss=0.3332 | acc=0.8591 | valid_loss=0.3076 | acc=0.8680 |
| Epoch 03 | train_loss=0.2557 | acc=0.8974 | valid_loss=0.2854 | acc=0.8776 |
| Epoch 04 | train_loss=0.2150 | acc=0.9162 | valid_loss=0.2826 | acc=0.8816 |
| Epoch 05 | train_loss=0.1851 | acc=0.9289 | valid_loss=0.3170 | acc=0.8720 |
| Epoch 06 | train_loss=0.1604 | acc=0.9400 | valid_loss=0.2983 | acc=0.8828 |
| Epoch 07 | train_loss=0.1385 | acc=0.9516 | valid_loss=0.3199 | acc=0.8776 |
| Epoch 08 | train_loss=0.1221 | acc=0.9571 | valid_loss=0.3572 | acc=0.8800 |
| Epoch 09 | train_loss=0.1053 | acc=0.9647 | valid_loss=0.3715 | acc=0.8764 |
| Epoch 10 | train_loss=0.0909 | acc=0.9720 | valid_loss=0.4148 | acc=0.8756 |
| Epoch 11 | train_loss=0.0779 | acc=0.9776 | valid_loss=0.4477 | acc=0.8608 |
| Epoch 12 | train_loss=0.0649 | acc=0.9822 | valid_loss=0.4969 | acc=0.8696 |
| Epoch 13 | train_loss=0.0522 | acc=0.9887 | valid_loss=0.5282 | acc=0.8620 |
| Epoch 14 | train_loss=0.0441 | acc=0.9907 | valid_loss=0.5920 | acc=0.8624 |
| Epoch 15 | train_loss=0.0356 | acc=0.9941 | valid_loss=0.6228 | acc=0.8560 |
| Epoch 16 | train_loss=0.0291 | acc=0.9961 | valid_loss=0.6689 | acc=0.8540 |
| Epoch 17 | train_loss=0.0253 | acc=0.9965 | valid_loss=0.7757 | acc=0.8576 |
| Epoch 18 | train_loss=0.0229 | acc=0.9968 | valid_loss=0.7894 | acc=0.8556 |
| Epoch 19 | train_loss=0.0176 | acc=0.9979 | valid_loss=0.8399 | acc=0.8532 |
| Epoch 20 | train_loss=0.0147 | acc=0.9984 | valid_loss=0.8651 | acc=0.8508 |


```
In [ ]: import matplotlib.pyplot as plt
```

```
plt.figure()
plt.plot(history["train_loss"], label="train_loss")
plt.plot(history["valid_loss"], label="valid_loss")
plt.xlabel("epoch"); plt.ylabel("loss"); plt.legend(); plt.title("Loss Curves")
plt.show()

plt.figure()
plt.plot(history["train_acc"], label="train_acc")
plt.plot(history["valid_acc"], label="valid_acc")
plt.xlabel("epoch"); plt.ylabel("accuracy"); plt.legend(); plt.title("Accuracy Curves")
plt.show()
```



Accuracy Curves



```
In [ ]: test_loss, test_acc = run_epoch(test_loader, model, criterion, optimizer=None)
print(f"TEST | loss={test_loss:.4f} acc={test_acc:.4f}")
```

```
TEST | loss=0.3290 acc=0.8700
```

6.2 Recurrent neural networks

- What we did in the IMDB example: an entire movie review was transformed into a single vector and processed in one go. --> *feedforward networks*
- As you are reading the present sentence, you are processing it word by word while keeping memories of what came before.
- A *recurrent neural network* (RNN) processes sequences by iterating through the sequence elements and maintaining a *state* containing information relative to what it has seen so far.
 - The network internally loops over sequence elements.

 No description has been provided for this image

```
# Pseudocode RNN
state_t = 0 # the state at t
for input_t in input_sequence: # iterates over sequence elements
    output_t = f(input_t, state_t)
    state_t = output_t # the previous output becomes the state for the next iteration

# More detailed pseudocode for the RNN
state_t = 0
for input_t in input_sequence:
```

```
    output_t = activation(dot(W, input_t) + dot(U, state_t) + b)
    state_t = output_t
```

```
In [ ]: # Numpy implementation of a simple RNN
import numpy as np

timesteps = 100
input_features = 32
output_features = 64

inputs = np.random.random((timesteps, input_features)) # input data

state_t = np.zeros((output_features,))

W = np.random.random((output_features, input_features))
U = np.random.random((output_features, output_features))
b = np.random.random((output_features,))

successive_outputs = []
for input_t in inputs:
    output_t = np.tanh(np.dot(W, input_t) + np.dot(U, state_t) + b)
    successive_outputs.append(output_t)
    state_t = output_t

final_output_sequence = np.concatenate(successive_outputs, axis=0)
```



A recurrent layer in PyTorch

```
import torch
import torch.nn as nn

rnn = nn.RNN(
    input_size=32,      # number of input features per timestep
    hidden_size=64,     # size of the hidden state (output features)
    batch_first=True    # input and output tensors have shape (batch, time, features)
)
```

- `RNN` processes batches of sequences.
 - Inputs `(batch_size, timesteps, input_features)`
 - The forward pass returns two tensors:
`output, hidden = rnn(x)`
 - `output` : contains the hidden state for each timestep
 - `shape = (batch_size, timesteps, hidden_size)`
 - `hidden` : contains the hidden state for the last timestep
 - `shape = (num_layers, batch_size, hidden_size)`

```
In [ ]: import torch
import torch.nn as nn

rnn = nn.RNN(
    input_size=32,      # number of input features per timestep
    hidden_size=64,     # size of the hidden state (output features)
    batch_first=True    # input and output tensors have shape (batch, time, features)
)
```

```

x = torch.randn(8, 10, 32) # (batch, timesteps, input_features)
output, hidden = rnn(x)

full_sequence = output          # all timesteps
last_output   = hidden[-1]       # Last timestep only
print(full_sequence.shape) # (8, 10, 64)
print(last_output.shape) # (8, 64)

torch.Size([8, 10, 64])
torch.Size([8, 64])

```

In []:

```

import torch
import torch.nn as nn

class SentimentRNN(nn.Module):
    def __init__(self, vocab_size, embed_dim=64, hidden_size=64, num_layers=1, pad_idx=None):
        super().__init__()
        # Embedding Layer
        self.embedding = nn.Embedding(
            vocab_size,
            embed_dim,
            padding_idx=pad_idx if pad_idx is not None else 0
        )

        # Simple RNN Layer
        self.rnn = nn.RNN(
            input_size=embed_dim,
            hidden_size=hidden_size,
            num_layers=num_layers,
            batch_first=True
        )

        # Output Layer
        self.fc = nn.Linear(hidden_size, 1) # binary classification (logit output)

    def forward(self, x):
        # x: (B, T)
        emb = self.embedding(x) # (B, T, E)

        # Handle padding (mask and sequence lengths)
        if hasattr(self.embedding, "padding_idx") and self.embedding.padding_idx is not None:
            pad_idx = self.embedding.padding_idx
            lengths = (x != pad_idx).sum(dim=1).cpu() # (B,)
            packed = nn.utils.rnn.pack_padded_sequence(
                emb, lengths, batch_first=True, enforce_sorted=False
            )
            packed_out, hidden = self.rnn(packed)
            # hidden: (num_layers, B, hidden_size)
            last_hidden = hidden[-1] # last layer's hidden state
        else:
            output, hidden = self.rnn(emb)
            last_hidden = hidden[-1] # (B, hidden_size)

        logits = self.fc(last_hidden).squeeze(1) # (B, )
        return logits

```

In []:

```

!pip install torchinfo -q
from torchinfo import summary

model = SentimentRNN(vocab_size=10000).cpu()

```

```
summary(model, input_data=torch.randint(0, 10000, (1, 20)).long())
```

Out[]:

```
=====
Layer (type:depth-idx)           Output Shape      Param #
=====
SentimentRNN                   [1]                  --
|Embedding: 1-1                [1, 20, 64]        640,000
|RNN: 1-2                      [20, 64]            8,320
|Linear: 1-3                   [1, 1]              65
=====
Total params: 648,385
Trainable params: 648,385
Non-trainable params: 0
Total mult-adds (Units.MEGABYTES): 11.29
=====
Input size (MB): 0.00
Forward/backward pass size (MB): 0.02
Params size (MB): 2.59
Estimated Total Size (MB): 2.61
=====
```

- Handling variable-length sequences with `nn.utils.rnn.pack_padded_sequence`
 - When working with text or sequential data, different samples often have different sequence lengths.
 - To form a batch, all sequences must be padded to the same length (e.g., using `<pad>` tokens).
 - However, RNNs will still process those padding tokens, which adds unnecessary computation and noise.
 - Ex.
 - "I love it" -> [12, 45, 88, 0]
 - "This movie is great" -> [13, 56, 44, 99]
 - "Bad" -> [77, 0, 0, 0]
 - The RNN will compute over all timesteps (including `<pad>`), even though most of those zeros carry no meaning.
 - `pack_padded_sequence`
 - PyTorch provides a utility that allows RNNs to ignore padded timesteps efficiently:
`from torch.nn.utils.rnn import pack_padded_sequence`
 - ```
packed = pack_padded_sequence(
 embedded_batch, lengths, batch_first=True, enforce_sorted=False
)
output, hidden = rnn(packed)
 ◦ embedded_batch : (batch_size, seq_len, embed_dim)
 ◦ lengths : a tensor or list containing the true lengths of each sequence
 ◦ batch_first=True : ensures batch dimension comes first ((B, T, E))
 ◦ enforce_sorted=False : lets you use unsorted batches safely
 ◦ Once packed, the RNN will automatically skip over padding positions, processing only the valid tokens in each sequence.
```

```
In []: import torch
from torch.nn.utils.rnn import pad_sequence, pack_padded_sequence, pad_packed_sequence

Example: three variable-length sequences (token IDs)
seqs = [
 torch.tensor([1, 2, 3, 4]), # Length = 4
 torch.tensor([5, 6, 7]), # Length = 3
 torch.tensor([8, 9]) # Length = 2
]

Pad them to the same length (right-padding with 0)
padded = pad_sequence(seqs, batch_first=True, padding_value=0)
print("Padded batch (shape =", padded.shape, "):\n", padded)

Compute true lengths
lengths = torch.tensor([len(s) for s in seqs])
print("\nSequence lengths:", lengths.tolist())

Pack the padded batch
packed = pack_padded_sequence(padded, lengths, batch_first=True, enforce_sorted=False)
print("\nPacked data representation:\n", packed.data)
print("Packed batch_sizes:", packed.batch_sizes)

Unpack (restore to padded form)
unpacked, unpacked_lengths = pad_packed_sequence(packed, batch_first=True)
print("\nUnpacked (back to padded):\n", unpacked)
```

```
Padded batch (shape = torch.Size([3, 4])):
tensor([[1, 2, 3, 4],
 [5, 6, 7, 0],
 [8, 9, 0, 0]])
```

```
Sequence lengths: [4, 3, 2]
```

```
Packed data representation:
tensor([1, 5, 8, 2, 6, 9, 3, 7, 4])
Packed batch_sizes: tensor([3, 3, 2, 1])
```

```
Unpacked (back to padded):
tensor([[1, 2, 3, 4],
 [5, 6, 7, 0],
 [8, 9, 0, 0]])
```

- Stacking several recurrent layers

```
In []: !pip install torchinfo -q
from torchinfo import summary

model = SentimentRNN(vocab_size=10000, num_layers=2).cpu()

summary(model, input_data=torch.randint(0, 10000, (1, 20)).long())
```

```

Out[]: =====
Layer (type:depth-idx) Output Shape Param #
=====
SentimentRNN [1] --
|---Embedding: 1-1 [1, 20, 64] 640,000
|---RNN: 1-2 [20, 64] 16,640
|---Linear: 1-3 [1, 1] 65
=====
Total params: 656,705
Trainable params: 656,705
Non-trainable params: 0
Total mult-adds (Units.MEGABYTES): 21.94
=====
Input size (MB): 0.00
Forward/backward pass size (MB): 0.02
Params size (MB): 2.63
Estimated Total Size (MB): 2.65
=====
```

- Revisit the IMDB movie review classification problem

```

In []: model = SentimentRNN(
 vocab_size=len(stoi),
 embed_dim=64,
 hidden_size=32,
 num_layers=1,
 pad_idx=PAD_IDX
).to(DEVICE)

criterion = nn.BCEWithLogitsLoss()
optimizer = optim.Adam(model.parameters(), lr=1e-3)

EPOCHS = 20
history = {"train_loss": [], "train_acc": [], "valid_loss": [], "valid_acc": []}
best_val_acc, best_state = 0.0, None

for epoch in range(1, EPOCHS + 1):
 tr_loss, tr_acc = run_epoch(train_loader, model, criterion, optimizer)
 va_loss, va_acc = run_epoch(valid_loader, model, criterion, optimizer=None)

 history["train_loss"].append(tr_loss)
 history["train_acc"].append(tr_acc)
 history["valid_loss"].append(va_loss)
 history["valid_acc"].append(va_acc)

 print(f"Epoch {epoch:02d} | train_loss={tr_loss:.4f} acc={tr_acc:.4f} | valid_loss={va_loss:.4f} acc={va_acc:.4f}")
 if va_acc > best_val_acc:
 best_val_acc = va_acc
 best_state = {k: v.detach().cpu().clone() for k, v in model.state_dict().items()}


if best_state is not None:
 model.load_state_dict(best_state)
```

|          |                   |            |                   |            |
|----------|-------------------|------------|-------------------|------------|
| Epoch 01 | train_loss=0.6890 | acc=0.5364 | valid_loss=0.6836 | acc=0.5520 |
| Epoch 02 | train_loss=0.6539 | acc=0.6098 | valid_loss=0.6582 | acc=0.6120 |
| Epoch 03 | train_loss=0.5898 | acc=0.6946 | valid_loss=0.5918 | acc=0.7040 |
| Epoch 04 | train_loss=0.5781 | acc=0.7079 | valid_loss=0.6152 | acc=0.6660 |
| Epoch 05 | train_loss=0.5296 | acc=0.7472 | valid_loss=0.6037 | acc=0.6888 |
| Epoch 06 | train_loss=0.5080 | acc=0.7661 | valid_loss=0.5592 | acc=0.7328 |
| Epoch 07 | train_loss=0.4903 | acc=0.7755 | valid_loss=0.5809 | acc=0.7276 |
| Epoch 08 | train_loss=0.5262 | acc=0.7454 | valid_loss=0.5755 | acc=0.7172 |
| Epoch 09 | train_loss=0.5365 | acc=0.7251 | valid_loss=0.6352 | acc=0.6524 |
| Epoch 10 | train_loss=0.5160 | acc=0.7444 | valid_loss=0.6072 | acc=0.6968 |
| Epoch 11 | train_loss=0.4510 | acc=0.7958 | valid_loss=0.5641 | acc=0.7464 |
| Epoch 12 | train_loss=0.4197 | acc=0.8174 | valid_loss=0.5636 | acc=0.7444 |
| Epoch 13 | train_loss=0.3842 | acc=0.8406 | valid_loss=0.5530 | acc=0.7604 |
| Epoch 14 | train_loss=0.3667 | acc=0.8510 | valid_loss=0.5637 | acc=0.7512 |
| Epoch 15 | train_loss=0.3448 | acc=0.8616 | valid_loss=0.5583 | acc=0.7588 |
| Epoch 16 | train_loss=0.3256 | acc=0.8735 | valid_loss=0.5655 | acc=0.7592 |
| Epoch 17 | train_loss=0.3190 | acc=0.8752 | valid_loss=0.5808 | acc=0.7576 |
| Epoch 18 | train_loss=0.3111 | acc=0.8797 | valid_loss=0.5740 | acc=0.7636 |
| Epoch 19 | train_loss=0.3989 | acc=0.8183 | valid_loss=0.6351 | acc=0.6776 |
| Epoch 20 | train_loss=0.4470 | acc=0.7952 | valid_loss=0.6240 | acc=0.6532 |


```
In []: import matplotlib.pyplot as plt
```

```
plt.figure()
plt.plot(history["train_loss"], label="train_loss")
plt.plot(history["valid_loss"], label="valid_loss")
plt.xlabel("epoch"); plt.ylabel("loss"); plt.legend(); plt.title("Loss Curves")
plt.show()

plt.figure()
plt.plot(history["train_acc"], label="train_acc")
plt.plot(history["valid_acc"], label="valid_acc")
plt.xlabel("epoch"); plt.ylabel("accuracy"); plt.legend(); plt.title("Accuracy Curves")
plt.show()
```



### Accuracy Curves



## Applications of RNNs

- | | No description has been provided for this image
- | | No description has been provided for this image
- | | No description has been provided for this image
- | | No description has been provided for this image

## Flexibility of RNNs

- | | No description has been provided for this image

## Understanding the LSTM and GRU layers

- There are two other popular recurrent layer: `LSTM` and `GRU`.
- `SimpleRNN` is difficult to learn long-term dependencies.
  - This is due to the *vanishing gradient problem*.
- Long Short-Term Memory (LSTM)
  - proposed by Hochreiter and Schmidhuber in 1997
  - It adds a way to carry information across many timesteps.

- **SimpleRNN**
  -
- **SimpleRNN + additional data flow (carry)**
  - An additional data flow that carries information across timesteps
    - $C_t$  where  $C$  stands for *carry*
  - It will be combined with the input connections and the recurrent connection via a dense transformation.
  - Then, it will affect the state being sent to the next timestep.
- **LSTM**
  - How is the next value of the carry dataflow computed?
    - It involves three distinct transformations, which all have the form of the following:  
 $y = \text{activation}(\text{dot}(\text{state}_t, U) + \text{dot}(\text{input}_t, W) + b)$
    - All three transformations have their own weight matrices.  
 $\text{output}_t = \text{activation}(\text{dot}(\text{state}_t, U_o) + \text{dot}(\text{input}_t, W_o) + \text{dot}(C_t, V_o) + b_o)$
    - We obtain the new carry state ( $c_{t+1}$ ).  
 $c_{t+1} = i_t * k_t + c_t * f_t$
  - Details of **LSTM**
    - **LSTM** diagram
      -
    - Computations involved in **LSTM**
      -
- **GRU (Gated Recurrent Unit)**
  - **GRU** diagram
    -
  - Computations involved in **GRU**
    -

## LSTM example in PyTorch

In [ ]:

```
import torch
import torch.nn as nn

class SentimentLSTM(nn.Module):
 def __init__(self, vocab_size, embed_dim=64, hidden_size=64, num_layers=1, pad_idx=None):
 super().__init__()
 # Embedding layer
 self.embedding = nn.Embedding(
 vocab_size,
 embed_dim,
 padding_idx=pad_idx if pad_idx is not None else 0
)

 # LSTM layer
 self.lstm = nn.LSTM(
 input_size=embed_dim,
 hidden_size=hidden_size,
 num_layers=num_layers,
 batch_first=True
)

 # Output layer
 self.fc = nn.Linear(hidden_size, 1) # binary classification (logit output)

 def forward(self, x):
 # x: (B, T)
 emb = self.embedding(x) # (B, T, E)

 # Handle padding (mask and sequence lengths)
 if hasattr(self.embedding, "padding_idx") and self.embedding.padding_idx is not None:
 pad_idx = self.embedding.padding_idx
 lengths = (x != pad_idx).sum(dim=1).cpu() # (B,)
 packed = nn.utils.rnn.pack_padded_sequence(
 emb, lengths, batch_first=True, enforce_sorted=False
)
 packed_out, (hidden, cell) = self.lstm(packed)
 # hidden: (num_layers, B, hidden_size)
 last_hidden = hidden[-1] # Last Layer's hidden state
 else:
 output, hidden = self.lstm(emb)
 last_hidden = hidden[-1] # (B, hidden_size)

 logits = self.fc(last_hidden).squeeze(1) # (B,)
 return logits
```

In [ ]:

```
model = SentimentLSTM(
 vocab_size=len(stoi),
 embed_dim=64,
 hidden_size=32,
 num_layers=1,
 pad_idx=PAD_IDX
).to(DEVICE)

criterion = nn.BCEWithLogitsLoss()
optimizer = optim.Adam(model.parameters(), lr=1e-3)

EPOCHS = 20
history = {"train_loss": [], "train_acc": [], "valid_loss": [], "valid_acc": []}
```

```

best_val_acc, best_state = 0.0, None

for epoch in range(1, EPOCHS + 1):
 tr_loss, tr_acc = run_epoch(train_loader, model, criterion, optimizer)
 va_loss, va_acc = run_epoch(valid_loader, model, criterion, optimizer=None)

 history["train_loss"].append(tr_loss)
 history["train_acc"].append(tr_acc)
 history["valid_loss"].append(va_loss)
 history["valid_acc"].append(va_acc)

 print(f"Epoch {epoch:02d} | train_loss={tr_loss:.4f} acc={tr_acc:.4f} | valid_loss={va_loss:.4f} acc={va_acc:.4f}")
 if va_acc > best_val_acc:
 best_val_acc = va_acc
 best_state = {k: v.detach().cpu().clone() for k, v in model.state_dict().items()}

if best_state is not None:
 model.load_state_dict(best_state)

```

| Epoch | train_loss | acc    | valid_loss | acc    |
|-------|------------|--------|------------|--------|
| 01    | 0.6766     | 0.5595 | 0.6332     | 0.6412 |
| 02    | 0.6217     | 0.6840 | 0.6209     | 0.6672 |
| 03    | 0.6479     | 0.6232 | 0.6429     | 0.6380 |
| 04    | 0.5682     | 0.7108 | 0.5201     | 0.7396 |
| 05    | 0.4917     | 0.7706 | 0.5129     | 0.7552 |
| 06    | 0.5017     | 0.7592 | 0.4681     | 0.7964 |
| 07    | 0.4641     | 0.7834 | 0.4919     | 0.7828 |
| 08    | 0.4540     | 0.7896 | 0.4468     | 0.8052 |
| 09    | 0.4283     | 0.8141 | 0.4423     | 0.8060 |
| 10    | 0.3534     | 0.8544 | 0.4261     | 0.8096 |
| 11    | 0.3267     | 0.8695 | 0.4059     | 0.8320 |
| 12    | 0.2984     | 0.8837 | 0.4137     | 0.8268 |
| 13    | 0.2879     | 0.8904 | 0.4037     | 0.8296 |
| 14    | 0.2768     | 0.8947 | 0.5504     | 0.7836 |
| 15    | 0.3201     | 0.8686 | 0.4482     | 0.8112 |
| 16    | 0.2597     | 0.9013 | 0.4590     | 0.8104 |
| 17    | 0.2457     | 0.9081 | 0.4090     | 0.8428 |
| 18    | 0.2438     | 0.9069 | 0.4154     | 0.8420 |
| 19    | 0.2056     | 0.9264 | 0.4127     | 0.8388 |
| 20    | 0.1937     | 0.9315 | 0.4244     | 0.8440 |



In [ ]: `import matplotlib.pyplot as plt`

```

plt.figure()
plt.plot(history["train_loss"], label="train_loss")
plt.plot(history["valid_loss"], label="valid_loss")
plt.xlabel("epoch"); plt.ylabel("loss"); plt.legend(); plt.title("Loss Curves")
plt.show()

plt.figure()
plt.plot(history["train_acc"], label="train_acc")
plt.plot(history["valid_acc"], label="valid_acc")
plt.xlabel("epoch"); plt.ylabel("accuracy"); plt.legend(); plt.title("Accuracy Curves")
plt.show()


```



- **Exercise:** implement RNN or LSTM + mean pooling over timesteps + Linear

## Bidirectional RNNs

- A standard RNN (or an LSTM/GRU) processes a sequence chronologically, from the first token to the last. This is called a unidirectional model.
- The hidden state at any time step  $t(h_t)$  only contains information about the inputs from the past ( $x_1, \dots, x_t$ ). It has no knowledge of what is coming next in the sequence.
- Consider "The apple pie was delicious."
  - To correctly understand that "apple" refers to a food item (and not the company), it is helpful to know that the next word is "pie". A standard RNN making a decision at the word "apple" does not have this future context.
- A Bidirectional RNN solves this problem by processing the sequences in two directions at once:
  - A Forward RNN: from left to right
  - A Backward RNN: from right to left



In [109...]

```

import torch
import torch.nn as nn

class BiLSTMSentiment(nn.Module):
 def __init__(self, vocab_size, embed_dim=100, hidden_size=128, num_layers=1, pad_idx=0, dropo
 super().__init__()
 self.embedding = nn.Embedding(vocab_size, embed_dim, padding_idx=pad_idx)
 self.lstm = nn.LSTM(
 input_size=embed_dim,
 hidden_size=hidden_size,
 num_layers=num_layers,
 batch_first=True,
 bidirectional=True, # <- bidirectional LSTM
 dropout=0.0 if num_layers == 1 else dropout
)
 # Concatenate forward/backward Last hidden → 2*hidden_size
 self.head = nn.Sequential(
 nn.Dropout(dropout),
 nn.Linear(2 * hidden_size, 1) # binary Logit
)

 def forward(self, x):
 # x: (B, T) with PAD indices
 emb = self.embedding(x) # (B, T, E)

```

```

Pack to ignore PAD steps in the LSTM
pad_idx = self.embedding.padding_idx if self.embedding.padding_idx is not None else 0
lengths = (x != pad_idx).sum(dim=1).cpu() # (B,)
packed = nn.utils.rnn.pack_padded_sequence(emb, lengths, batch_first=True, enforce_sorted=False)

packed_out, (h_n, c_n) = self.lstm(packed)
h_n: (num_layers*2, B, H). Take last layer's forward/backward: [-2], [-1]
last_fwd = h_n[-2] # (B, H)
last_bwd = h_n[-1] # (B, H)
feat = torch.cat([last_fwd, last_bwd], dim=1) # (B, 2H)

logits = self.head(feat).squeeze(1) # (B,)
return logits

```

In [111...]

```

model = BiLSTMSentiment(
 vocab_size=len(stoi),
 embed_dim=64,
 hidden_size=32,
 num_layers=1,
 pad_idx=PAD_IDX,
 dropout=0.2
).to(DEVICE)

print(model)

criterion = nn.BCEWithLogitsLoss()
optimizer = optim.Adam(model.parameters(), lr=1e-3)

EPOCHS = 20
history = {"train_loss": [], "train_acc": [], "valid_loss": [], "valid_acc": []}
best_val_acc, best_state = 0.0, None

for epoch in range(1, EPOCHS + 1):
 tr_loss, tr_acc = run_epoch(train_loader, model, criterion, optimizer)
 va_loss, va_acc = run_epoch(valid_loader, model, criterion, optimizer=None)

 history["train_loss"].append(tr_loss)
 history["train_acc"].append(tr_acc)
 history["valid_loss"].append(va_loss)
 history["valid_acc"].append(va_acc)

 print(f"Epoch {epoch:02d} | train_loss={tr_loss:.4f} acc={tr_acc:.4f} | valid_loss={va_loss:.4f} acc={va_acc:.4f}")
 if va_acc > best_val_acc:
 best_val_acc = va_acc
 best_state = {k: v.detach().cpu().clone() for k, v in model.state_dict().items()}

if best_state is not None:
 model.load_state_dict(best_state)

```

```
BiLSTMSentiment(
 (embedding): Embedding(10000, 64, padding_idx=1)
 (lstm): LSTM(64, 32, batch_first=True, bidirectional=True)
 (head): Sequential(
 (0): Dropout(p=0.2, inplace=False)
 (1): Linear(in_features=64, out_features=1, bias=True)
)
)
Epoch 01 | train_loss=0.6784 acc=0.5624 | valid_loss=0.6283 acc=0.6516
Epoch 02 | train_loss=0.5455 acc=0.7248 | valid_loss=0.4817 acc=0.7752
Epoch 03 | train_loss=0.4203 acc=0.8113 | valid_loss=0.3994 acc=0.8224
Epoch 04 | train_loss=0.3604 acc=0.8466 | valid_loss=0.3934 acc=0.8324
Epoch 05 | train_loss=0.2980 acc=0.8791 | valid_loss=0.3613 acc=0.8468
Epoch 06 | train_loss=0.2554 acc=0.8995 | valid_loss=0.3454 acc=0.8576
Epoch 07 | train_loss=0.2254 acc=0.9129 | valid_loss=0.3518 acc=0.8608
Epoch 08 | train_loss=0.1910 acc=0.9275 | valid_loss=0.3612 acc=0.8632
Epoch 09 | train_loss=0.1662 acc=0.9377 | valid_loss=0.3526 acc=0.8508
Epoch 10 | train_loss=0.1444 acc=0.9472 | valid_loss=0.3575 acc=0.8560
Epoch 11 | train_loss=0.1189 acc=0.9580 | valid_loss=0.3841 acc=0.8584
Epoch 12 | train_loss=0.1060 acc=0.9623 | valid_loss=0.3926 acc=0.8584
Epoch 13 | train_loss=0.0850 acc=0.9714 | valid_loss=0.4128 acc=0.8664
Epoch 14 | train_loss=0.1269 acc=0.9538 | valid_loss=0.4401 acc=0.8472
Epoch 15 | train_loss=0.0778 acc=0.9752 | valid_loss=0.4826 acc=0.8644
Epoch 16 | train_loss=0.0540 acc=0.9844 | valid_loss=0.4786 acc=0.8596
Epoch 17 | train_loss=0.0497 acc=0.9853 | valid_loss=0.4757 acc=0.8556
Epoch 18 | train_loss=0.0401 acc=0.9891 | valid_loss=0.5452 acc=0.8624
Epoch 19 | train_loss=0.0282 acc=0.9938 | valid_loss=0.5906 acc=0.8616
Epoch 20 | train_loss=0.0490 acc=0.9838 | valid_loss=0.5361 acc=0.8568
```


In [112...]

```
import matplotlib.pyplot as plt


plt.figure()
plt.plot(history["train_loss"], label="train_loss")
plt.plot(history["valid_loss"], label="valid_loss")
plt.xlabel("epoch"); plt.ylabel("loss"); plt.legend(); plt.title("Loss Curves")
plt.show()

plt.figure()
plt.plot(history["train_acc"], label="train_acc")
plt.plot(history["valid_acc"], label="valid_acc")
plt.xlabel("epoch"); plt.ylabel("accuracy"); plt.legend(); plt.title("Accuracy Curves")
plt.show()
```

Loss Curves



Accuracy Curves



## 6.4 Sequence processing with convnets

- For sequence processing, time can be treated as a spatial dimension like the height or width of a 2D image.

- 1D convnets can be competitive with RNNs on certain sequence processing tasks.

- It has cheaper computational cost.

## Understanding 1D convolution for sequence data

- We can use 1D convolutions, extracting local 1D patches (subsequences) from sequences.

|  No description has been provided for this image

- Such 1D convolution layers can recognize local patterns in a sequence.

## 1D pooling for sequence data

- Similar to 2D pooling operations: 2D average/max pooling
- It extracts 1D patches (subsequences) from an input and outputting the maximum value (max pooling) or average value (average pooling).
- It is used for reducing the length of 1D inputs (subsampling).

## Implementing a 1D convnet

In [ ]:

```
import torch
import torch.nn as nn
import torch.optim as optim

--- Conv1D Sentiment Model ---
class SentimentCNN(nn.Module):
 def __init__(self, vocab_size, embed_dim=128, pad_idx=0):
 super().__init__()
 self.embedding = nn.Embedding(vocab_size, embed_dim, padding_idx=pad_idx)

 self.conv_block = nn.Sequential(
 nn.Conv1d(in_channels=embed_dim, out_channels=32, kernel_size=7),
 nn.ReLU(),
 nn.MaxPool1d(kernel_size=5),
 nn.Conv1d(in_channels=32, out_channels=32, kernel_size=7),
 nn.ReLU()
)

 # Global max pooling over time
 self.global_max_pool = nn.AdaptiveMaxPool1d(1)

 # Fully connected output layer
 self.fc = nn.Linear(32, 1)

 def forward(self, x):
 # x: (B, T)
 emb = self.embedding(x) # (B, T, E)
 emb = emb.transpose(1, 2) # (B, E, T) → required by Conv1d

 features = self.conv_block(emb) # (B, 32, L)
```

```

 pooled = self.global_max_pool(features).squeeze(-1) # (B, 32)
 logits = self.fc(pooled).squeeze(1) # (B,)
 return logits

--- Model initialization ---
model = SentimentCNN(vocab_size=len(stoi), embed_dim=128, pad_idx=PAD_IDX).to(DEVICE)
print(model)

```

```

SentimentCNN(
 (embedding): Embedding(10000, 128, padding_idx=1)
 (conv_block): Sequential(
 (0): Conv1d(128, 32, kernel_size=(7,), stride=(1,))
 (1): ReLU()
 (2): MaxPool1d(kernel_size=5, stride=5, padding=0, dilation=1, ceil_mode=False)
 (3): Conv1d(32, 32, kernel_size=(7,), stride=(1,))
 (4): ReLU()
)
 (global_max_pool): AdaptiveMaxPool1d(output_size=1)
 (fc): Linear(in_features=32, out_features=1, bias=True)
)

```

```

In []: # criterion = nn.BCEWithLogitsLoss()
optimizer = optim.RMSprop(model.parameters(), lr=1e-4)

criterion = nn.BCEWithLogitsLoss()
optimizer = optim.Adam(model.parameters(), lr=1e-3)

EPOCHS = 10
history = {"train_loss": [], "train_acc": [], "valid_loss": [], "valid_acc": []}
best_val_acc, best_state = 0.0, None

for epoch in range(1, EPOCHS + 1):
 tr_loss, tr_acc = run_epoch(train_loader, model, criterion, optimizer)
 va_loss, va_acc = run_epoch(valid_loader, model, criterion, optimizer=None)

 history["train_loss"].append(tr_loss)
 history["train_acc"].append(tr_acc)
 history["valid_loss"].append(va_loss)
 history["valid_acc"].append(va_acc)

 print(f"[CNN] Epoch {epoch:02d} | "
 f"train_loss={tr_loss:.4f} acc={tr_acc:.4f} | "
 f"valid_loss={va_loss:.4f} acc={va_acc:.4f}")

 if va_acc > best_val_acc:
 best_val_acc = va_acc
 best_state = {k: v.detach().cpu().clone() for k, v in model.state_dict().items()}

Load best weights
if best_state is not None:
 model.load_state_dict(best_state)
 print(f"Loaded best CNN model (val_acc={best_val_acc:.4f}).")

Evaluate on test set
test_loss, test_acc = run_epoch(test_loader, model, criterion, optimizer=None)
print("=" * 60)
print(f"[CNN] Test Loss: {test_loss:.4f} | Test Accuracy: {test_acc:.4f}")
print("=" * 60)

```


```
[CNN] Epoch 01 | train_loss=0.6301 acc=0.6212 | valid_loss=0.4997 acc=0.7588
[CNN] Epoch 02 | train_loss=0.4041 acc=0.8181 | valid_loss=0.3944 acc=0.8244
[CNN] Epoch 03 | train_loss=0.2785 acc=0.8861 | valid_loss=0.3615 acc=0.8372
[CNN] Epoch 04 | train_loss=0.1908 acc=0.9282 | valid_loss=0.3740 acc=0.8476
[CNN] Epoch 05 | train_loss=0.1231 acc=0.9605 | valid_loss=0.4028 acc=0.8556
[CNN] Epoch 06 | train_loss=0.0700 acc=0.9822 | valid_loss=0.4521 acc=0.8516
[CNN] Epoch 07 | train_loss=0.0340 acc=0.9945 | valid_loss=0.5002 acc=0.8548
[CNN] Epoch 08 | train_loss=0.0148 acc=0.9992 | valid_loss=0.5521 acc=0.8580
[CNN] Epoch 09 | train_loss=0.0063 acc=0.9999 | valid_loss=0.6019 acc=0.8572
[CNN] Epoch 10 | train_loss=0.0031 acc=1.0000 | valid_loss=0.6410 acc=0.8556
Loaded best CNN model (val_acc=0.8580).
```

```
=====
[CNN] Test Loss: 0.5811 | Test Accuracy: 0.8463
=====
```


```
In []: import matplotlib.pyplot as plt
```

```
plt.figure()
plt.plot(history["train_loss"], label="train_loss")
plt.plot(history["valid_loss"], label="valid_loss")
plt.xlabel("epoch"); plt.ylabel("loss"); plt.legend(); plt.title("Loss Curves")
plt.show()

plt.figure()
plt.plot(history["train_acc"], label="train_acc")
plt.plot(history["valid_acc"], label="valid_acc")
plt.xlabel("epoch"); plt.ylabel("accuracy"); plt.legend(); plt.title("Accuracy Curves")
plt.show()
```



Accuracy Curves



In [ ]: