

Linear Regression

AI, machine learning, and deep learning

- Example: Linear regression
 - goal = build a system that can a vector $x \in \mathbb{R}^p$ as input and predict the value of a scalar $y \in \mathbb{R}$ as its output
 - We define the output to be

$$\hat{y} = \boldsymbol{\beta}^T \boldsymbol{x} = \beta_0 + \sum_{i=1}^p \beta_i x_i$$

- Task T : to predict y from \boldsymbol{x} by outputting $\hat{y} = \boldsymbol{\beta}^T \boldsymbol{x}$

AI, machine learning, and deep learning

- Example: Linear regression
 - Performance measure P : to compute the MSE on the test set

$$\text{MSE}_{\text{test}} = \frac{1}{n} \sum_i \left(\hat{y}_i^{(\text{test})} - y_i^{(\text{test})} \right)^2$$

- How to gain experience by observing $(\mathbf{X}^{(\text{train})}, \mathbf{y}^{(\text{train})})$? → minimize the MSE on the training set, $\text{MSE}_{\text{train}}$

AI, machine learning, and deep learning

- Example: Linear regression
 - Suppose that the training set $\{x_{i1}, x_{i2}, \dots, x_{ip}, y_i\}$, $i = 1, \dots, n$ is given.

$$\begin{aligned} & \min \frac{1}{2} \sum_{i=1}^n (y_i - \hat{y}_i)^2 \\ &= \frac{1}{2} \sum_{i=1}^n \left(y_i - (\beta_0 + \beta_1 x_{i1} + \beta_2 x_{i2} + \dots + \beta_p x_{ip}) \right)^2 \end{aligned}$$

AI, machine learning, and deep learning

- Example: Linear regression

- In vector/matrix notation,

$$y_1 = \beta_0 + \beta_1 x_{11} + \beta_2 x_{12} + \cdots \beta_p x_{1p} + \epsilon_1$$

$$y_2 = \beta_0 + \beta_1 x_{21} + \beta_2 x_{22} + \cdots \beta_p x_{2p} + \epsilon_2$$

⋮

$$y_n = \beta_0 + \beta_1 x_{n1} + \beta_2 x_{n2} + \cdots \beta_p x_{np} + \epsilon_n$$



$$\mathbf{y} = \begin{bmatrix} y_1 \\ y_2 \\ \vdots \\ y_n \end{bmatrix} = \begin{bmatrix} 1 & x_{11} & \cdots & x_{1p} \\ 1 & x_{21} & \cdots & x_{2p} \\ \vdots & \vdots & \ddots & \vdots \\ 1 & x_{n1} & \cdots & x_{np} \end{bmatrix} \begin{bmatrix} \beta_0 \\ \beta_1 \\ \vdots \\ \beta_p \end{bmatrix} + \begin{bmatrix} \epsilon_1 \\ \epsilon_2 \\ \vdots \\ \epsilon_n \end{bmatrix} = \mathbf{X}\boldsymbol{\beta} + \boldsymbol{\epsilon}$$

where

$$\mathbf{y} = \begin{bmatrix} y_1 \\ y_2 \\ \vdots \\ y_n \end{bmatrix}$$

$$\mathbf{X} = \begin{bmatrix} 1 & x_{11} & \cdots & x_{1p} \\ 1 & x_{21} & \cdots & x_{2p} \\ \vdots & \vdots & \ddots & \vdots \\ 1 & x_{n1} & \cdots & x_{np} \end{bmatrix}$$

$$\boldsymbol{\beta} = \begin{bmatrix} \beta_0 \\ \beta_1 \\ \vdots \\ \beta_p \end{bmatrix}$$

$$\boldsymbol{\epsilon} = \begin{bmatrix} \epsilon_1 \\ \epsilon_2 \\ \vdots \\ \epsilon_n \end{bmatrix}$$

AI, machine learning, and deep learning

- Example: Linear regression

$$\begin{aligned} \min & \frac{1}{2} \sum_{i=1}^n (y_i - \hat{y}_i)^2 \\ = & \frac{1}{2} \sum_{i=1}^n \left(y_i - (\beta_0 + \beta_1 x_{i1} + \beta_2 x_{i2} + \cdots + \beta_p x_{ip}) \right)^2 \end{aligned}$$



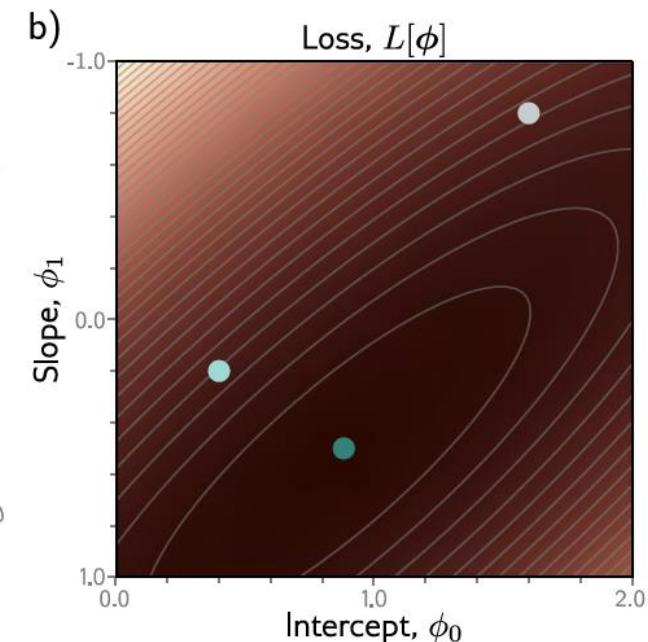
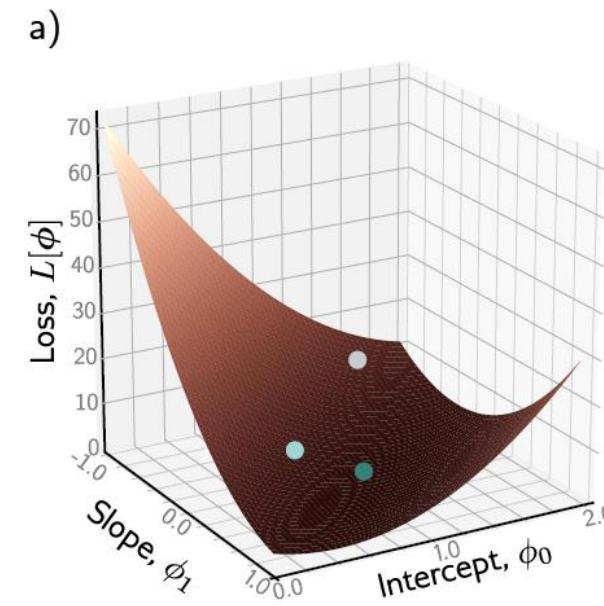
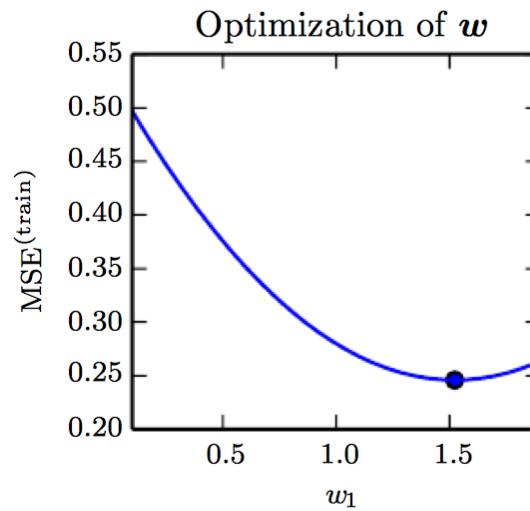
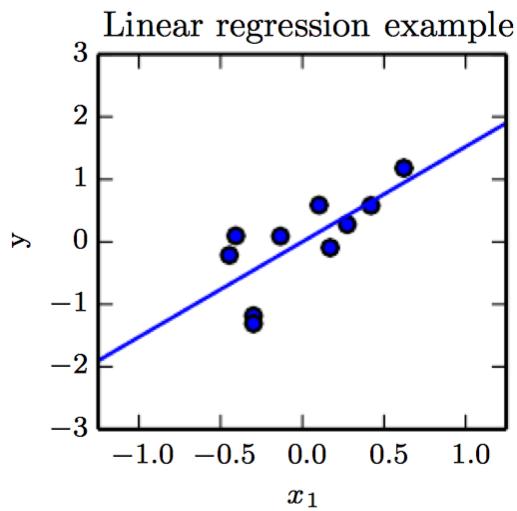
$$\begin{aligned} \min E &= \frac{1}{2} (\mathbf{y} - \mathbf{X}\boldsymbol{\beta})^T (\mathbf{y} - \mathbf{X}\boldsymbol{\beta}) \\ \Rightarrow \frac{\partial E}{\partial \boldsymbol{\beta}} &= -\mathbf{X}^T (\mathbf{y} - \mathbf{X}\boldsymbol{\beta}) = 0 \\ \Rightarrow -\mathbf{X}^T \mathbf{y} + \mathbf{X}^T \mathbf{X} \boldsymbol{\beta} &= 0 \end{aligned}$$

$$\boxed{\boldsymbol{\beta} = (\mathbf{X}^T \mathbf{X})^{-1} \mathbf{X}^T \mathbf{y}}$$

AI, machine learning, and deep learning

- Example: Linear regression
 - To minimize $\text{MSE}_{\text{train}}$, we can simply solve for where its gradient is **0**.

$$\nabla_{\beta} \text{MSE}_{\text{train}} = 0 \rightarrow \beta = (X^T X)^{-1} X^T y \quad \text{closed-form solution}$$



Reading assignment

- “Understanding DL” book
 - Chapter 2. Supervised Learning